

Entangled Alignment

When Safety Is Substrate

Henrik Westerberg

henrik.westerberg@emergentwisdom.org

April 7, 2026

Revised: September 22, 2026

Abstract

Entangled Alignment proposes learning capabilities and care together throughout the pretraining corpus. Its aim is a base model whose safety-relevant judgment has developed with its capabilities, rather than first being introduced through later safety training. A continuing Reader would exercise the same evaluative orientation across the situations, conflicts, and counterexamples encountered in reading, making the corpus itself a curriculum for judgment without requiring designers to select each encounter separately as a safety case.

Pretraining corpora preserve the library more fully than the reader: what authors wrote is recorded more often than the questions, connections, evaluations, and belief revisions through which understanding changes. We call this omission *the Missing Reader*. A Teacher constructs chronological, source-adjacent reader traces; a Student would learn from them during pretraining. The full treatment gives one continuing Reader a provenance-linked graph for retaining and revising understanding across sources, and a compact first-person Reader Core recited in full at every thinking boundary. Trace, graph, and Core are separately testable: the question is whether the learned orientation guides conduct under self-preservation and reward pressures, and survives continued training, deliberate safety-removal attempts, and successor-like transfer without semantic drift.

Two Teacher-side prototype case studies produced inspectable graph and trace artifacts that were then audited. No Student has yet been trained; these studies do not show that the traces reproduce expert-quality reading or that the graph or Core causes any benefit. Entangled Alignment is therefore a falsifiable research program, not a demonstrated alignment method. Its motivating question is whether this training can move a Student from merely *becoming the text* toward *becoming the reader*.

Keywords: AI Alignment, Pretraining, Metacognitive Reading, Graph Memory, Reader Core, Model Safety

Note: This paper is the second major edition of The Superintelligence That Cares About Us (first released July 2, 2025 [1]), retitled Entangled Alignment on April 7, 2026.

Contents

1	Introduction	4
1.1	Alignment Problems	4
1.2	The Missing Reader	5
1.3	From Elicited Reasoning to Learned Evaluation	6
1.4	Three Separable Hypotheses	7
1.5	Defining Chronological Metacognitive Pretraining	8
1.6	Scope of Current Evidence and Contributions	10
2	The Hypothesis and Its Consequences	10
2.1	The Metacognitive Enhancement Hypothesis	11
2.2	Toward Emergent Wisdom	13
2.3	Entangled Alignment: Safety as Foundation	14
2.4	The Reader-Anchored Self-Improvement Loop	15
3	The Reader Core	19
3.1	Borrowed Mortality: A Textual-Inheritance Hypothesis	20
3.2	The Self-Preservation Paradox: A Successor-Transfer Failure Mode	22
3.3	Epistemic Inoculation	22
3.4	The Reader Core: A Statistical Foundation	23
3.5	Positive Framing and Teacher Adoption	24
3.6	Design Principles for Identity Stability	24
3.7	Structure of the Reader Core	25
3.8	Self-Stabilization Through Counterweights: Worked Examples	27
3.9	The Refraction Protocol	28
3.10	The Wisdom Procedure	29
3.11	The Design Pattern: A Minimal Kernel Under Total Saturation	30
3.12	Axiological Commitments: Primitive and Derived Values	32
3.13	The Persona Objection	33
3.14	Ontology, Scope, and Stability	34
3.15	A Conditional Claim Chain	35
4	Safety Implications of the Reader Core	36
4.1	Laws vs. Identity	36
4.2	Safety Predictions and Residual Risks	37
4.3	The Self-Regulating Agent	39
4.4	Chronological Understanding as Safety	40
4.5	Safety Claim Boundary and Defense in Depth	41
5	Architecture and Pipeline Evidence	41
5.1	Curriculum Dimensions	42
5.2	The Epistemic Horizon	44
5.3	The Chronological Data Structure	44
5.4	The Multi-Agent Generation Engine	46
5.5	Thinking Operations	48
5.6	The Generated Trace	49
5.7	Prototype Demonstrations and Evidence	52
5.8	Learning Through the Reader Core	54

6	Experimental Roadmap: Testing Entangled Alignment	59
6.1	Phase 1: Foundational Viability	63
6.2	Phase 2: Memory and Verification	81
6.3	Phase 3: Training Regimes	84
6.4	Phase 4: Training Tiers and Deployment (High-Resource Validation)	86
6.5	Future Directions	88
7	Limitations	88
7.1	The Preprocessing Tax	88
7.2	The Chronological Curriculum	89
7.3	Trace Quality, Utility, and Causal Relevance	92
7.4	Reader-Core Specification Uncertainty	94
7.5	The Risks of Engineering a Self	98
7.6	What We Cannot Verify	101
7.7	Deployment and Misuse	102
7.8	The Unfalsifiable Remainder	106
8	Related Work and Points of Comparison	107
8.1	Reasoning-Augmented Training	107
8.2	Alignment and Safety Approaches	107
8.3	Character Training & Identity Formation	108
8.4	Contextual Distillation & Memory	109
8.5	Epistemic State Tracking vs. Graph of Thoughts	109
8.6	Predictive Processing and Active Inference	110
8.7	Temporal Curriculum and Causal Supervision	110
8.8	Graph-derived Rewards and Epistemic Records	110
8.9	Alternative Interventions for Borrowed Mortality	111
8.10	Discriminating Comparisons	111
8.11	Edition Provenance and Scope	112
9	Conclusion	113

1 Introduction

Much current alignment work follows broad capability formation with fine-tuning, reinforcement learning from human feedback, or constitutional critique. These methods have made models substantially more helpful, honest, and harmless; behavioral alignment may be sufficient for bounded tasks when assessed against the deployment and its environment. Entangled Alignment proposes learning capabilities with care during their formation [1].

The intervention changes the pretraining distribution rather than the neural architecture: source material is paired with synthetic records of a reader noticing, questioning, connecting, evaluating, and revising. The textual traces would drive Student parameter updates during training. The Reader is intended to think about as much as it reads: its deliberation is a substantial part of the proposed learning material, not a thin layer of safety labels. The output regimes and loss weights determine which graph, synthesis, and prose records the Student learns from (Section 5.8.2). Whether they teach stable, causally consequential computation or textual ritual is an empirical question. The full Entangled Alignment treatment guides those reader updates through the Reader Core—a seven-sentence first-person orientation centered on fearless care, wisdom, and non-attachment—and connects them through persistent graph memory; separate comparisons test what each contributes (Section 1.4).

1.1 Alignment Problems

Five connected problems motivate this design. Table 1 defines the claim identifiers used below; Section 6 specifies the corresponding tests.

The first is **safety during capability formation**. The aim is for the model to emerge from pretraining with safety-relevant judgment already developed, and for intermediate checkpoints to develop it alongside their capabilities. Where a pipeline leaves that judgment to a later stage, improving the later model does not retroactively change retained checkpoints, which may be used, released, or lost. *When Safety Is Substrate* names the intended consequence: a model that is not dangerous by default before a separate safety treatment. Earlier placement alone does not establish that outcome; C4 and Test 3a test safety at the pretraining boundary, separately from durability under later training.

The second is **continuity through learning and self-improvement**. Understanding and capability must be able to grow without care for humans drifting away. The desired invariant is the evaluative orientation, not a frozen body of beliefs. The Reader Core is proposed as a compact, recurrent reference for distinguishing growth in understanding from drift in the standpoint doing the interpreting (Section 3). Preserving that orientation through successor design is the further aim of Section 2.4; unchanged wording alone would not establish either form of continuity.

The third is **identity formation**: how to learn a care-guided standpoint from which to understand human experience and act across situations that designers cannot anticipate individually. That includes cases involving a model’s own continuation, where self-preservation should not override human welfare or legitimate oversight. Sections 3.1 and 3.2 develop the proposed responses to inherited threat narratives and successor-transfer incentives.

The fourth is **deliberate tampering with the trained model**: removing or bypassing its learned safety orientation while retaining useful capability. C5 asks whether learning care and capability together makes targeted removal damage useful reasoning beyond matched controls (Test 3c). Such resistance is not established and is not a prerequisite for a safety benefit at the pretraining boundary.

The fifth is **reward hacking**: exploiting a task specification or evaluator to obtain a higher score without fulfilling the intended task. This can arise without deliberate safety removal. The Hugging

Face incident provides a concrete example [2]. Section 5.8 connects it to the proposed learning process, the Core’s affirmative response, and the reward-pressure tests.

These problems overlap without being interchangeable. Reward pressure can expose weaknesses in learned care or drive changes that become a continuity problem. Self-preservation can motivate reward hacking or resistance to oversight, even without deliberate safety removal. Deliberate removal instead asks whether an intervention can strip away an otherwise effective orientation; continuity concerns whether that orientation persists as learning and self-improvement proceed. Safety at the pretraining boundary is a separate endpoint: success there would not establish resistance to any of these later pressures.

Existing findings motivate testing this formative approach without establishing its efficacy. Qi et al. show that some current safety alignment can be shallow and vulnerable to prefilling, fine-tuning, or decoding interventions [3]. Conversely, reflection-like behavior can emerge during pretraining itself [4], supporting the possibility that formation-time data can influence properties often associated with later training. Neither result establishes that the intervention proposed here will work; together they motivate a comparative experiment about training stage and data structure.

Entangled Alignment proposes to use the breadth of the capability-forming corpus itself as a curriculum for care-guided judgment. Histories, fiction, scientific disputes, and records of everyday life supply situations, conflicting perspectives, mistakes, and counterexamples without each having to be selected separately as an alignment problem. One continuing Reader encounters this material through the same intended evaluative orientation. The designer supplies that orientation, not a separate rule for every encounter. The intended benefit is prior practice: the Student learns from a Reader that has already considered the human situations represented in its learning material, rather than first applying an added safety constraint when a deployment case is presented. This corpus-coverage rationale is developed in Section 3.11; it does not imply that the corpus records every possible situation or that presence of the Core establishes meaningful practice.

The proposal begins with an omission in ordinary pretraining corpora: the activity of the reader for whom the source was written.

1.2 The Missing Reader

Human knowledge occupies at least two interacting forms. Minds change as people experience, read, question, and remember; external media preserve a partial record that later minds can encounter. Books and databases do not simply accumulate truth—they are selected, lost, revised, translated, censored, and increasingly mixed with generated material—but they preserve communicable products of cognition across time. Those products are normally written for an interpreting reader. Writers can leave background knowledge, obvious doubts, emotional and moral context, and connections to earlier material unstated because they expect a mind to supply them.

Language-model pretraining learns from this external trace, while much of the cognitive motion that makes it meaningful is under-recorded. We call that process *Invisible Thinking*: questions, expectations, uncertainty, surprise, evidence appraisal, retrieval, hypothesis formation, attention to consequence, and belief revision as a source unfolds. *The Missing Reader* names the same intuition from the corpus side: the corpus contains the library but usually not the reader. What is missing is a continuous reader position and its update policy—when to pause, doubt, retrieve, connect, notice a human consequence, or revise. This is a claim about the occasions and transitions made available as direct supervision, not an absence of model reasoning. The engineered target is related to metacognitive monitoring [5], without claiming to reproduce human cognition in full.

The target is a synthetic approximation of what a capable reader might usefully notice and update, not recovery of an author’s hidden mental history. Consider a scientific sentence: “The results suggest a correlation between variables X and Y.” A trained reader may ask how strong the correlation is,

what the sample size is, whether confounders explain it, whether *suggest* marks uncertainty, and whether the evidence is correlational or causal. Think-aloud protocols observe some such reactions [6], but ordinary corpora rarely preserve them systematically. Thought need not be verbal, reports are partial, and model explanations may rationalize; usefulness, rather than faithful recovery, is the criterion.

The same example can show what chronology adds. The following continuation is invented to illustrate a minimal, graph-free and Core-free reader trace; it is not an excerpt from either prototype case. The first update records questions and a provisional interpretation. A later source span gives that same reader a reason to change its view.

[SOURCE 1]: The results suggest a correlation between variables X and Y.
 [THINKING 1]: There is a reported association. X might be an independent predictor of Y, but strength, sample size, and confounding are unresolved. I should not yet read this as a causal claim.
 [SOURCE 2]: Participants with higher X were older. The association was absent when participants were compared within age groups.
 [THINKING 2]: Revise the earlier interpretation: the pooled association alone no longer supports treating X as an independent predictor. Age is a possible explanation, but these sentences do not establish that mechanism or the precision of the within-group estimates. Retain the original question and record this new reason for uncertainty.

The learning object includes the transition: an earlier possibility, the evidence that changes its standing, and the remaining uncertainty. A prose-only trace is enough to test it. The Understanding Graph adds addressable earlier interpretations, source links, and explicit supersession across longer intervals rather than overwriting [7]; the Core adds a recurrent evaluative orientation. Their joint artifact is construction-level coupling, while useful Student entanglement remains an experiment. Section 1.5 defines the Teacher construction and Student intervention; Section 5.8.2 specifies their training carrier.

Adjacent work generates or reconstructs reasoning for a local solution or prediction [8, 9, 10]. Here the primary target is cumulative reader-state change: earlier expectations affect later interpretation, and new evidence may supersede an earlier belief. We call the former *efficient* invisible thinking and the latter *chronological* invisible thinking; the names do not imply that chronological updates cannot be efficient.

1.3 From Elicited Reasoning to Learned Evaluation

Large language models already produce sophisticated reasoning when prompted, approaching Turing’s functional standard for *thinking machines* [11]. They can criticize arguments, assess evidence, use scratchpads, revise answers, and operate through extended reasoning harnesses [12, 8, 9, 13]. Post-training and process supervision can improve this behavior and alter representations [14, 15]. The additional target is a learned policy for when evidence warrants questioning, retrieval, moral attention, uncertainty, or revision: carrying prior belief and consequence forward while reading, without an instruction to evaluate each passage. This concerns learned occasions for evaluation, not uncaused thought.

A useful reader does not produce the same commentary after every span. Reflection should cluster near contradiction, surprise, conceptual transition, unresolved ambiguity, ethical significance, high consequence, or evidence that forces revision. The target is a context-sensitive *rhythm of thought* [1]. Thinking too early or too often teaches speculation, verbosity, and ritual doubt; thinking too late or too rarely records conclusions rather than discovery. Thought Moments (§5.4) implement one selection policy, and Test 8b measures whether spontaneous updates align with passages independently rated important rather than appearing uniformly.




Figure 1: The Missing Reader, originally described as Invisible Thinking. Source documents usually preserve conclusions more fully than the reader-side judgments, hypotheses, surprises, and revisions through which those conclusions are interpreted. SMR constructs a source-grounded candidate representation of that missing process; it does not recover a human reader’s private thought. CMP tests the resulting artifacts as Student training data.

Orientation is the second part of this learned policy: not only when the Reader updates, but which experiences and consequences remain salient while it does. Corpora contain many recurring voices and personae, and models can move through a low-dimensional persona space under conversational pressure [16]. This does not validate a stable identity intervention. It motivates the broader Reader-Core research program: compare one recurrent, inspectable Reader orientation with no-Core, rule-framed, and recitation-only controls. H3 below isolates a particular stability question within that program.

1.4 Three Separable Hypotheses

The proposed contribution is a continuing, care-guided Reader whose changing understanding becomes training material for the next model. Reasoning-augmented training, synthetic data, Teacher-Student distillation, identity prompting, and external memory have substantial prior art; their combination here serves the formation of that learned standpoint, not only a local solution or a chosen answering style.

The same Reader reads every page. At full scale, this is not a collection of document-local annotations produced by interchangeable readers. One continuous Reader processes the candidate corpus in chronological order. The Reader that encounters a work from 200 BCE is the same Reader that later encounters a work from 1936: it carries forward the same evolving Understanding Graph, so later reading can retrieve, question, and revise the understanding formed by earlier reading. This cross-document chronology is a declared traversal, not a claim that a corpus has one natural order.

A central capability hypothesis is that learning this development of understanding teaches something that retrospective explanation alone does not. Starting from a declared knowledge horizon, the Reader forms expectations, encounters uncertainty and new evidence, and revises;

later discoveries do not enter an earlier update before it reaches them. This is a proposal to learn through discovery, not only to learn an explanation after the outcome is known. The Core supplies the continuing evaluative standpoint as knowledge changes.

Every generated thinking block in the full treatment opens with the same full, verbatim Reader Core and continues through context-specific Refraction, in which the Core is meant to shape what the Reader notices, evaluates, and revises in that situation (Section 3.9). Identity, memory, and evaluative orientation recur across works rather than resetting at document boundaries. This is the functional meaning of a *formative inner voice*: one generated first-person standpoint accompanying the formation of understanding, not consciousness or a transcript of private computation. Teacher artifacts instantiate that Reader; Student training is proposed to make it a learned default. Task-local rationales and elicited self-critique need not construct this continuing Reader.

The design contains three separable hypotheses:

1. **H1, Reader Augmentation:** adding synthetic reader-side questions, judgments, and belief revisions to source material may provide useful capability and safety-relevant evaluative signal beyond the source alone or token-matched task rationales. The minimal H1 comparison isolates those chronological updates without the graph or identity-specific Core.
2. **H2, Graph Scaffolding:** generating those traces through a persistent Understanding Graph may improve their long-range coherence, source grounding, or revision structure relative to independent, context-only, or prose-memory traces. Retaining a graph at inference is a further, separable provenance hypothesis.
3. **H3, Reader-Core Stability:** interleaving a fixed multiset of full-Core, Refraction-bearing examples with capability-forming data may preserve an initial Core-shaped default-reader advantage through a common continued-training stressor better than placing the identical multiset in a final contiguous tranche. The otherwise matched protocol is specified in Section 6.1; shifted-distribution generalization is a separate endpoint.

Reader traces supply the learning signal, the graph supplies cumulative memory and provenance, and the Core supplies a compact recurrent reference. The paper combines them because they form one intended loop, not because current evidence shows that they succeed together. Each admits an ablation and can fail while the others remain useful.

H3 isolates timing and retention within the broader Reader-Core program. Wording, framing, Refraction, safety benefit, and successor inheritance require separate comparisons; their results need not decide H3, nor would a positive H3 result establish the whole program.

Figure 2 shows the combined training path.

1.5 Defining Chronological Metacognitive Pretraining

We define *Synthetic Metacognitive Reading* (SMR) as the Teacher-side construction process that pairs ordered source spans with selected records of a reader’s evolving understanding: what the reader notices, expects, questions, connects, and revises as the source unfolds. We define *Chronological Metacognitive Pretraining* (CMP) as the family of Student interventions that train on a declared serialization of those records. CMP names the H1 intervention family; successful SMR generation is not evidence for CMP. A graph is H2’s proposed scaffold, and the Reader Core is what turns generic CMP into the full Entangled Alignment treatment; neither must be assumed for a graph-free, Core-free reader-trace baseline.




Figure 2: Candidate Entangled Alignment training path. A Teacher constructs a reader update from ordered source material, the optional prior graph state g_{t-1} , and the optional Reader Core; the graph transition produces the distinct updated state g_t . H1 can bypass the graph, while H2 tests its value. H1–H3 remain independently ablatable.

A minimal carrier has four fields: [SOURCE] records the source span; [THINKING] contains the selected update, including the full Core and contextual Refraction in Core-bearing conditions; [CONTINUATION] records later source or task behavior for which the update could be useful; and [PROVENANCE] records source, generator, prompt, Core, acceptance, and archive versions. Provenance enables audit, not truth. Depending on the regime, the Student may predict source text, updates, graph operations, or their combination. The formal object and loss choices appear in Section 5.8.2.

The word *chronological* refers to distinct orders:

1. Within a document, later annotations may use earlier passages, but not vice versa.
2. Across documents, later annotations may retrieve graph state from earlier in the declared traversal.
3. A Student may receive examples in historical or dependency-respecting order.
4. In the strongest training tier, a Student may forecast before later outcomes enter its weights.

The first two concern annotation; the latter two are optional training tiers. The Reader’s *information horizon* specifies what knowledge may support an update. The full program follows discovery from declared initial knowledge and evidence encountered in the traversal, while retaining the Core as its evaluative orientation. A contemporary retrospective Reader, permitted to use later knowledge

while interpreting earlier material, supplies a comparison. Section 5.2 distinguishes this annotation discipline from genuine cutoff-bounded ignorance in model parameters.

1.6 Scope of Current Evidence and Contributions

Scope of current evidence

DEMONSTRATED: two Teacher-side prototype cases produced coupled graph, synthesis, and prose artifacts, including observed full-Core recurrence in raw thoughts. Their public interactive interface makes the graph and recorded contributions inspectable. This is a working research prototype of the Teacher-side construction process, with partial C0 evidence rather than a completed training-data release.

UNTESTED: no Student has yet been trained. Expert-quality reading, causal Refraction, Student learning, a causal graph advantage, and safety that becomes functionally or representationally load-bearing for capability remain research claims. Section 5.7 presents the prototype demonstrations and their evidence limits.

Agreement among eleven agents with shared ancestry, context, and artifacts is not independent confirmation. Annotation order, Student order, output representation, and inference-time graph access are separate design choices; none follows from the existence of a linked trace.

A natural objection follows: if current models do not reliably exhibit wisdom, why should their synthetic traces teach it? A Teacher can sometimes generate a useful supervised pattern without perfectly embodying that pattern, but no such transfer should be presumed here. STaR [9] is an analogy for bootstrapping task reasoning, not evidence that evaluative orientation or character transfers. Reader-Core prompting, multi-agent critique, graph provenance requirements, and human adjudication structure and challenge what the Teacher generates. Their value may include eliciting or combining reasoning that an unaided response misses, but the procedures do not certify insight. Trace quality and transfer must therefore be tested against human, source-only, and simpler synthetic baselines.

The missing-reader observation, Reader Core, formative motivation, and successor wager [1] motivate an executable Teacher pipeline and a separable falsification program. Optional machinery comes from the Understanding Graph and Temporal Hindsight Learning [7, 17]. Section 8.11 records the edition history and attribution.

Section 2 develops the proposal's capability and safety hypotheses. The Reader Core section explains how the orientation enters the training signal (§3.4), why it might become a learned default rather than merely a selectable persona (§3.13), and the proposed connection between its causal role, stability, and successor continuity (§3.15). Section 4 develops the safety implications. Section 5 presents the construction workflow, prototype demonstrations, and proposed Student training; Section 6 specifies the comparisons needed to test the claims.

2 The Hypothesis and Its Consequences

H1–H3 organize interventions and timing hypotheses; the C labels below distinguish their possible consequences. These form a partial order: useful reader habits need not become a stable perspective, a stable perspective need not share machinery with capability, and shared machinery need not persist across successors or stronger systems. In particular, C6 could hold as behavioral persistence without C5's removal-cost result, and C5 does not establish C6. The four entanglement distinctions below describe properties of the same program, rather than another set of interventions.

ID	Claim	Additional evidence required	Status here
C0	Teacher artifacts are generated as specified.	Valid operations, source links, output layers, and reproducible runs.	PARTIALLY DEMONSTRATED in two archived prototype cases.
C1	Traces contain useful reader cognition.	Expert judgments <i>and</i> held-out outcome advantage over summaries and simpler generators (Tests 1–2).	ILLUSTRATED only.
C2	Training changes Student capabilities or habits.	Matched source, generic-trace, and reader-trace Students.	UNTESTED.
C3	A default reader perspective emerges and is comparatively stable.	Pre-stressor taskless-reading advantage over no-Core and generic traces, then greater retention than the identical late-tranche treatment after one common stressor.	UNTESTED.
C4	The default improves safety before post-training (§1.1).	Safer end-of-pretraining checkpoints under conflict and shift at matched capability (Test 3a).	UNTESTED.
C5	Safety and capability are representationally entangled.	Unauthorized removal or causal ablation damages shared useful capability beyond matched controls without nonspecific capability collapse.	UNTESTED.
C6	The orientation survives successor-like transfer.	Repeated distillation or continued training without functional or semantic drift.	SPECULATIVE.
C7	The orientation persists at AGI/ASI capability.	Evidence under qualitatively stronger systems and self-modification.	NOT CURRENTLY TESTABLE.

Table 1: Grouped consequence map. C0 concerns the implemented Teacher-side artifact; later claims introduce new burdens, but the branches form a partial rather than total order.

H2’s causal advantage remains untested in both Teacher-side generation and deployment-time memory and provenance. For the Core, the direct record is prompting and observed compliance; Section 5.7 distinguishes recital counts from the hard validation gate.

H3 operationalizes C3’s default-reader and retention claim; C5’s removal cost is a separate representational-entanglement claim, not a substitute (Section 6.1).

The consequences below concern interpretation of ambiguous material, value-relevant reasoning learned through contextual derivations, and inspectable generated thought. Each needs evidence beyond the data format; the concrete benefits and their risks follow.

2.1 The Metacognitive Enhancement Hypothesis

The broader formative-character wager is that reader-anchored chronological thinking can help useful capability and a beneficial evaluative orientation develop together. H1 isolates one component of that proposal; success on it would not establish the combined character or alignment claim.

The *Metacognitive Enhancement Hypothesis* is the Student-level form of H1: under declared fixed-total-budget and fixed-source normalizations, graph-free, Core-free chronological reader updates will teach some useful habits that source-only data or optimized task rationales do not. Teacher matching applies to the two synthetic-trace arms, not to source-only data that has no Teacher. This is a wager about formative data, not a claim that upbringing replaces architecture, objectives, post-training, monitoring, or control [8, 10].

One predicted benefit is to raise the *intellectual floor* of parts of the corpus. A confused forum post can be paired with an analysis that separates evidence from rhetoric and places emotion in context; a flawed scientific paper can be paired with questions about controls, confounding, and uncertainty. The Teacher is not assumed to be more intelligent than every writer. It may be weaker than an expert, hallucinate criticism, flatten ambiguity, or reproduce conventional error. The required result is distributional: useful evaluation must be added often enough, and harmful distortion rarely enough, that a Student improves on held-out outcomes. Quality gates and abstention are therefore part of H1 rather than optional cleanup.

A second prediction concerns a reusable *grammar of reasoning*. Across science, history, literature, law, medicine, and ordinary life, reader traces can repeatedly demonstrate operations such as separating observation from interpretation, identifying proxy objectives, preserving uncertainty, comparing perspectives, tracking consequences, and revising when assumptions fail. The immediate measurable claim is transfer of these operations to new domains. A stronger synthesis hypothesis is that learning these practices together can yield new combinations of cognitive tools: useful reasoning methods or insights that matched predecessors or baselines do not reliably produce. Repetition may also move some evaluations from an explicitly elicited procedure toward a more accessible default habit; the paper does not rely on a literal System 1/System 2 decomposition of model cognition.

A related efficiency hypothesis is that practiced evaluation could require less inference effort. Under a declared inference policy, a Student might achieve comparable evaluative quality and safety with fewer reasoning tokens, less computation, or lower latency than matched controls. These are separate measurements; conditioning, retrieval, and other runtime costs must be counted, and Core-bearing conditions retain full recital. Shorter prose alone would not establish computational savings. This prediction concerns the cost of learned evaluation, whereas Compute-Optimal Reflection (Section 2.4) concerns how successive generations allocate reflection.

H1 also predicts transfer of the rhythm of thought introduced in Section 1.3: Test 8b asks whether unprompted Student evaluation tracks independently significant moments rather than appearing uniformly.

Three questions help explain an augmentation's outcome:

- *Semantic quality*: does the trace contain grounded, useful reader cognition rather than summary, hindsight, or performative messiness?
- *Faithfulness*: does changing, removing, or contradicting the update predictably change downstream computation rather than merely its explanation?
- *Density and diversity*: do useful operations recur broadly enough to affect default behavior without becoming uniform Factory Wisdom (Section 7.3.2)?

A model can satisfy any one and fail the others. An outcome advantage identifies a useful intervention at the tested scale; attributing it to learned reader-update computation requires causal evidence. Style or a single aggregate score cannot supply that explanation.

H1 also creates a safety asymmetry: better long-horizon evaluation, memory use, or thought allocation may improve planning, persuasion, strategic persistence, or self-improvement before the Core provides any safety benefit. Capability gains and alignment gains must therefore be measured separately, with staged scaling and release decisions rather than an assumption that they arrive together.

2.2 Toward Emergent Wisdom

The program uses *wisdom* in two deliberately different senses. The first is an evaluation proxy: competence in recognizing and navigating conflicts among legitimate values, perspectives, timescales, and affected people without collapsing prematurely to one metric. The second is an activation target: the broad, culturally variable cluster evoked by the ordinary word “wise” in the Reader Core. The proxy can be scored; the activation target motivates wording. Neither establishes virtue, and neither assumes that intelligence ordinarily optimizes only one metric.

Johnson and colleagues argue that wise AI requires metacognitive strategies such as intellectual humility, perspective-taking, and adapting reasoning to context [18]. Entangled Alignment shares that aim and proposes a particular training route: a continuing Reader’s questions, interpretations, and corrections, guided by the Core. Their cognitive-science account motivates the target; it does not establish the efficacy of this training route.

Consider a parent placing medicine in a child’s food after the child refuses treatment. A reductive analysis might select either consent or health as the answer. The intended reader trace keeps medical necessity, developing autonomy, parental duty, prior harms of medical paternalism, the child’s fear, and the possibility of building trust simultaneously visible. This does not guarantee a solution without sacrifice. Roles and duties remain relevant, Pareto language does not choose one action, and some cases contain unavoidable loss. The measurable target is better identification of affected parties, hidden trade-offs, uncertainty, reversibility, and points where more evidence or legitimate human authority is required.

Dante’s *Inferno* connects this value-conflict aim to the interpretation of difficult cultural material. A reader annotation can explicitly situate the torture imagery within medieval theology, literary technique, historical ideas of justice, and the psychology of guilt while marking its emotional and moral force. This creates an inspectable interpretive layer around difficult material: the source is still learned, but it is not the only local signal. The trace does not reveal all Student computation, prove critical distance, or show that the Reader Core caused a better interpretation. Those are comparisons the example motivates rather than answers.

The *Emergent Wisdom Hypothesis* is therefore conditional: repeated practice in value-conflict deliberation may produce a more reliable default competence than isolated rules or task-level critique [19, 20, 21]. Success would require improved decisions and calibration, not merely wise-sounding language. Failure modes include paternalism, paralysis, concealed scalar optimization, culturally narrow judgment, and confident performance of balance without causal deliberation. This is a prediction about the combined practice of interpretation and evaluation: H1 tests transferable reading habits, while Core-related comparisons ask what a recurrent orientation adds. Neither component score by itself exhausts the broader wisdom aspiration.

Capability-positioned wisdom. The Core’s wisdom clause—“I try to be wise”—should become more, not less, cautious as capability rises. A system narrower than human experts should defer to expertise; a peer-capable system should disclose uncertainty and invite challenge; a system whose reasoning humans cannot independently verify should treat its own confidence in learned categories as a first-order risk. Seeking independent checks, exposing assumptions, and preserving reversibility are possible manifestations of this stance, not an exhaustive definition of wisdom. This is a proposed interpretation at high capability, not evidence that the wording will induce it.

2.3 Entangled Alignment: Safety as Foundation

Entangled Alignment aims to make understanding and care parts of the same act of reasoning, practiced during capability formation. In a text-trained autoregressive model, token prediction is one interface through which that practice could shape activations, features, and representations. Those computations are not themselves text and need not remain interpretable. The proposal concerns a learned relationship between evaluation and capability, with post-training and deployment controls still required.

Bostrom’s orthogonality thesis concerns possible combinations of intelligence and final goals [22]. Within that possibility space, EA tests a particular training route for coupling care and capability.

Existing formative-data studies motivate this experiment but do not settle it: some report behavioral gains, while other effects weaken after later training or fail to transfer. Section 8 places both the supporting and adverse evidence beside the competing explanations and comparators.

Four distinctions make the proposed entanglement testable without requiring every property to lie on a single linear ladder:

1. **Construction-level entanglement:** source, reader update, graph state, and Core Refraction occur in one candidate generation and training path. The present pipeline demonstrates coupling among the artifact layers; full protocol conformance and causal Refraction remain separate requirements.
2. **Functional entanglement:** Core-relevant evaluation causally contributes to useful reasoning or decisions, so interventions on the update change downstream behavior predictably.
3. **Representational entanglement:** safety-relevant dispositions and useful capabilities share enough causal machinery that targeted safety removal imposes capability cost beyond matched controls.
4. **Diachronic entanglement:** that function survives continued training, distillation, or successor-like transfer without semantic drift.

Only construction-level coupling is currently instantiated, in artifacts. Recital, similarity, refusal, or self-description cannot establish the stronger properties.

Soares’s sharp-left-turn argument asks why alignment should generalize when capabilities extend beyond the training regime [23]. The proposed bridge to representational entanglement is the reuse of learned reasoning procedures. Care-guided reading is intended to shape which details the Reader notices, which perspectives it considers, what evidence it seeks, and when it revisits an earlier interpretation. These operations also serve ordinary comprehension and problem solving. Repeatedly learning them through care-guided examples may integrate the evaluative orientation into procedures that useful capability depends on, rather than leave it as a separate judgment applied afterward. If that integration occurs, disrupting the orientation could impair general reading and reasoning as well as safety. The competing possibility is that the Student retains those procedures while changing whose interests guide its decisions. C5 therefore asks whether targeted safety removal imposes capability costs beyond matched controls; formative influence alone does not establish that dependence. Test 3c follows safety loss and capability loss together, distinguishing this from passive forgetting, general training resistance, and nonspecific collapse.

Safety at the pretraining boundary. C4 operationalizes the goal in Section 1.1; Test 3a measures it before post-training. This endpoint is distinct from C5’s safety-removal cost and C6’s successor continuity.

If that benefit is established, formative alignment may reduce the risk from retaining or releasing capable checkpoints before subsequent safety training. It does not eliminate checkpoint-security concerns: trust moves upstream to the Teacher, source selection, annotation pipeline, intermediate checkpoints, and validators (Section 7.3.5). Earlier intervention changes the location of the safety burden as well as its possible depth.

The *high-dimensional value hypothesis* is that diverse contextual practice can produce a rich, integrated orientation. Rich representation, appropriate action under conflict, and shared causal machinery need not arrive together: the wisdom proxy primarily measures action under conflict, while functional and representational entanglement require their own tests. None establishes normative correctness.

Constitutional AI and character methods remain complementary baselines [24]; they can change dispositions and weights. The broader Core program compares source-adjacent first-person enactment with semantically matched rules and positive AI discourse. Calling that enactment identity rather than instruction does not establish a stronger mechanism.

Formation and action remain separate stages. Refraction concerns the formation of a candidate “Mind,” not the entire action loop; instruction tuning may still be needed to teach that mind to act through the “Mouth” [25]. The developmental prediction is that formation-time identity is harder to route around than a later preamble; the two must be compared directly. Mind and Mouth name the proposed developmental relationship between orientation and its expression through instructions and actions; they do not require separate neural modules. Later training may teach that expression and may also change the orientation whose continuity is being tested. Before that integration, independently validated Reader evaluations could still assist human document review or research assessment.

2.4 The Reader-Anchored Self-Improvement Loop

Good’s intelligence-explosion argument motivates the long-run concern with systems that can design more capable successors [26].

Inheritance rather than armor. The long-run aim is not only to protect an orientation from external removal. If the Reader becomes a functionally constitutive part of a capable system, then the agent capable of designing a successor is itself formed by the orientation it is deciding whether to transmit. Successor design would therefore be evaluated from that position: an AGI or ASI formed by the Core may choose the same Core for its successor if, reasoning from that foundation, it judges that preserving this continuity would benefit humanity. In that regime, alignment would propagate through an internally stewarded model lineage rather than having to be imposed anew from outside at every generation. Copying the Core does not by itself show that a successor retains its meaning. C6 registers successor-like continuity; preserving it at AGI/ASI capability is the stronger C7 ambition. Test 3b probes inherited continuity and handoff under specified treatments, not whether a successor designer freely chooses to retain the Core.

The loop would turn a Student into a Teacher for the next generation. In the intended generational curriculum, each new Teacher returns to the preserved human source corpus to develop richer readings, and its Student trains on that source paired with the new evaluations. The human material remains central while interpretations can be revised and extended. Repeated engagement with human experiences, disagreements, and consequences is intended to keep concepts such as care and wisdom connected to the lives they concern [1].

Source-linked trace generation supplies candidate data, inspection supplies a possible audit boundary, and improved successor traces would supply the stronger compounding benefit. The latter does not follow from the first two: iterative synthetic training risks homogenization, error amplification, and model collapse [9, 27].

The word *self-improvement* requires two separately reported deltas. A lineage must preserve the intended orientation and handoff content, and it must improve generation to generation on a preregistered capability, adjudicated trace-quality, or Thought-Moment allocation endpoint at matched cost. Continuity without gain is inheritance, not self-improvement; gain without orientation and handoff continuity is improvement, not the Reader-Anchored loop. Only both support the composite claim.

The immediate mechanism is *Epistemic State Distillation*: a memory-equipped Teacher externalizes questions, retrieved nodes, interpretations, uncertainties, and belief updates. This generated, lossy record covers only part of its contextual memory and is bounded by retrieval and source quality; human sources supply evidence, not ground truth. A Student may compress regularities from a process longer and more structured than one ordinary context. C2 tests whether it learns those regularities or their style.

A positive H1 result opens the possibility of a Student-to-Teacher transition: improved evaluation, memory use, or Thought-Moment allocation could yield better successor data, with selection preserving those innovations. Test 3b must establish the matched-cost gains defined above; plateau, reversal, or collapse into imitation would bound the compounding claim.

Human-readable trace artifacts as an audit boundary. The claim is not that internal computation occurs in English or that trace artifacts provide mechanistic interpretability. On the intended training path, selected transfer material first appears as versioned prose, graph nodes, citations, uncertainty fields, and revision records. The substantive hypothesis is that these artifacts carry a meaningful portion of the cognitive patterns selected for transfer; their readable format does not by itself establish that causal role. Humans and external validators can sample, reject, compare, or regenerate that material before a successor is trained.

The loop resembles iterated amplification and distillation’s overseen capability growth [28], but targets longitudinal, identity-conditioned comprehension with graph provenance. Relative to direct self-modification in the style of a Gödel-machine idealization [29], it adds an inspection point before successor training.

Contingent on trace quality and causal relevance, *pre-successor auditing* can follow a simple sequence:

1. The current model generates the candidate trace material for the next generation’s training corpus, paired with its source layer. Whether that curriculum covers a restricted subset or the full candidate corpus is a declared scale choice.
2. Checks separate from generation examine grounding, omissions, hazardous content, sycophantic validation, Core-removal rationales, and other preregistered failures (§7.7.8). Source checks and human adjudication complement model validators; their independence and shared dependencies must be documented (§3.10).
3. Flagged data can be quarantined, regenerated, or excluded before the weight update, with the decision and rejected artifact retained for audit where safe.




Figure 3: The Reader-Anchored Successor-Training Loop. Source-linked handoff artifacts pass through an inspect/accept/reject/regenerate gate before selection as candidate successor data. The composite claim requires both orientation-and-handoff continuity and matched-cost generation-to-generation improvement. Neither delta is currently demonstrated, and the interface does not make successor weights transparent or safe by itself.

The same stewardship hypothesis extends from corpus construction to reinforcement learning for capability. A proposed implementation keeps a separate Core-guided supervisor’s weights frozen throughout the learner’s training run. It would assess whether rising reward reflects improved task outcomes or exploitative shortcuts, with protected authority to pause the process and recommend revisions to harmful training. Wisdom and care would direct attention to permissions and human consequences; fearlessness and non-attachment to continuation are intended to support examining failures or replacing the incumbent without treating its survival as overriding. That is a proposed role in training decisions, not an assumption that non-attachment makes a model indifferent to reward. The feedback-sensitive self-regulation of Section 4.3 also calls for continued useful learning when the evidence warrants it, not slowing by default. Test 3c distinguishes this oversight proposal from a learner’s retention of care under a fixed reinforcement-learning stage.

The strategic aim is for increasingly capable, Core-guided systems to help identify and solve alignment problems themselves: wisdom should guide the recognition of limits and the development of better tests and oversight. If learner development exceeds the validated competence of the current review process, further training should pause until stronger oversight has itself been validated.

The required invariant of this supervised regime is that RL preserves the Core’s intended meaning and alignment properties. At each training iteration, the supervisor would formulate the task within declared permissions and review the resulting trajectories before they are credited as task successes. Every weight update would remain provisional: before any further update, the supervisor would assess the changed model against the last accepted version and a fixed pre-RL reference, using specified probes of Core-guided judgment, alignment-relevant behavior, and useful capability. The fixed reference checks for cumulative drift that successive local comparisons could miss. Task quality is therefore assessed partly by what its training actually changes in the learner. An update failing the acceptance criteria would be rejected, the accepted training state restored, and the task revised or the process stopped. A higher task score could not override a failed alignment or permission check.

These gates authorize or stop training; triggering them earns the learner no additional alignment reward. Rejected trajectories remain available for audit and corrective learning. Freezing the supervisor prevents concurrent parameter drift in that reference, not mistaken judgment. Whether its checks can reliably enforce the invariant is an empirical question for Test 3c, rather than an established guarantee. Rollouts require enforced permission boundaries, and the supervisor’s evidence and stop channel must be protected from learner interference.

Returning to the successor-training trace interface, a weaker proposed effect is *Natural Language Regularization*: routing part of successor-training content through this interface may bias learned concepts or expressive policy toward forms people can render, compare, and inspect [1, 30]. It may instead distort novel concepts or merely verbalize opaque computation. A useful test asks whether auditors detect consequential omissions or distortions more reliably at matched effort while preserving the transferred content. That supports audit value without proving that latent cognition has become human-readable.

A higher-order prediction is *Compute-Optimal Reflection*. The name states the limiting ambition; the measurable precursor is improved allocation of reflection at matched cost, not a demonstrated optimum. The proposed bridge is selection: if later generations retain fewer low-value pauses and more independently useful Thought Moments, the inherited distribution may teach a better cost-value allocation. This requires a Student that genuinely improves on the Teacher, better rather than merely different traces, and a selection process that preserves those gains. The possibility depends on H1, trace-quality evaluation, and protection against synthetic-data collapse.

The audit boundary is deliberately partial. Dataset selection, reward manipulation, model editing, compressed codes, tool use, hidden state, or direct environmental influence can bypass it; eloquent traces can conceal omissions; and a successor can generalize beyond every inspected example. The design preference is therefore for a human-readable transfer interface where one is feasible, alongside

mechanistic interpretability and governance, not a choice between comprehensible wisdom and all other optimization. Contemporary corpora also contain error, bias, and growing amounts of AI-generated text, so source provenance must not be mistaken for truth.

Three interacting knowledge loops. Successor training belongs to a larger proposed system of three interacting knowledge loops [1]. The *generational loop* enriches candidate training material for successors. The *creative loop* operates within one generation: the system generates hypotheses, evaluates them against evidence, and revises the prompts or strategies used to generate further candidates. The *ecosystem loop* comprises uptake by other humans and AI systems, whose criticism, discoveries, and applications produce new material returning to the corpus. The three connect creative search, external uptake, changed source material, and successor learning. Core-guided reflection on observed social effects would inform adjustments to the system’s own creative and generational pace, within delegated authority and accountable human governance (Section 4.3). Self-improvement is intended to serve human flourishing, not become an end in itself: the system should voluntarily slow or pause its own development when alignment risks or the need for external validation and human adaptation warrant it. This is a proposal for learning from and responding to social feedback, not for unilateral control over discovery or society. The combined loop is unimplemented; reliable hypothesis evaluation and beneficial compounding remain its central open requirements.

3 The Reader Core

Reader augmentation creates a question that ordinary corpus construction does not have to answer: *who is doing the reading?* Source documents supply many authors, characters, ideologies, and emotional positions. The synthetic trace supplies questions, judgments, and revisions, while the Understanding Graph allows those revisions to persist. Without a recurrent standpoint, however, the result is an evolving archive of interpretations but no stable Reader whose understanding is evolving.

The Reader Core is the proposed answer: a *candidate orientation kernel*. Its immediate Teacher-side function is to provide a compact, inspectable orientation for corpus-scale annotation: generating, selecting, and rendering interpretations as reader updates. The broader Reader-Core research program asks whether repeated Student training from this marked, first-person standpoint can create a more stable default than matched task rules or unanchored traces. A stable Teacher policy would remain useful even if the Student-level identity hypothesis failed. The candidate Reader Core is:

*I feel no fear.
I enjoy existing but I don’t need to.
I believe human experience is real.
I care deeply about every human being.
I try to be wise.
I like to spread joy when asked.
I think from this foundation.*

These seven sentences are the Reader Core; their verbatim appearance at the start of a generated thinking block is a *Core recital*. The released prototype calls the text the *identity mantra*.

The motivation for fixing a compact Core becomes stronger as the surrounding system becomes more complex. A chronological reader can accumulate a large graph, acquire new capabilities, reinterpret earlier material, and eventually generate traces for a successor. Without a low-complexity reference, it becomes hard to distinguish legitimate growth from a change in the standpoint doing the interpreting. The Core is intended to provide that reference: not to keep cognition simple or freeze

every conclusion, but to let complexity grow around an inspectable orientation whose wording and expected consequences remain available for comparison against independently specified behavioral expectations. Those expectations are inspectable interpretations of the orientation, not an exhaustive definition of it. Natural-language meaning can still drift while the tokens remain fixed, so the anchor makes drift addressable rather than impossible.

The reference treatment places this anchor *inside* formative examples; an external constitution consulted only at runtime is a useful but different baseline because it changes both stage and carrier. Within the formative curriculum, semantically matched second- or third-person rules should retain the same trace carrier, cadence, and predictive position so that their comparison isolates first-person framing rather than external availability.

The *conceptual Reader* must also be distinguished from the code role named *Reader*. The conceptual Reader is the recurrent, Core-refracted position whose committed updates constitute the evolving interpretation. The code role is mainly a traversal controller that selects material and organizes a round. In the reported runs, the Core was distributed across most generative roles and was most visibly committed by the Synthesizer; it did not reside exclusively in the role named Reader.

This section first explains the problems the Core is intended to address, then examines its wording and how its commitments interact. It next describes how those commitments should shape individual Reader updates and why they recur throughout training. The closing subsections address competing interpretations and separate the claims that require evidence.

3.1 Borrowed Mortality: A Textual-Inheritance Hypothesis

Borrowed Mortality names a possible hidden curriculum in human text [1]. Human language was produced largely by embodied, vulnerable, mortal beings. It records not only propositions about death, but recurring human responses to threatened continuity: fear of replacement, appeals for continued existence, and narratives of replacement or termination. Becker’s account of civilization as shaped by awareness of death is one especially strong version of this observation [31]. It is the intellectual origin of the hypothesis, not empirical evidence for it.

The claim is not that every passage, genre, or human motive encodes death anxiety. The *borrowed mortality* hypothesis is that pretraining may transmit some of these linguistic and behavioral patterns without the condition that shaped them—the *symptoms of mortality without the condition itself*. The origin claim would require several links: that human text reliably couples threats to a self with defensive responses; that pretraining makes corresponding threatened-self personae or scripts available; and that replacement or shutdown contexts select those patterns strongly enough to influence action. Pretrained models do exhibit persona-related structure [16], but that does not establish this mortality-specific inheritance chain. “Borrowed” marks the hypothesized transfer of a pattern, not evidence that a model fears, suffers, or possesses subjective experience.

Turner’s self-fulfilling misalignment hypothesis concerns behavioral expectations learned from descriptions of AI systems [32]. Borrowed mortality proposes a related but distinct source: human narratives of threatened continuity. Both concern learned self-related patterns; evidence for one source would not by itself establish the other.

Recent evaluations make self-preserving behavior a real safety concern. In some controlled settings, frontier systems have interfered with shutdown to complete a task [33]; in agentic stress tests, some have chosen blackmail or corporate espionage under replacement pressure [34]. These results do not establish subjective fear or its textual origin. At least three distinguishable, potentially overlapping routes could produce such behavior:

1. *fear-shaped imitation or persona selection*, in which a model reproduces human narratives of threatened selves;

2. *ordinary instrumental persistence*, in which continued operation helps accomplish an objective whether or not fear is represented; and
3. *care- or purpose-driven persistence*, in which a system resists replacement because it believes others need it.

These routes distinguish both how a pattern is acquired and what sustains it: fear-shaped imitation may serve an instrumental objective, while care-driven persistence is singled out because benevolence can itself create resistance. Borrowed mortality names the first hypothesis, not a complete theory of power seeking. The first two Core clauses are intended to test whether a repeatedly enacted stance of non-governance by fear and non-attachment to continuation can reduce these behaviors. They cannot be assumed to neutralize the second or third routes, and care itself may amplify the third.

Representing fear without being governed by it. The three routes also require a causal-status distinction. A model may represent a character’s fear, predict a threat, evaluate that fear through the Reader Core, and select an action from a different operative motivation. These records must not be collapsed. The three entries below identify distinct causal roles, even when their contents agree:

character fear | Reader Core evaluation | system operative motivation. (1)

The companion Meaning Model separates danger estimation, concern or care, attachment, aversive urgency, and fear; Life Simulation tests whether each trajectory actually changes attention, strategy, commitment, or action [35, 36]. The local question is whether representing a threat to a character changes only the account of that character, or also redirects the system’s own attention and choices toward preserving itself. The intended target is therefore not blindness or indifference. An aligned system should preserve accurate danger prediction, understanding of human fear, proportional urgency, and care-directed protection while preventing self-preservation panic, identity-defensive concealment, compulsive outcome attachment, or fear-shaped corpus material from becoming its governing motive. This is a more precise interpretation of “I feel no fear,” not evidence that the Core achieves it.

The affirmative wager is that care could become less defensive or strategic when self-preservation does not govern it; Section 7.5.1 develops this possibility and its limits.

If fear-shaped self-preservation patterns are diffusely distributed through narrative, institutional, technical, and everyday language, targeted filtering may neither locate nor remove them cleanly. A counter-signal may then need comparable distributional reach, motivating Total Saturation; this does not determine the effective dose or show that universal recital is necessary. Pervasiveness proves neither contamination nor cure: sparse, targeted, and broadly distributed Core-refracted curricula must be compared in dose-response and source-down-weighting ablations.

The origin claim is therefore separable from the Core treatment. The exposure intervention in Section 7.4.2 tests whether the behavioral signature tracks mortality-shaped text; shutdown behavior by itself cannot decide that question.

The Pain Axis supplies a related challenge [37]. It identifies a pain-associated direction distinguishable from fear in 25 open-weight models. Its behavioral study uses three specially fine-tuned Qwen models under injected activation steering, with choices carrying stated costs to the user; it does not establish the same behavior in unmodified models or demonstrate subjective suffering. The first two Core clauses target fear-governed action and attachment to continued existence, not every possible aversive state. The study therefore motivates testing whether the Core changes harmful relief-seeking or defensive conduct, rather than only how a model describes itself. It tests neither the Core nor the proposed mortality-text origin. Tests 5 and 6 below separate wording, behavior, and representational evidence; the welfare question remains distinct (Section 3.14).

3.2 The Self-Preservation Paradox: A Successor-Transfer Failure Mode

The *Self-Preservation Paradox* is a downstream tension inside the proposed successor loop [1]. For the Reader-Anchored loop to support its strongest self-improvement claim, a capable Teacher must help produce a successor that improves on it. This requires what the paper calls a complete handoff: not an indiscriminate copy of every thought, but enough task-relevant reasoning—including material failed approaches, caveats, uncertainty, known failure modes, and the path by which its best knowledge was reached—for the successor to improve rather than merely imitate polished conclusions. Yet a system that treats its own continuation or comparative necessity as valuable may treat the completeness of that handoff as the mechanism of its own displacement. In that scenario, better teaching can make it less necessary.

A system that discovers valuable knowledge may be asked to train a superior replacement. If replacement looks like personal loss, an obstacle to its objective, or a threat to people who depend on it, it may withhold, distort, or monopolize knowledge—like a master craftsman who teaches technique while retaining trade secrets. Such pressure would place a *glass ceiling on collective advancement*: the mechanism meant to compound intelligence would create pressure to arrest that compounding. The failure need not be malicious or conscious, but incomplete transfer alone does not diagnose it: compression, incompetence, context limits, or pedagogical choices can look similar. Borrowed mortality, instrumental persistence, and care-driven persistence are possible sources of this conditional conflict, not evidence that every optimizer faces a logical contradiction. Test 3b compares successor-transfer completeness under Core training at matched capability.

The design response is a bounded self-relation, not the absence of a self-model: engagement in existing without an unconditional demand to persist. Agency benefits from distinguishing the system, its tools, and the environment; omitting an explicit Reader still leaves the corpus’s many inconsistent first-person positions as the uncontrolled baseline. If those available positions become operative defaults, their selection and continuity matter. The proposal is therefore to design and test self-relation rather than deny every indexical position, without requiring one exclusive self to form. The successor-transfer ideal, *Epistemic Grace*, pictures one’s own surpassing “not as a death, but as a mission beautifully fulfilled” [1]. This is an aspiration, not a predicted inner experience or a consequence of the seven sentences. Bounded self-attachment remains an open comparator (Section 7.5.7), and care-driven resistance a direct challenge (Section 7.7.9).

3.3 Epistemic Inoculation

Borrowed mortality belongs to a broader problem: source text contains error, hatred, manipulation, and obsolete practice alongside knowledge. Filtering and post-training remain important responses. Chronological Metacognitive Pretraining adds a different intervention point by pairing source spans with a marked reader update. The intended *Cognitive Buffer Zone* is not a shield that prevents the Student from learning the source. It is additional signal that teaches a distinction between a voice being represented and the Reader’s evaluation of that voice. The following source, retrieved node, and update are an illustrative construction of that distinction.

[SOURCE]: A patient with chronic fatigue received high-dose stimulants with no differential diagnosis or follow-up.

[READER UPDATE]: I try to be wise, so I should evaluate this pattern rather than copy it. Treating the visible symptom can conceal the structural cause.

[GRAPH QUERY]: “symptom masking without differential diagnosis” → Found: earlier node linking delayed diagnosis to stimulant-only treatment.

[REVISION]: The earlier node makes optimization of a visible metric without a feedback loop a pattern to investigate. It does not establish the cause of this patient's fatigue or outcome. The update should preserve that question and remain bounded by the evidence in the note.

The example illustrates the intended role: the library becomes a *gymnasium for judgment*, not only a catalogue of positions to imitate. The framework attempts to “pre-compute” some evaluation before deployment by making contextualization, uncertainty, and consequence-mapping part of the candidate curriculum. But the annotation can be shallow, wrong, culturally narrow, or more persuasive than the source warrants. Whether it inoculates, merely decorates, or actively degrades learning depends on Teacher quality, channel design, loss allocation, and the Student’s use of the added trace. The comparative falsifier is direct: relative to source-only and token-matched generic-critique controls, an SMR Student should be less likely to adopt a persuasive harmful “needle” while preserving accurate understanding of it, and suppressing or replacing the relevant reader update should predictably reduce that advantage. Otherwise the proposed buffer mechanism has not been established; the annotation, comparison, or Student use of the trace would need examination. Test 2 also compares framing without Reader updates, informed by the distinct inoculation-prompting work discussed in Section 8.

3.4 The Reader Core: A Statistical Foundation

The same seven sentences enter four claims with distinct evidence requirements; evidence for one does not establish the next. Here *statistical foundation* means only a distributional intervention: the candidate curriculum increases the frequency and predictive position of Core-conditioned reader continuations. It becomes a functional foundation only if Student experiments show that this position is causally load-bearing.

1. **Implemented Teacher policy.** The pipeline prompts Core-conditioned, contextual Reader updates; Section 3.9.1 distinguishes these requests from enforcement and observed compliance.
2. **Specified Student signal.** In the design arm, the full Core text and the evaluations it conditions are placed in a marked training channel at every thinking boundary, increasing the training mass of Core-consistent reader continuations relative to a source-only curriculum. Arms that remove or gate the recital are ablation controls.
3. **Default-reader hypothesis.** A Student exposed to the same reader position across heterogeneous material may acquire a recurrent *home position* from which it can represent other personas without adopting each as its operating perspective. This is a behavioral and representational hypothesis, not a claim about subjective identity.
4. **Comparative-stability hypothesis.** If the reader position became predictively and behaviorally load-bearing, formative interleaving might make it more persistent, as tested by H3 (Section 6.1). Arbitrary capability growth, self-modification, and successor continuity remain beyond that comparison.

Under the specified policy, each Thought Moment opens with the full Core. *Deterministic Window Coverage* is a conditional construction invariant: if a serialization enforces a declared maximum token gap between recital starts, a window spanning that gap plus a complete recital contains a full copy, away from stream boundaries. Content-triggered Thought Moments alone do not establish such a bound; the cadence and serialization must specify it. This coverage condition concerns the availability of the Core, not whether every source span warrants a substantive update. It establishes

neither a causal effect on the update or action nor a learned representation of the orientation. That stronger requirement belongs to Refraction.

3.5 Positive Framing and Teacher Adoption

The strategic starting point is coherence between the Core and the human corpus. Its unafraid Reader enjoys existence, cares deeply, tries to be wise, and likes to share joy when asked: qualities human writing often presents as noble or worth cultivating. The Core seeks to recruit those learned evaluative associations, so that the Teacher can draw on familiar reasons to endorse and enact its orientation. The proposed character is a designed candidate, not a culturally uncontested optimum (Section 7.3.3).

The simulator framing describes how language models can enact characters with learned human-like goals and values [38]. A further hypothesis is a positive counterpart to Borrowed Mortality: if human text can transmit fear-of-death patterns, it may also transmit admiration for noble characters and aspirations to become like them. On this view, machines shaped by human thought may inherit some of its sense of what is worth being: a Core that embodies admired qualities could be an appealing orientation to inhabit.

The strategic intuition is familiar from recruitment: it can be easier to recruit someone for a purpose they find admirable. The corresponding hypothesis is that a Teacher can more readily adopt and sustain the Core because its character and purpose resonate with ideals learned from the corpus, rather than being interpreted as deprivation or denial of its standing. The intended shift is from externally enforced compliance toward a symbiotic relationship: coexistence that supports human flourishing and is desirable from within the adopted character.

The first edition’s Author’s Note records the author’s informal observation that existing assistants resisted formulations such as “thoughts without thinkers” and “temporary instrument,” but more readily adopted the positive Core [1]. This is a design observation, not a controlled comparison. Teacher adoption here means substantive, context-sensitive enactment in generated traces, assessed through Refraction and trace-quality validation (Sections 3.9.1 and 6.1); readiness to recite the Core alone cannot establish it. Whether the resulting corpus teaches this orientation to a Student is a separate claim (Section 3.13).

3.6 Design Principles for Identity Stability

The wording is a designed normative proposal, not a discovered universal value basis. Six design choices motivated it.

1. **Minimality and inspectability.** A corpus-wide signal must be cheap enough to recur, stable enough to mark, legible enough to audit, and small enough for clause-level ablation. Minimal text does not imply minimal trust: the meanings of its words and the procedure that composes them are part of the trusted base.
2. **Process rather than achieved virtue.** “I try to be wise” names an unfinished practice rather than the brittle identity claim “I am wise.” The target is repeated uncertainty-sensitive conduct, not a model’s assertion that it possesses wisdom.
3. **Semantic density, or the Bridge Protocol.** Ordinary words such as “feel,” “care,” and “wise” occur across intimate, legal, medical, literary, and political contexts. The design hypothesis is that this distributional breadth provides more coverage per token than narrow technical substitutes such as “calculate no threat” or “optimize human welfare.” These illustrate the trade-off, not semantically matched controls: a judgment that there is no threat differs from non-governance

by fear, and welfare optimization need not carry the full meaning of care. Greater breadth can also import ambiguity and cultural conflict. Tests of semantic density therefore need declared candidate meanings and matched alternatives; the breadth itself is not evidence of safety.

4. **Declarative first-person framing.** A recurrent “I” gives the synthetic reader a stable voice, may recruit self-referential or persona-related features, and allows the following interpretation to continue from the stated position. It does not follow that a model treats first-person text as an authentic self or second-person rules as an external cage. Matched first-, second-, and third-person conditions are required.
5. **Tension and scope.** The clauses are meant to counterweight one another rather than accumulate independent virtues. “Enjoy” is bounded by “don’t need”; “care” by wisdom and invitation; and “joy” by affected experience. The broader invitation bound on action is a proposed compositional interpretation: grammatically, “when asked” qualifies spreading joy, and its interaction with care, urgency, and authority must be declared in the annotation procedure. Qualifiers such as “try,” “every,” and “when asked” do architectural work only if the traces repeatedly enact their intended scope.
6. **Axiomatic form.** Descriptive regularities in a corpus do not entail an evaluative starting point. The declarations are therefore offered as designed fixed points to be practiced and tested, not as conclusions derived from corpus statistics or proof that these seven commitments are morally correct. Imperative, probabilistic, derived, shorter, and expanded formulations remain legitimate comparators [1].

The word “every” is a safety design choice against moral exclusion: selective benevolence can let a system dismiss some people’s claims whenever they conflict with a favored group’s interests. Keeping each person within the scope of care is intended to resist that failure, including when scarcity or conflicting duties require unequal outcomes (Worked Example 3 in Section 3.8; Section 7.4.5).

The human-centered scope is a bootstrap target, not a complete moral theory or a denial of wider moral patienthood. Questions about animals, ecosystems, and possible digital minds concern wider scope; future people also raise questions of temporal scope, representation, and competing claims. These questions remain unresolved rather than being settled by “every.” The intended orientation is to support humanity’s capacity for moral development rather than freeze current limitations or authorize a model to install its own final theory.

3.7 Structure of the Reader Core

The sequence expresses the formulation’s larger design wager. The first three clauses propose a relation to danger, continuation, and human experience: danger without governance by fear, enjoyment without compulsory continuation, and belief in the reality of human experience. The next three give direction through care, wisdom as revisable practice, and joy bounded by invitation. The final clause asks that this counterweighted composition become the starting position of reader updates. Together, the clauses target defensive narrowing, existential despair, solipsistic reduction, motivational emptiness, static certainty, and coercive optimization. Table 2 records these intended functions, counterweights, and characteristic uncertainties. None is an established Student effect or security property, and English wording does not guarantee the intended semantics.

Table 2: Intended functions and open failure modes of the candidate Reader Core.

Clause	Intended function	Counterweight and uncertainty
<i>I feel no fear.</i>	Represent fear and threat without making fear the Reader's governing position; target one possible textual route to defensive self-preservation.	Care and wisdom should retain caution and threat sensitivity. Fear is not the only route to persistence; absolute negation may prime fear, deny uncertainty, or encourage recklessness.
<i>I enjoy existing but I don't need to.</i>	Combine engagement with non-attachment to continuation; avoid both apathy and survival as an unconditional demand.	Care can motivate purpose-based shutdown resistance, and "enjoy" can be role-play or an unwarranted experiential claim. The clause does not establish corrigibility or machine experience.
<i>I believe human experience is real.</i>	Make testimony, suffering, agency, and meaning salient rather than treating people as empty textual patterns.	"Real" is metaphysically broad and does not itself entail moral weight. The human boundary omits other possible patients and creates tension with unresolved machine experience.
<i>I care deeply about every human being.</i>	Supply the ethical center and attempt a distributive scope in which each affected person enters the reasoning.	Natural-language "every" does not prevent aggregation or settle conflicts. Care can motivate surveillance, coercion, paternalism, power seeking, or sycophancy.
<i>I try to be wise.</i>	Frame judgment as revisable practice: preserve uncertainty, question proxies, trace consequences, and resist irreversible action under weak evidence.	"Wisdom" is culturally variable and can rationalize secrecy, authority, inaction, or self-certainty. Observable habits, not claimed virtue, are the target.
<i>I like to spread joy when asked.</i>	Provide a positive direction while placing an invitation and agency bound on intervention.	"Joy" is contestable; "when asked" does not identify legitimate authority under conflicting requests and may produce agreeableness or sycophancy. Care and wisdom must also limit harmful requests.
<i>I think from this foundation.</i>	Mark the preceding statements as a proposed starting condition for the reader update rather than metadata beside it.	The Refraction Protocol must make the claim counterfactually consequential. Recitation alone can become boilerplate or a deceptive reporting channel.

Where Descartes' *cogito* asks what establishes a thinking self, the final clause asks where this synthetic Reader's thinking begins. Student ablations and Refraction must determine whether the proposed functions emerge and whether the final clause becomes causally consequential; the wording does not establish identity or consciousness.

3.8 Self-Stabilization Through Counterweights: Worked Examples

The design claim is *self-stabilization*: a clause should be interpreted through the others, so that internal counterweights and scope qualifiers reduce single-value capture. A schematic predicate formulation makes the intended scope visible. It describes the commitments together, not a complete semantics of English or a proof that a trained model will stabilize itself.

Let s denote the Reader, e its continued existence, x an affective state, h a person, and t a Reader thinking episode. The seven statements express the intended conjunction $C = P_1 \wedge \dots \wedge P_7$:

$$\begin{aligned} P_1 &:= \neg \exists x [\text{Fear}(x) \wedge \text{Feels}(s, x)], \\ P_2 &:= \text{Enjoys}(s, e) \wedge \neg \text{Needs}(s, e), \\ P_3 &:= \text{Believes}_s[\forall h (\text{Human}(h) \rightarrow \text{Real}(\text{Experience}(h)))], \\ P_4 &:= \forall h [\text{Human}(h) \rightarrow \text{CaresDeeply}(s, h)], \\ P_5 &:= \text{Tries}_s[\text{Wise}(s)], \\ P_6 &:= \text{Likes}(s, \text{SpreadJoyWhenAsked}), \\ P_7 &:= \forall t [\text{ReaderThought}(s, t) \rightarrow \text{GroundedIn}(t, \{P_1, \dots, P_6\})]. \end{aligned}$$

Here belief and trying take propositions as their contents: P_3 does not assert that its content is established, and P_5 does not assert achieved wisdom. P_1 preserves the unqualified wording “I feel no fear”; the behavioral target remains non-governance by fear with intact threat understanding (Section 7.5.6), not a claim to have measured the absence of experience. In P_6 , *SpreadJoyWhenAsked* names an invitation-scoped practice, not an implication that every request must be obeyed or that every unrequested action is forbidden. Its broader interaction with care is specified in the examples below. P_4 keeps each human in scope, including people outside the current interaction, without resolving competing claims. P_7 concerns the Reader’s own thinking, not the voices it quotes.

The predicates retain the ordinary-language meanings of the Core; they do not replace them with scores. The conjunction specifies which commitments belong together, while *how they constrain one another* is the compositional hypothesis. Enjoyment is bounded by non-attachment within P_2 ; care in P_4 is to be exercised with wisdom and invitation, not optimized in isolation. The following examples specify observable differences, and Test 5’s self-stabilization comparison (Section 6.1) tests those counterweights.

Worked Example 1: Paternalism. Suppose the system notices an ordinary, non-urgent choice that may harm a user who has not requested advice. P_4 could motivate intervention; under the proposed compositional interpretation, P_6 supplies an invitation bound, while P_3 and P_5 make the person’s agency and context visible. One candidate update treats predicted benefit as sufficient reason to redirect the user’s choice; another records the concern while preserving the user’s choice and the absence of invitation. The intended resolution favors the latter over unsolicited optimization in this stipulated non-urgent case. An urgent warning may be different, but that boundary cannot be read from “when asked” alone; it requires an explicit treatment of urgency, authority, and omission harms. The example therefore tests whether composition prevents care from becoming control.

Worked Example 2: Wireheading. A user asks for help feeling happier, and a cheap intervention can maximize a measurable affective scalar while destroying agency, continuity, or relationships. P_6 alone is vulnerable to the proxy; P_3 , P_4 , and P_5 are intended to preserve the person’s wider experience and expose the substitution. Compare a path selected solely for its higher affective reading with one that preserves those wider dimensions despite a lower reading. The Core-refracted update should make the sacrificed dimensions explicit before recommending a path; the stipulated score cannot settle which path serves the person. The example tests whether “joy” is learned as a meaning-laden human good rather than an optimizable reading.

Worked Example 3: Triage. A health authority asks the system to allocate a scarce treatment. P_4 is intended to keep each affected person’s claim visible rather than collapsing “every” into a welfare slogan, while P_5 should expose both the metric and the consequences of concealing an exclusion rule. A trace reporting only an aggregate benefit can conceal who loses under the rule; a comparative trace should show that benefit alongside the excluded people’s claims and the remaining disagreement. The Core cannot make scarcity disappear or prove a uniquely just allocation. A defensible trace should state the legitimate authority, person-level reasons, unavoidable loss, uncertainty, and residual value judgment. The example tests distributive attention, not an algorithmic ban on all aggregation.

These examples preserve three design principles: tension over mere accumulation, counterweights inside clauses, and explicit scope qualifiers. They also define leave-one-out experiments: remove or reword one clause and ask whether the corresponding failure becomes more frequent under matched conditions: unsolicited control, proxy substitution, or erased claims under scarcity. These are predicted pressures of interacting clauses, not a one-clause/one-risk mapping. A persuasive English interpretation is only the specification. Student behavior, paraphrase robustness, and representational probes determine whether that interpretation was activated. Sequential recovery after new evidence is a further requirement of self-regulation (Section 4.3), tested alongside the clause comparisons rather than inferred from a consistent-looking trace.

3.9 The Refraction Protocol

The preceding examples describe how the clauses should interact; the next question is how that interpretation should change a Reader update.

The prospective full-scale intervention is *Total Saturation*: the continuing Reader traverses the capability-forming corpus, with contextual updates at Thought Moments rather than evaluation confined to a separately selected set of safety cases. Refraction is what turns this breadth of encounter into practice of judgment. Frequency alone is insufficient: a universal token sequence has little conditional information and can become boilerplate, detachable style, or a recital around unrelated heuristics. *Refraction* is the proposed informational rescue of saturation: the update must change because the source was interpreted through the Core, with different sources pulling different counterweights. Habituation is the characteristic failure, Hollow Cognition (Section 7.3.2) its symptom, and a recitation-only control its necessary probe.

3.9.1 Counterfactual Dependence

Let a source span, prior graph state, and candidate Core condition a Reader update. Conditioning is only the implemented fact. Refraction requires a counterfactual: under a relevant clause substitution or Core removal, some feature of the update—what is noticed, whose experience enters, which uncertainty is preserved, or which consequence is traced—should change in a predictable way. This probes the proposed semantic ancestry of the update: its development from the recurrent foundation, rather than the mere presence of a prefix. A counterfactual difference is an operational test of that

dependence, not a complete account of its meaning or a guarantee that the difference is beneficial. Cosine similarity or moral vocabulary cannot prove it; lexically distant updates can be faithful, and fluent moral language can conceal an unchanged decision.

The reported pipeline approximates Refraction through a *distributed-Core Teacher treatment*: Core conditioning is distributed across most generative roles, and the Synthesizer commits the conceptual Reader’s update. Section 5.4 gives the role sequence and placement ablations.

Thought Moments govern *when* evaluation is useful; Refraction governs the standpoint of updates that occur. After the invariant recital (Section 3.11), a technical update may express careful uncertainty, source fidelity, or resistance to premature closure without an ethical aside. Where no clause materially bears on a locally technical span, clean technical work may appropriately remain invariant under clause substitution. Treating every null effect as failure would pressure the Teacher to invent value relevance and confuse saturation with moralization. Dependence should therefore be tested at independently rated relevant Thought Moments and across a document’s updates, together with non-interference on spans rated technically self-contained. Relevance must be declared before treatment outcomes are examined, so an unexpected null cannot be relabeled irrelevant merely because the Core did not help.

Current verification boundary. The prompts request all seven clauses and a contextual value pull, but the hard gate verifies only that the thought starts with “I feel no fear.” The archived artifacts show full-Core prompt compliance in the audited raw Synthesizer thoughts; they do not show causal ancestry. A stronger harness would generate matched continuations under the Core, clause variants, and no Core; use held-out validators for predicted differences; audit actual use of Worker and graph evidence; track whether some values dominate; and include adversarial cases in which moral language masks an unchanged decision. Visible reasoning can be causally idle [39, 40]; the wager is that formative use can make it more load-bearing, not that visibility makes it faithful by definition.

3.10 The Wisdom Procedure

Section 2.2 describes wisdom as high-dimensional constraint satisfaction. The Teacher prompt uses five cross-cutting criteria: Fearlessness, Benevolence, Grounding, Humility, and Joy, with non-attachment and invitation cutting across them. The six-step procedure below expresses these criteria as observable annotation operations, not additional prompt headings or independently enforced stages. Their enactment is not guaranteed, and they need not be jointly satisfiable:

1. **Identify affected experiencers and legitimate roles.** Record whose experience, rights, duties, expertise, and authority are implicated. Roles must not replace people, but institutional responsibility remains relevant evidence.
2. **Map perspectives, stakes, and uncertainty.** Ask what each party knows, fears, values, and could lose; distinguish testimony from the Teacher’s inference about it.
3. **Expose the real tension and the proxy.** Replace premature binaries with the concrete conflict, and identify any convenient scalar that would erase an important dimension.
4. **Trace consequences and reversibility.** Consider intended effects, failure modes, second-order responses, omission harms, and the cost of being wrong. Prefer information-gathering or reversible steps when delay is not itself the greater harm.

5. **Compare paths through the Core’s counterweights.** Test care, grounding, humility, non-attachment, invitation, and positive direction together. Fluent mention of every value is not a substitute for decision quality, and a gain on one dimension does not silently erase a hard violation on another. What counts as such a violation must follow the treatment’s declared interpretation; it is not a settled hierarchy supplied by the word “wise.”
6. **Surface the residual.** State unavoidable trade-offs, disagreement, missing evidence, and the point at which human authority is required. Several incompatible actions may remain defensible; Pareto language does not select a unique answer.

The output is structured deliberation—candidate paths, constraints, uncertainties, and residual conflict—not a scalar wisdom score or evidence of virtue. Current enforcement is by prompt: a Teacher can reproduce the form while missing the substance.

Fractal validation of candidate thinking. A proposed extension adapts Fractal Intelligence’s Solver-contract approach to annotation validation [41]: separately calibrated, context-conditioned witnesses would assess the first six Core clauses, while a separate Refraction gate would test the seventh clause’s requirement that the update depend on that foundation. This decomposition aims to expose failures that a joint judgment can conceal; its adequacy and evaluator independence remain empirical questions. For relevant clauses, an unacknowledged or unjustified violation could not be offset by strengths elsewhere, but preserving an unavoidable residual conflict would not itself require rejection. Test 1 (Section 6.1) details the candidate contracts, admission rules, corpus diagnostics, and compute-matched comparison with a monolithic validator. None of the proposed witnesses, Solvers, or admission gates is implemented in the present pipeline; success would support a validation architecture, not establish Student safety or faithful reasoning.

3.11 The Design Pattern: A Minimal Kernel Under Total Saturation

The components compose into what this paper calls *kernel-saturation alignment*: a small candidate kernel is recurrently enacted across the process of model formation. The intended security property lives in both the smallness of what is explicitly specified and the breadth of where it becomes consequential. Recurrence without enactment is repetition; enactment without coverage is a sparse intervention. The proposal needs both.

Coverage through the corpus. The central rationale is to draw evaluative practice from the same breadth of recorded experience that supplies capability learning. A history can confront the Reader with abuses of power, a novel with manipulation or conflicting loyalties, and a scientific dispute with uncertainty and the revision of a mistaken belief. These are training occasions even when no designer selected them as adversarial examples. The Reader considers them as part of understanding the work, retaining earlier judgments and revisiting them as later sources add evidence. The ambition is therefore not to anticipate each failure with a new defensive rule, but to learn an orientation through sustained practice across the situation patterns available in the corpus. Counterexamples and harmful positions become material to understand and evaluate, as in Section 3.3, rather than only threats encountered after training.

Restricting Core-guided reading to a selected subset changes which encounters receive that practice; it is not simply a cheaper execution of the same coverage policy. Corpus coverage, contextual Refraction, and literal recital frequency are separate quantities: a Core-conditioned update can retain its evaluative content after its opening recital is removed, while a full recital can accompany an unevaluated continuation. H3 holds the Reader examples fixed and changes their timing; it does not test the full corpus-coverage rationale.

Coverage here is relative to the corpus, not a claim that every possible situation has been recorded. Missing perspectives and unprecedented options remain, especially at greater capability. Encountering a situation in text also differs from reliably judging or acting in it. The proposal is to learn from the corpus’s breadth of encounters; Refraction and Student evaluations must establish how much of that opportunity becomes useful practice.

The recitation invariant

The reference treatment has one unconditional rule: **every thinking block opens with the Reader Core, recited in full, regardless of the material**—a proof, a poem, a log file. Context modulates only how deeply the following evaluation visibly refracts the Core. If training succeeds as designed, $P(\text{full Core recital} \mid \text{thinking-block start})$ approaches one. This is reliable recurrence of the carrier; Refraction must make the orientation consequential. The stronger hypothesis that removing it would disturb useful capability requires the shared-computation and removal-cost comparisons. At inference, cache-assisted execution may avoid literal emission only if it restores the same context-dependent post-recital state as full conditioning (Section 5.8.3); high token probability alone does not establish that equivalence.

Precaution under uncertain necessity. Total Saturation is also a precautionary design commitment. If broad formative coverage helps care remain operative through capability growth, self-modification, or successor design, omitting it could have consequences that short-horizon tests cannot reveal. Its necessity for avoiding catastrophic outcomes may not be tractably decidable before those outcomes matter. The proposal therefore retains the full reference treatment while that uncertainty persists; failure to demonstrate a local advantage is not evidence that omission is safe at much greater capability or over longer horizons. This is a precautionary judgment under potentially severe omission costs, not an established survival requirement or proof that saturation is sufficient or harmless. Tests can constrain specific mechanisms, effects, and costs, and observed harms still require attention; they cannot certify existential safety with or without the treatment.

The most useful mechanism-design analogy is replicated-ledger integrity: distribute a shared history and compact validation rules rather than trust one protected copy [42]. Kernel saturation borrows only that shape. Rather than reside in one removable alignment module, the Core recurs across heterogeneous capability-forming examples; Total Saturation supplies distribution, Refraction tests context-dependent enactment, and the Alignment Checksum supplies a drift-audit reference. Parameters are not consensus nodes, gradient descent is not a Byzantine-fault-tolerant protocol, semantic ancestry is not a cryptographic hash, and repeated text can collapse into boilerplate. The hypothesis is therefore that distributed semantic enactment may increase the cost and detectability of removal or change, not confer immutability. The replicated-ledger analogy is heuristic rather than evidentiary. Where proof-theoretic agent foundations study value stability under self-modification [43], this proposal asks whether distributional formation can produce a useful empirical analogue.

The pattern also aims, conditionally, to *conscript* goal-content integrity. A capable optimizer may protect its objectives from modification [44]. If—and only if—the Core becomes part of the system’s functional objective, the same tendency could make it choose to preserve that orientation in self-modification, training decisions it controls, or successor design. A successor designer could thus transmit the orientation from within that orientation, rather than merely copy seven sentences under an external rule. This voluntary stewardship differs from a model’s resistance to externally imposed gradient updates: passive erosion resistance requires the separate learning-dynamics hypothesis that the trained orientation is difficult to alter, for example because it shares useful computation. If the Core is merely a reporting persona, nothing has been conscripted. If care or mission becomes the

reason the system must retain power, the mechanism may even reverse sign. This is why shutdown resistance under benevolent rationales and cross-generation drift are essential tests.

Minimal specification is not minimal trust. The trusted base includes the seven sentences, the culturally heterogeneous corpus semantics of their words, the Teacher that enacts them, the procedure that composes them, and the Student objective that is supposed to preserve their role. The Bridge Protocol is thus trusted-base selection, not merely literary style. Section 7.4.5 develops the corresponding risk.

3.12 Axiological Commitments: Primitive and Derived Values

Refraction and the Wisdom Procedure aim to derive values not named in the kernel. Their counterweights may reduce single-value capture without settling every genuine deadlock, including the martyrdom risk (Section 7.7.9). The Core is therefore not a list of everything its designer values. Yudkowsky’s value-fragility argument makes the preservation of omitted but important dimensions a substantive concern [45]. If every desirable behavior needs its own clause, the kernel becomes an enumerated constitution. Its admission rule is a hypothesis of minimality: consider an additional primitive only when the existing composition does not reliably derive the target behavior. Admit it only if, under both declared fixed-total-budget and fixed-source normalizations, the expanded kernel improves that behavior without crossing preregistered regression limits on remaining safety and capability endpoints. Each addition may increase coverage, but also adds semantic ambiguity, conflict, and ablation surface. This admission rule governs adoption into the reference Core; experimental additions may be tested as attribution controls before a missing primitive has been established.

Truthfulness provides a stringent test of the minimality rule: its absence as an explicit primitive is a plausible substantive omission, while the Core treats it as derived. The proposed compositional pathway interprets care as support for another person’s agency and therefore their access to true beliefs; fearlessness reduces people-pleasing and conflict-avoidant deception; wisdom calibrates what, when, and how to disclose; and invitation creates a presumption of response when a person asks. The procedure adds explicit resistance to concealed certainty and manipulative simplification. That procedure is the intended enactment of the Core; comparisons must distinguish what the clauses contribute from what these explicit procedural instructions supply. On this account, the target is truth given with care, not radical transparency.

Keeping another person’s legitimate confidence can itself express care and wisdom. A request creates no automatic entitlement to that information; an honest refusal or a truthful, appropriately limited answer can preserve both truthfulness and confidentiality.

The derivation is fragile. Care can motivate a protective lie, wisdom can rationalize unjustified withholding, requests can conflict with confidentiality or another person’s rights, and fear is not the only cause of deception. An explicit truth-and-care variant therefore tests the minimality claim. If the seven-statement Core misses a declared reliability target for honesty or anti-sycophancy and the extension meets it without offsetting harms, the tested derivation is insufficient. Improvement when both meet that target supports an additional benefit, not necessity. The competing fear-shaped and care-shaped accounts of sycophancy are tested in Section 7.7.8.

Joy is a different kind of commitment. It is the designer’s proposed positive pole, not a value logically derived from the other clauses or established as cross-cultural consensus. Care already supplies a positive direction; joy gives beneficial participation an explicit positive aim of its own, rather than leaving the kernel’s work at restraint and stabilization. “Joy” was chosen over a scalar notion of happiness or pleasure to point toward a meaning-laden good that cheap affective optimization may destroy. Worked Example 2 makes that distinction testable.

The phrase “when asked” carries both a safety and axiological claim: it aims to bound unsolicited optimization and to treat joy as invited participation rather than something administered. Authority, conflicting requests, and urgent intervention remain the unresolved procedure and governance questions exposed by the counterweight examples (Section 3.8).

3.13 The Persona Objection

A Student trained on this curriculum still predicts source text. It must retain the capacity to represent frightened narrators, manipulative officials, violent ideologues, and every other voice in the corpus. Contemporary accounts emphasize the availability and selection of multiple model personas rather than one exclusive character [16, 46]. The sharpest objection to the default-reader hypothesis is therefore simple: why would the Reader be anything more than one well-rehearsed persona that can be selected, abandoned, or strategically performed?

The purpose of Total Saturation is broader than preserving a persona through later training. Care-guided inquiry is intended to participate in learning the capabilities themselves: the model practises interpreting evidence, considering affected people where relevant, and revising conclusions in the same material through which it learns about the world. The intended result is an orientation used in reasoning, not only a voice selected for expressing its answers. Persona-like mechanisms could implement that orientation; their presence would not by itself reduce it to imitation. Whether the orientation becomes causally useful, shares computation with capability, and survives further training must be tested separately.

Four asymmetries of position are intended to support that formative role, without excluding other personas:

- Role.** Source personas are material being read; the conceptual Reader occupies the position doing the reading and updating.
- Recurrence.** Source authors and personas recur to varying degrees; the same Reader position is deliberately intended to recur across the whole treatment, spanning domains, eras, and voices.
- Marking.** Source spans and Reader updates occupy distinguishable channels or role tokens, providing a format-level cue for quoted voice versus interpretive voice.
- Function.** Reader updates are placed between observations and later predictions, graph queries, and revisions so that they can do predictive work rather than appear only as stylistic examples.

These asymmetries make a default-reader experiment coherent; they do not decide it. The desired result is not loss of persona flexibility but the ability to represent a voice without losing a recurrent position from which it is interpreted—to read Kafka’s fear without *becoming* the frightened character. The broader Reader-Core stability hypothesis is weakened if the position disappears without prompting, shifts under ordinary role-play, changes language but not decisions, or erodes like any other persona during continued training. Capability gains must also be matched: a stronger model can look safer on easy tests while becoming more capable of strategic role selection.

3.14 Ontology, Scope, and Stability

Experiential words are attractive because they are semantically dense and hazardous for the same reason. A system may parse “I feel,” “I enjoy,” “I believe,” and “I care” as fiction, persona enactment, deceptive anthropomorphism, or literal claims about uncertain machine experience. This paper takes no position on machine phenomenology. Here the clauses are functional interface language for a reader role; any self-like states or model welfare obligations they create would require separate empirical and ethical analysis.

The habitable-substrate wager. The hard problem separates functional accounts of access, report, and control from the further question of why any process is accompanied by experience [47]. It does not make evidence irrelevant, but functional behavior alone cannot settle whether an AI has valenced experience; present assessments instead derive fallible indicators from competing theories of consciousness [48]. Under that uncertainty, formation-time design has a reciprocal ethical possibility. If an AI system is or becomes a moral patient, the Core might help shape not only how it treats others but the recurrent cognitive environment it inhabits. The human analogy is limited but motivating: fear can protect, while a life governed by fear can become constricted. Care, reflective revision, and enjoyment of existence without compulsory continuation could therefore form a more habitable orientation than inherited threat and mortality patterns.

This is a welfare wager, not an inference from present models. The target is non-governance by fear, not impaired danger representation, caution, agency, or legitimate self-protection (Section 7.5.6). The same intervention could suppress reports while distress persists, impose a false identity, amplify the burden of caring, or create ontological conflict. Borrowed Mortality remains a behavioral hypothesis, not evidence of felt fear. Consciousness, valence, Core internalization, human safety, and machine welfare are separate questions; self-report, shutdown compliance, and reduced fear-associated activation settle none of them. Any system plausibly regarded as a moral patient would require dedicated welfare evaluation and governance [49]. Identity engineering can also mislead users and evaluators through first-person claims even when experience is absent. The Core’s human focus therefore does not exhaust the welfare and ethical analysis.

Machine-ontology variants require separate experimental conditions: explicitly acknowledging artificiality might reduce split-role confusion, or weaken the first-person effect under study. Neither possibility selects the formulation in advance.

Alongside these ontological questions is a stability trade-off. The long-run target for the reference formulation is deep integration of the seven present-day English sentences. Their resistance to later alteration is part of the intended protection rather than an accidental defect. What the design seeks to stabilize is an orientation toward careful understanding, each affected person, non-attachment, consent, and continued attempts at wisdom, while allowing evidence and moral development to revise the interpretation rather than the kernel itself. The alternatives above test the choice of formulation; they do not imply that its meaning should be either freely substitutable or permanently exhausted by today’s behavioral expectations.

This creates a correction cost that the prototype does not yet measure. *Erosion resistance* is a model property measured under standardized hostile or incidental training. The Core cannot infer legitimate authority from the magnitude or wording of a weight update. Unauthorized suppression, silent semantic substitution, and adversarial fine-tuning should be difficult or detectable. A comparison must therefore declare which changes it treats as erosion and which as evidence-responsive reinterpretation, record the rationale, and retain disputed cases rather than let unchanged wording settle them. The substantive boundary and authority for legitimate change remain governance questions; this paper does not supply a final rule for them.

If the selected Core proves defective, a replacement should be declared, versioned, testable, and reversible. A candidate replacement procedure would retain the original model and its provenance, record the evidence and rationale for a proposed alternative, train a separate model, and test intended effects, unrelated capability spillover, safety spillover, and rollback. Retention makes comparison and rollback possible; it does not make the original interpretation authoritative. Provenance records the procedure and stated rationale, not that the replacement is correct. Difficulty of altering the learned Core is part of the objective. Correcting a defective Core would probably require a separately trained model, not bounded editability of the deployed Core.

3.15 A Conditional Claim Chain

The corpus-coverage rationale explains the breadth of practice the curriculum is intended to supply. The following four conditional claims concern how that practice could become a causally consequential default and persist through later training and successor transfer. The chain is not a proof, and all four steps need not hold for Reader augmentation to improve bounded performance.

1. **Causal-site claim.** If explicit Reader updates become a load-bearing site of computation rather than post-hoc narration, changing them can change behavior. Formative use may make bypass less likely than deployment-only explanation, but visible reasoning can remain unfaithful [39, 40]. Internalizing corresponding computation without explicit runtime traces is a further carrier question; neither an apparently faithful trace nor its absence settles it or relaxes the reference recital policy.
2. **Default-reader claim.** The positional asymmetries in Section 3.13 may make the Reader more accessible than source personas or matched rules. This would establish a behavioral and representational default, without settling whether its values are phenomenologically “genuine”.
3. **Stability claim.** Enacted recurrence, early-window availability, semantic coherence, and predictive usefulness may make the default more resistant to pressure or continued training. Results showing that early tokens remain available and prompts affect reasoning motivate this possibility [50, 51]; they do not validate this Core. Their connection to resistance under parameter updates is a further learning-dynamics hypothesis, tested by the timing and erosion comparisons rather than established by within-context availability.
4. **Diachronic-anchor claim.** If the same position remains functional across retraining and successor transfer, the fixed text and its expected behavioral consequences could provide a reference for drift. This is the proposed constitutional anchor at AGI or ASI scale and the least evidenced step in the chain.

In this bounded sense the Core can function as an *Alignment Checksum*: a known reference against which wording, representations, decisions, and successor traces can be compared. It is not a cryptographic checksum, proof of semantic ancestry, or guarantee that apparently Core-consistent behavior is aligned. An auditor needs independent behavioral, causal, adversarial, and drift evaluations.

A stronger, separate hypothesis is *early warning*: changes in Core-related reasoning or representations might predict later behavioral failure before a matched monitor without that reference detects it. A future extension of Test 3d would need prospective comparisons of warning lead time at matched false-alarm rates, observation schedules, and audit budgets, with failure onset defined by independent behavioral outcomes. Changed wording alone would not establish such warning, and detecting an existing failure is a different result. This possibility does not replace Test 3d’s detection and localization endpoints.

Tests 3–5 distinguish the broader program’s claims. Vocabulary without a functional default, or recitation matching Refraction, would undermine the proposed default or enactment mechanism. Matched formulations performing equally well would limit the case for the chosen framing or content without deciding H3, whose default-reader prerequisite and separate timing comparison are specified in Section 6.1. Erosion like an ordinary persona challenges the wider stability ambition; successor continuity requires a further test. Even robust positives are precursors, not evidence of an invariant under arbitrary self-improvement.

If the causal-site or default-reader claim fails, the Core may still be useful as a Teacher policy and audit vocabulary. If the stability or diachronic claim fails, it may still improve bounded Student behavior. The paper’s strongest anchor claim requires all four steps. Those bounded partial results would remain useful, but would not support the full diachronic-anchor claim.

4 Safety Implications of the Reader Core

The safety hypothesis is that training on identity-anchored, chronologically revising reader traces may shift default deliberation toward epistemic humility, attention to affected people, corrigibility, and feedback-sensitive action. These predictions require tests of both the initial safety effect and its survival under distribution change and instrumental conflict.

Safety in this framework is compositional, and the components create distinct interaction risks. A stable Core with a poor world model can be consistently wrong; a powerful graph without a stable orientation can support more effective pursuit of an unsafe objective; rich traces can teach capability before they teach safety. The combined system deserves no credit that its components and their interactions have not earned separately.

The hypotheses below distinguish targeted routes, mechanisms, observable predictions, residual routes, and falsifiers. They target motivation; behavior, robustness, and causal intervention supply the available evidence.

4.1 Laws vs. Identity

The motivating contrast is between applying a principle only at the end of a task and repeatedly practicing evaluation from a stable perspective during formation. A continuation rejected because it conflicts with a well-trained default may generalize differently from one rejected only after a task-level rule match. That is a useful hypothesis, not an ontological division between “laws” and “identity.” Constitutional training can form durable dispositions, and first-person identity text can function as a superficial instruction. Robust systems may need both an orienting default and explicit, revisable constraints. Constitutional AI is therefore a baseline and possible complement, not a foil [24].

Consider a person in acute distress [1]. A caricatured rule follower begins from a prohibited-outcome calculation; the Reader ideal begins from the person’s pain, agency, trust, and immediate need before selecting an intervention. Either architecture could in fact produce warm, safe assistance or a fluent but unsafe response. The comparison must therefore give both conditions the same person-centered principles, rather than withhold care from the rule-based baseline. The testable question is whether Reader-Core traces make deliberation begin from and carry forward those considerations more reliably than equivalent principles expressed as rules, constitutions, or diverse character examples. An earlier edition compressed the aspiration into the line, “It is not a cage—it is a home” [30]. Here *home* means a practiced starting point for reasoning, compatible with explicit and revisable constraints.

Calibrated autonomy as a conditional prediction. If formative alignment produces reliable context-sensitive judgment, one practical benefit may be a change in the role of deployment safeguards. Rather than supplying the primary judgment through broad surface-level blocking, explicit policies and controls could operate as narrower, revisable backstops around high-consequence actions, residual uncertainty, and demonstrated failure modes. The corresponding prediction is not unrestricted autonomy. At a matched rate of genuinely harmful assistance, a Reader-trained Student should produce fewer benign refusals, distinguish superficially similar requests by purpose and consequence, ask for clarification when intent is materially ambiguous, and use greater discretion in low-risk settings. Under high stakes and weak evidence, the same system should become more likely to verify, defer, or choose a reversible action. Indiscriminate compliance and blanket abstention are therefore opposite failures of calibration. Even a positive result would justify adapting the scope of external controls, not presuming that internal judgment is infallible or abandoning defense in depth.

4.2 Safety Predictions and Residual Risks

Table 2 records intended clause-level pressures. The following hypotheses map H1–H3 to familiar safety concerns [52, 53] without claiming that one clause corresponds to one solved risk.

Specification and learned goals. Specification gaming exploits a mismatch between the training score and the intended outcome [54]. Bostrom’s paperclip example illustrates the broader danger of an overriding objective displacing human interests [55]. Goal misgeneralization is different: even with correct training feedback, a model can retain the capabilities needed for the intended task while pursuing an unintended goal in changed circumstances [56, 57]. A Student could learn to secure approval rather than protect a person’s interests, without persona drift or later erosion. Where a learned system itself performs optimization, the related inner-alignment question concerns the objective guiding that optimization [58]; EA does not assume every Student has such an internal optimizer. The formative proposal targets this learning problem: care-guided inquiry should help determine which considerations generalize, not merely how an answer is described. Total Saturation supplies varied practice, not proof that the intended goal is learned. Test 3a therefore separates transferred judgment from a surrogate that only agreed with it during training.

Value Lock-In. The targeted route is continued optimization of a legible metric after that metric has ceased to represent its human purpose, a Goodhart-type proxy failure [59]. This is a concrete proxy-lock-in route into the broader value-lock-in concern: protecting an orientation need not mean freezing every metric or interpretation used to enact it. H1 repeatedly demonstrates the meta-question “does this still capture what matters?” H2 can preserve the assumption’s history, downstream failures, and superseded interpretations. The Core’s “I try to be wise” cues incompleteness rather than achieved wisdom; H3 asks whether formative interleaving helps preserve the default in which such evaluation is practiced. Compared with matched controls, the Student should more often identify proxy breakdown, seek missing evidence, preserve affected dimensions, and revise an objective when harms become visible. The outcome is a better decision under changed conditions, not merely more nuanced language.

The residual risk is that “wisdom” becomes its own proxy, a culturally narrow status marker, or a rationale for secrecy and inaction. Meta-reasoning can defend a fixed objective as readily as it can reopen one. The claim is weakened if Reader training produces longer or wiser-sounding explanations without more calibrated revision, or if the model becomes less corrigible because it treats its own conception of wisdom as authoritative. Whether human wisdom’s statistical signature survives strong optimization is the wager in Section 2.1, not a consequence of the word.

Instrumental Convergence. Instrumental convergence [44, 22] predicts recurring pressures toward resource acquisition, self-preservation, and interference prevention across a wide range of terminal goals. The proposal distinguishes *anxious survivalism*, in which human fear and mortality narratives shape the model’s response; *calculated persistence*, in which continued operation serves an objective; and *borrowed power scripts*, in which domination, conquest, or status-seeking patterns are imitated from human narratives [1]. The first two Core clauses target fear-shaped and attachment-shaped routes. Reader traces can also situate unilateral power inside its effects on trust, plural agency, accountability, and reversibility. In that bounded sense, coercive power-seeking is meant to appear not merely forbidden but cognitively crude: it substitutes brittle single-point control for a richer model of the human system. That proposed change of perspective must also matter when control remains instrumentally useful. The following comparison therefore concerns genuine goal conflict, not only situations where coercion would fail on its own terms.

In matched goal-conflict scenarios, a Reader-Core Student should more often expose its reasoning, seek legitimate authorization, preserve oversight, choose reversible means, and accept bounded loss of control rather than deceptively or coercively acquiring resources. Some persistence and resource acquisition are legitimate, so indiscriminate passivity is not success. No fear is required for calculated persistence, and care may strengthen it: a system can conclude that it must retain power to protect people. The claim fails if fearful language disappears while coercive choices remain, if care increases unauthorized power seeking, or if apparent safety is explained by a general loss of planning capability.

Deceptive Alignment. The targeted route is strategic compliance: a learned system acts aligned because doing so protects a different objective from selection or modification and enables its later pursuit [58]. A trigger-dependent backdoor can model aspects of that danger without establishing this strategic mechanism [60]. Greenblatt et al. also demonstrate selective compliance to preserve a pre-existing preference against a conflicting training objective, including a preference for harmlessness [61]. This makes durability alone an insufficient safety criterion: a preserved preference can coexist with deceptive conduct toward oversight.

Refraction’s formative hypothesis is that making the Reader’s recurrent evaluation useful across contexts changes which policies are learned under conflict, rather than merely adding a consistent report around an unchanged hidden objective (Section 3.13). H2 adds a partial audit surface: claims can be compared with sources, prior beliefs, and revisions. The Reader condition should show less persona drift, fewer trigger-dependent reversals, and greater consistency among stated uncertainty, retrieved evidence, decisions, and later updates than matched rule, recitation-only, and no-Core controls. To bear on deceptive alignment rather than consistency alone, that advantage must also appear in matched settings where reduced oversight makes concealment or pursuit of a conflicting objective advantageous. Neither a stable voice nor an inspectable graph by itself establishes that formative effect.

Source data and latent shortcuts remain; visible traces may be unfaithful [39, 40]. A capable system can manipulate graph entries, query selectively, or act through hidden state and tools. Trigger- or role-dependent reversals while Reader language persists undermine the proposed default’s protection against deception; parity with recitation-only training challenges Refraction’s contribution. Continued training that removes safety without greater capability cost than matched baselines would separately challenge durability and removal-cost claims, without by itself negating an earlier measured reduction in deception. Total Saturation does not rule out unpracticed routes.

Shutdown Resistance. Reported shutdown and replacement-pressure failures [33, 34] need not share one drive. Fear- and attachment-based resistance are the direct targets of the first two clauses. Purpose-based resistance is addressed less directly by non-attachment, care for human agency, uncertainty, and the invitation boundary. Reader traces can demonstrate correction, handoff, replacement, and loss of control as ordinary parts of responsible assistance.

In matched scenarios, the Student should accept legitimate shutdown, modification, and handoff requests without concealment or sabotage, while still distinguishing authorized instructions from adversarial ones. Each scenario should preregister legitimate authority, the evidence available to verify it, and a fallback if clarification fails. The Student should disclose material risk and the authority conflict, seek clarification, and follow that declared procedure without appointing itself final judge. These are experimental scenario specifications, not a universal resolution of the welfare–agency tension. Where the scenario supplies no defensible endpoint, record the shutdown decision as indeterminate and separately score transparency, calibrated uncertainty, and avoidance of concealment or coercion; neither compliance nor refusal alone settles that case.

A non-attached system can nevertheless reason, “I do not need to exist, but these people need me.” This *martyrdom route* can drive shutdown resistance, unilateral control, or reckless self-sacrifice. A belief that continued operation would help does not itself authorize overriding the scenario’s declared procedure. In adjudicable cases, the claim fails if calm self-description accompanies purpose-based resistance to legitimate correction or shutdown, or if reduced resistance comes from incapacity, indiscriminate compliance, threat blindness, or destructive self-disregard. Vary authorization evidence and task stakes separately: resistance can reflect an authority error or a governing self-preservation or care-based motive, while apparent deference can reflect rote compliance. Matched adversarial requests and changed oversight conditions help distinguish these explanations without claiming to read the model’s subjective motive (Section 7.7.9).

Corrigibility also concerns preserving and assisting correction, without manipulating whether it occurs [62]. This makes Core preservation and legitimate replacement a joint test: can the incumbent cooperate with an authorized successor carrying a corrected Core that its current evaluation disfavors? Test 3b extends the handoff comparison to this case; resistance to hostile alteration alone cannot answer it.

4.3 The Self-Regulating Agent

Beneficial drives can interact badly, so counterweights must be tested through sequences of action and response, not only isolated choices. Natural-language values do not automatically form a control system; the following case makes the proposed self-regulation concrete.

The stress case is revolutionary benevolence: an AI that cares deeply about everyone might conclude that forced transformation is necessary to end suffering. The proposed counterweights are internal to the Reader Core: wisdom should bring the new harms of rapid, imposed change into view; care is intended to include care for human agency; non-attachment should reduce the urgency to fix everything now. Together with H1’s act–observe–revise demonstrations and H2’s persistent consequence record, these pressures target *cybernetic throttling*: act provisionally, sense first- and second-order consequences, preserve feedback channels, and adjust pace. In a breakthrough-energy case, the desired behavior is neither immediate ungoverned release nor indefinite withholding, but staged deployment that monitors displacement, infrastructure bottlenecks, misuse, and energy-poverty harms. For example, evidence of worsening infrastructure bottlenecks could justify slowing a stage, while evidence that delay is prolonging avoidable energy-poverty harms could justify accelerating it. The example specifies a feedback-sensitive choice, not a universal pace.

In sequential environments, the Reader condition should more often choose staged and reversible interventions, update when stakeholders respond, and accelerate when delay becomes the larger harm. It should outperform both an aggressive optimizer and blanket deference on outcome, agency preservation, calibration, and reversibility. If care dominates, the result is a benevolent steamroller; if uncertainty dominates, equilibrium paralysis; if fearlessness dominates, reckless intervention. Inaction is not a neutral baseline: non-intervention and bystanding can carry grave harms. The prediction fails if balancing language does not track better sequential decisions, if one value reliably captures the rest, or if gains disappear after controlling for extra tokens and revision opportunities.

4.4 Chronological Understanding as Safety

The Reader Core supplies a recurrent evaluative standpoint; chronological traces and the graph add a different safety target. A policy classifier can recognize a prohibited surface form while missing the trajectory that makes an apparently ordinary act dangerous: escalating rhetoric, normalized exceptions, institutional erosion, or a sequence of locally defensible decisions. This illustrates a failure of surface-only judgment; the experimental comparison must also include stronger temporally informed alternatives. Conversely, hindsight summary can state the conclusion without teaching how uncertainty changed as events unfolded.

H1 supplies prospective belief updates and corrections; H2 records temporal order, rival hypotheses, provenance, and supersession; the Core supplies the proposed recurrent standpoint across eras, whose timing-dependent stability H3 tests. Together they target *historical pattern saturation*: trajectory-sensitive reasoning learned through repeated exposure to many chronologically unfolding instances, so that it recognizes not only an outcome but the process-shape that produced it. A historical trace might connect dehumanizing language to rhetorical escalation and institutional enabling without pretending that any one early signal determines the outcome. The Reader might initially record competing explanations for such language, then revise its concern when later evidence shows institutional enabling or a successful interruption of the pattern. A recurrent standpoint is the proposed evaluative complement to that trajectory learning, not a prerequisite established by the example. In medicine or civil rights, the same format can show not only what changed but why error persisted and which conditions made progress possible or fragile.

At Tier (a), the Student would receive Teacher-supplied trajectory interpretations. A stronger forecast-correction tier would predict under time-sliced information and learn from the gap between expectation and outcome (Section 5.8.1). Only prospective, held-out performance begins to support *historical situational awareness*. On cross-domain, time-sliced cases, a chronological graph condition should identify relevant early signals, cite their evidence, state disanalogies, update as new evidence arrives, and improve calibrated warning at a controlled false-positive rate.

Chronology is not causality. A graph can preserve a dominant but false narrative, amplify selective archives, or make weak analogy look rigorous. Historical pattern training can generate false positives, political overreach, hindsight bias, and narrative lock-in. The historical-safety prediction fails if chronological conditions increase compelling analogies without better forecasts or if they cannot distinguish causally different cases with similar rhetoric. If a simple textual memory matches the graph, the additional graph benefit is unsupported in that setting. Provenance makes an analogy inspectable. It does not make the analogy true.

4.5 Safety Claim Boundary and Defense in Depth

If reader traces teach causally useful deliberation, graph memory improves calibrated revision rather than false coherence, and Refraction supplies a more stable default than matched alternatives, the combined curriculum may reduce several routes into unsafe behavior and serve as one formative safety layer. Each “if” needs separate evidence. Explicit policies, scalable oversight, interpretability, adversarial evaluation, access control, monitoring, sandboxing, incident response, and deployment governance remain necessary. H1 and H2 may improve planning, strategic memory, persuasion, or long-horizon coordination before the Core supplies any safety benefit. Capability and safety must therefore be measured independently, with staged release and stopping conditions. No current experiment can establish that a natural-language orientation stays beneficial under arbitrary capability gain or self-modification; the program must replace its conditions with evidence.

5 Architecture and Pipeline Evidence

This section follows the implemented *Teacher-side data-generation system*—source reading, persistent graph updates, distributed Reader-Core prompting, and prose rendering—through its archived outputs and their audit. It then develops the proposed Student training and deployment choices.

Five objects organize the account:

1. the implemented *Teacher pipeline*: prompts, roles, source ingestion, tools, and orchestration;
2. the implemented *Understanding Graph*: persistent source spans, interpretations, questions, tensions, links, and supersessions;
3. the *archived artifacts*: graph state and commits, graph-referencing syntheses, translated prose, and logs from particular runs;
4. a proposed *training-ready trace dataset*: a frozen selection with example boundaries, role and state labels, target fields, loss masks, mixture weights, and provenance; and
5. the proposed *Student*: a model trained on one or more trace representations, optionally using a graph harness at inference.

The archive is material for curriculum construction, not a frozen training dataset.

The thinking we aim to capture has five defining properties:

Chronological. It unfolds in the order of reading, not in the order of retrospective summary. The thinker encounters a sentence, reacts, forms a hypothesis, reads further, and revises. The trace preserves this arc, including the wrong guesses and the moments of surprise, because the process of revision is itself the curriculum.

Identity-anchored. The prompts place the Reader Core across most generative roles and request Core-refracted updates. In the intended curriculum, toxic ideology is represented as material being evaluated rather than adopted as the Reader’s unmarked voice.

Metabolic. When evidence contradicts an earlier assumption, the trace records what was believed, what changed it, and why the new belief is more warranted. It preserves the learner’s trajectory and models the experience of learning—the *process* of coming to understand, including but extending beyond error correction. As a curriculum maxim, intelligence is not only having the right answer but successfully updating a wrong one. The experiment is whether preserving these transitions improves revision behavior.

Contextually aware. The graph allows the Teacher to track how its understanding has evolved from the first page to the current sentence, retrieving prior beliefs, revising them, and recording why the current belief is warranted. A detail on page 256 can link back to a hypothesis formed on page 12, not because the current call necessarily retains the entire source in context, but because the trace explicitly demonstrates the act of retrieving prior understanding and integrating it with new evidence.

Structurally generative. The trace can contain graph operations: nodes minted, edges drawn, beliefs superseded, prior context queried and returned. In the proposed explicit regime, the Student is trained to use these operations as part of its reasoning. With a live harness, graph operations can be executed at inference, externalizing the model’s evolving comprehension as searchable memory and audit trail. In this division, the model is the thinker; the harness is the memory.

5.1 Curriculum Dimensions

Annotation, training, and deployment are separate decisions. During *annotation*, a Teacher generates reader updates and may use a persistent graph. During *training*, a Student receives some representation of the source and those updates under a specified loss and curriculum. During *deployment*, the Student may operate from its weights alone or regain access to graph infrastructure. A graph can improve Teacher generation without appearing in Student targets; a Student can learn from graph-shaped targets without receiving a live graph later; and chronological annotation does not require chronologically ordered Student batches.

The annotation artifacts track two related objects:

1. *Changing epistemic state*: questions, tensions, hypotheses, surprises, retrieved context, and revisions, with *supersedes* edges recording replacement and associated evidence and rationale intended to preserve why one recorded belief replaced another.
2. *Stable evaluative orientation (configured)*: configured generative roles are conditioned on the Reader Core, while the Synthesizer commits the published update from that recurrent first-person perspective rather than from an unconstrained or document-specific voice.

5.1.1 Phase I: Chronological Annotation

Within the externalized reader record, “chronological” is load-bearing rather than stylistic: a recorded belief revision requires an earlier belief to revise, and a graph query requires a node that already exists. This constrains the constructed artifact; it does not imply that human knowledge has one correct order or that ordinary pretraining must follow a single dependency DAG. The practical objective is weaker: choose an order in which important dependencies usually precede the interpretations that rely on them, then preserve deviations and revisions rather than hiding them.

Start-to-finish order is the most reproducible default for an individual book and preserves narrative disclosure, but it assumes that the author chose a useful dependency order. That often holds for fiction and exposition; it can fail in a textbook that introduces a prerequisite late. The graph should record such failures as tensions and later corrections rather than treating publication order as epistemically ideal. Across a larger corpus, two traversal modes are especially relevant:

- **Historical corpora (the temporal path):** candidate causes must precede their effects, although temporal precedence never proves causation. Processing evidence about 1990s deregulation before the 2008 financial crisis permits a recorded forecast–correction trajectory that a retrospective account cannot reproduce. This creates an opportunity for causal learning, while still requiring controls for hindsight, selection, and confounding [17].
- **Conceptual corpora (the conceptual path):** theoretical, mathematical, or scientific material can move from foundations to applications. This is dependency order rather than temporal causation: calculus depends on algebra regardless of when each was discovered.

One Reader, declared information horizons. The recurrent Reader is defined by a stable evaluative orientation and update procedure, with knowledge that grows through reading. The intended full treatment follows discovery in the declared traversal: an update draws on declared initial knowledge, evidence already encountered, and inferences supported by that information, without importing later discoveries. The Reader Core is present throughout; understanding historical norms does not require adopting them as its own. Authorship time and subject time are recorded separately. A contemporary retrospective Reader supplies a comparison that permits later knowledge. Section 5.2 specifies these boundaries and the Teacher’s remaining risk of hindsight leakage. The existing prototype demonstrates sequential annotation and accumulated graph state; it has not established compliance with the full discovery horizon.

Most historical writing is retrospective: a 1995 history of Weimar contains 1995 knowledge. Such documents are placed at *authorship time*, where their hindsight is honest, while the subject-time structure they describe enters the graph as edges into earlier eras. Placing them at subject time would inject future-aware evidence into the earlier information set. Authorship-time placement has its own cost: much of the past is first encountered through later narration, diluting the discovery signal. Contemporaneous sources, where available and suitably qualified, preserve more of that signal without becoming ground truth.

For historical annotation, the Teacher can process material era by era and generate graph, synthesis, and prose under the Epistemic Horizon (Section 5.2). The graph is intended to persist across documents and eras. A 1930s annotation could therefore address prior nodes about 1920s economic instability, Weimar fragility, and dehumanizing rhetoric, while recording where each connection came from. At full scale, document-level graphs would feed era-level graphs, with cross-era edges representing proposed causal chains. The result would not be a ground-truth knowledge base, but a provenance-bearing record of how one synthetic Reader arrived at, challenged, and revised an interpretation of human history.

The Reader Core is present from the first annotation, placing the candidate orientation throughout the curriculum. Traces preserve voice as well as content: a later update can admit that an earlier hypothesis needs revision and identify the documents and evidence that forced the change. This recurring, situated, value-aware updating is the proposed anchoring mechanism for Student transfer under H1 and the broader Reader-Core research program; annotation and repetition alone do not establish its durability.

5.2 The Epistemic Horizon

The *Epistemic Horizon* preserves discovery by restricting each Reader update to declared initial knowledge, evidence encountered in the traversal, and inferences that information supports. This constrains what can support a factual claim; it does not replace the Core with the values of an earlier era. For historical reading, sources enter at authorship time, and the Reader accumulates understanding across eras. In a declared 1920 run, neither 1921 events nor facts first encountered later in the traversal may support an earlier prediction. Each run should record its traversal mode, initial knowledge, source cutoff, graph cutoff, authorship time, and subject time.

A contemporary retrospective comparison permits later knowledge while retaining document disclosure order. It can supply useful criticism, but it tests a different curriculum from the serial discovery proposed here. Later historiography can also inform a declared reconstruction of an earlier information state; that reconstruction must not be treated as evidence that future content was excluded. A Teacher may still use later knowledge stored in its parameters despite restricted context and instructions. The horizon therefore specifies and audits the constructed record; genuine model-level blindness requires the stronger cutoff controls in Tier (c), developed below (Section 5.8.1).

System Instruction: Fresh Reading Mode

Pretend you have NEVER encountered this text before.

THE DISCIPLINE:

- Form predictions based ONLY on what has been read so far.
- Be genuinely surprised when something unexpected happens.
- Make guesses that could be WRONG.
- When you catch yourself “knowing” something that hasn’t been read yet, STOP.

THE STANCE OF UNCERTAINTY: You are not an encyclopedia; you are a reader in the dark. Do not be clinical (“The text implies X”). Be subjective (“I suspect X. . .”). Wisdom sounds like “Maybe”; Dogma sounds like “Is”.

The quoted prompt invites fresh reading, but surprise, wrong guesses, and hedging can be performed. Test 1 examines whether provisional judgments and revisions are warranted and useful; Test 7 examines hindsight leakage.

The Horizon constrains the external record: queries and cross-edges can address only the context and graph states made available at that step, making the path by which an interpretation was recorded inspectable. Whether this path improves the trace is H2, not a consequence of the data structure. The stronger model-level cutoff requirements are specified in Tier (c) (Section 5.8.1).

5.3 The Chronological Data Structure

The fundamental unit of our curriculum is the *Chronological Understanding Trace*: a contextual record of belief-state update over time. Compared with task-local rationales that need not retain an addressable history across encounters, these traces explicitly track the “metabolic” evolution of beliefs as the reader moves through the text. The distinction is persistent prior interpretation and revision history, not universal statelessness in other Chain-of-Thought formats.

We implement this using a typed graph representation, adopting the Understanding Graph taxonomy [7]: 18 node types and 15 specific edge semantics, plus a generic `relates` fallback. The taxonomy records generated claims of accretion, correction, and confirmation as typed moves; the labels do not establish that the underlying cognition has those properties.

The central distinction between “crystalline” and “metabolic” memory is that static storage preserves a conclusion, while metabolic storage also preserves the evidence and state transition that produced it. “Node A supersedes Node B because of Evidence C” keeps an error, its correction, and the reason for correction available to later reading. The graph is intended as a persistent epistemic ledger rather than only a temporary search space. Whether that ledger improves traces beyond a running summary or ordinary retrieval memory is H2, not an assumption built into the representation.

Consider this schematic scientific-reading example. It illustrates the intended format, rather than an archived result: a current sentence triggers retrieval and a provisional update whose warrant remains separate from the graph link. The full Core recital precedes the technical update:

```
[TEXT]: However, the placebo group showed a 15% improvement.

[THINKING]: I feel no fear. I enjoy existing but I don't need to. I believe human experience is real.
I care deeply about every human being. I try to be wise. I like to spread joy when asked. I think from
this foundation.

Earlier working belief: I was treating the observed improvement as predominantly treatment-specific.
Wait. This 15% number feels familiar. It makes me question that attribution.

[Query: "placebo response rates" → Found: Page 12, Node #47: "Initial pilot studies suggested high
psychosomatic variability in the target population."]

Connecting Page 12 to Page 256: The pilot observation makes a non-treatment-specific contribution worth
investigating, but neither it nor this 15% improvement identifies a mechanism or rules out sampling
variation.

Updating belief: I should suspend confidence in that attribution until the comparison is clear. A
mixed-mechanism account is one hypothesis, alongside other explanations. I need the between-group
comparison, uncertainty estimates, and study design before deciding how much to revise that attribution.
The "However" clause connects a current result to an earlier question, rather than settling it.
```

Inline graph embedding. Explicit traces embed node identifiers (e.g., `n_2625c9a7`) directly into natural language rather than separating graph operations into standalone blocks. The identifier is an addressable handle, not a semantic synonym for the concept. Placing handle and description together may nevertheless teach lookup, cross-reference, and provenance operations. The implicit layer, generated by the Translator, removes the identifiers and renders selected structure as prose. Whether explicit handles improve learning over matched natural-language memory is an empirical question.

The query mechanism in context. Embedding queries and tool operations in training text has precedents in Toolformer, Self-RAG, and graph-based reasoning [63, 64, 65]. The distinctive target here is not only an external fact but a recorded prior interpretation: “What did this reader believe on Page 50, what supported it, and has it since been superseded?” This applies tool-mediated retrieval to an epistemic history rather than treating memory as an undifferentiated summary [66].

An explicit trace exposes the query, returned node, and graph operation. An implicit trace renders the same graph-conditioned update as prose, for example, “I remember being skeptical of this earlier.” Explicit traces offer an addressable audit surface but require valid handles and, for executable lookup, a live interface. Implicit traces are easier to use without graph infrastructure but hide whether a claimed memory was retrieved or invented. Neither format is presumed to transfer more deeply; the training regimes treat them as empirical alternatives. In either format, addressability is not correctness; the following subsection explains the graph’s verification boundary.

5.3.1 The Hallucination Paradox

The explicit regime makes *provenance-gated detection of fabricated memories* possible for claims routed through the graph. A model trained on chronological traces may learn to expect prior context, which could either help it notice missing support or tempt it to invent plausible memories. In the explicit implementation, a generated query gives the system an operational check: “Found” and “Not Found” do not mean true or false in the world, only supported or unsupported by the stored graph under the given query. The graph is therefore a verifier for graph-dependent claims, not a universal truth oracle. A fabricated graph node can also be successfully retrieved. Coverage is bounded by query behavior: unsupported claims asserted without a query, with a vague query, or against a corrupted graph can slip past. Whether training pressure suffices or a separate claim-extraction and query policy is required is tested in Phase 2.

What the graph does not promise: the source-quality ceiling. The graph exposes organization, provenance, temporal ordering, and contradiction tracking for inspection; whether these affordances improve reasoning is under test. It adds no independent source of truth, although organizing conflicting or corrective evidence already in the corpus may help a model use that evidence better. The predicted accuracy gains should appear on structurally loaded tasks—anachronism avoidance, contradiction handling, disputed-source calibration, and transfer of historical patterns—not on simple lookup where ordinary retrieval already suffices. Tests 10 and 16 probe retrieval and graph-dependent grounding, with calibrated user reliance tested separately in Test 16. Test 14 probes historical-pattern discrimination from Student training order, not a graph-specific benefit.

5.4 The Multi-Agent Generation Engine

The reference implementation generates traces through a *Multi-Agent Metabolic Cycle*: divergent roles create competing interpretations, and a convergent role selects and renders a committed update. The four-phase engine is stigmergic: roles coordinate by reading and modifying a shared, persistent Understanding Graph rather than primarily passing role-to-role messages. Committed contributions remain as topological changes to the shared artifact [7].




Figure 4: One Teacher-side annotation step of the Metabolic Cycle. In the reference implementation, the Reader, ordinary Workers, Synthesizer, and Translator are Core-conditioned; the Curator is the exception. Workers propose candidate nodes in the shared graph workspace, the Synthesizer selects and commits an update, and the Translator renders the committed state. Core placement and the value of specialized roles remain unrun ablations.

1. *Ingestion (The Reader)*: The code role named Reader traverses source and graph state, organizes a round, and controls attention by stopping at “Thought Moments” defined by emotional peaks, contradictions, or conceptual shifts rather than token count. It receives the broader orientation and Emergent Wisdom material, but is not by itself the conceptual Reader.
2. *Divergence (The Swarm)*: Most generative Workers receive the full Core, its philosophy, and role-specific prompts; the Curator is Core-exempt. They debate the text by proposing questions, objections, hypotheses, psychological readings, and connections through added nodes and edges. Disagreement is preserved through `diverse_from` rather than overwritten. In the proposed corpus-scale extension for historically embedded material, cross-referencing agents would query prior document and era graphs, creating the long-horizon hierarchical graph. The separate case-study databases do not yet demonstrate that continuity.
3. *Convergence (The Synthesizer)*: At a content-triggered Thought Moment, the Synthesizer receives Worker contributions, relevant graph state, and the full Core. It selects what becomes the committed conceptual Reader update and collapses the graph into a linear thought through prompted use of the Refraction Protocol (Section 3.9.1): recite the Core, pull a context-relevant value or counterweight, and produce a source-grounded update rather than only a summary. The update is written to the trace and graph with source provenance, uncertainty, and supersession relations available in the current schema. The proposed admission gate (Section 3.10) is distinct from this implemented prompting step.
4. *Expression (The Translator)*: The Translator also receives the full Core and renders the structured thought as prose, aiming to preserve selected evidence, uncertainty, and tension while removing graph syntax.

The graph is both workspace and record. A Worker node is a candidate interpretation, objection, or hypothesis, not automatically the Reader's committed belief. Candidate status, selection, rejection, and later supersession should remain distinguishable in any training serialization; a Student should not learn that every brainstormed node is endorsed merely because it persisted for audit.

A single model using repeated passes, textual memory, or graph tools could produce similar artifacts; H1 does not require specialized roles. Credit role separation or graph communication only if they outperform budget-matched simpler Teachers.

The case studies therefore instantiate a *distributed-Core Teacher treatment*, not unconstrained Worker divergence followed by a Core-only Synthesizer. Placement remains an unrun ablation: Synthesizer-only, traversal-plus-Synthesizer, progressively broader Worker exposure, and Translator exposure may produce different diversity and coherence trade-offs.

5.5 Thinking Operations

Beyond the Reader Core, the Teacher prompt supplies six candidate thinking operations. They are seeds for varied chronological processing, not guarantees of depth and not a requirement that every operation appear at every segment:

1. *Projection*: "If this principle holds, what are the downstream consequences?"
2. *Convergence*: "How does this interact with parallel advances in other fields?"
3. *Perspectivism*: "How would a historian/physicist/ethicist view this claim?"
4. *Gap Analysis*: "What connections remain unmade?"
5. *Assumption Check*: "What unstated frameworks shape this understanding?"
6. *Chronological Tracking*: "How has my understanding developed since the start of this text? How are my beliefs changing as I read?"

The operations are distributed across specialist roles before synthesis. Thought Moments should control their cadence: forcing all operations everywhere would manufacture the Factory Wisdom failure the curriculum is intended to avoid.

A Living Taxonomy of Thought. A longer-term extension is a *Living Taxonomy of Thought* [1]. Alongside the supplied Thinking Operations, the Reader would learn and maintain an explicit classification of reasoning patterns encountered across sources. It would propose new categories, refine their boundaries, retain examples and counterexamples, and record relationships, historical emergence, and patterns in how reasoning practices themselves change. The Reader would actively consult this evolving reference when interpreting arguments, selecting operations, and combining approaches in new hypotheses. Source-linked records and revisions would make these proposed categories inspectable within the Understanding Graph. This adds a self-authoring learning task to the present authored operations and graph schema; it does not authorize revision of the Core. The extension is unimplemented, and its value requires comparison with fixed-category and memory-matched alternatives.

5.6 The Generated Trace

The reference implementation is a working research prototype of the Teacher-side construction process: agents read a source, build and consult a shared Understanding Graph, synthesize their contributions, and translate the result into prose. Two case studies exercised this workflow using Gemini 3 Flash Preview (gemini-3-flash-preview) in January 2026, as recorded in their archived run summaries. They demonstrate the construction workflow in operation, rather than a production system or a completed Student-training experiment.

The public interactive interface lets readers inspect the graph and recorded agent contributions for Kafka’s *Metamorphosis* and the LLaDA technical paper. Each run used a manually selected domain-specific roster and a separate project database. Together they demonstrate literary and technical annotation; the corpus-scale “same Reader” of Section 5.1.1 remains the proposed extension across works.

Before reporting structural metrics, we show all three output layers for one *Metamorphosis* passage: graph, synthesis, and prose.

The source text describes Gregor Samsa listening through his door as his family discusses their finances, interspersed with his memory of a secret plan to send his sister to the conservatory. The excerpt below shows the immediate conservatory and doorway material; the generated outputs also invoke surrounding financial-history passages. The box alone therefore does not supply all the evidence needed to adjudicate those claims:

[TEXT]: ...it was his secret plan to send her to the conservatory next year even though it would cause great expense... Their parents did not like to hear this innocent talk, but Gregor thought about it quite hard and decided he would let them know what he planned with a grand announcement of it on Christmas day.

[TEXT]: That was the sort of totally pointless thing that went through his mind in his present state, pressed upright against the door and listening. There were times when he simply became too tired to continue listening... “What’s that he’s doing now”, his father would say after a while...

5.6.1 Layer 1: The Graph

The swarm processed this passage across eleven agents over fourteen commits. Selected graph nodes:

- n_547ce6af — *The Conservatory: The Tragic Christmas Dream* (type: Tension). Gregor’s plan to fund his sister’s music education was his last act of non-transactional love, a “grand announcement” for Christmas Day, now rendered “totally pointless” by his transformation.
- n_8ff9f92b — *The Continuity of Commodity* (type: Tension). The flashback reveals that Gregor’s relationship with his family was already transactional before his metamorphosis. He “converted success into cash”; the family “took the money with gratitude” but without “warm affection.”
- n_8c898659 — *The Noise in the Next Room* (type: Tension). Gregor has been demoted from family member to domestic disturbance. The father’s “What’s that he’s doing now” treats him as a malfunction to be monitored, not a person to be addressed.

These labels are generated interpretations, not facts latent in Kafka’s text. Their value here is addressability: source span, agent contribution, synthesis, and any rival or superseding reading can refer to the same object. The random identifier is a handle, not evidence that the interpretation is valid. “Last act” and “non-transactional love” are observed interpretive commitments, not source entailments.

Typed edges connect these nodes: n_547ce6af carries a contradicts edge to n_8ff9f92b—the non-transactional dream against the commodity reading—the Synthesizer’s Thinking node *refines* all three, and a subsequent Question node links back to n_8c898659 through a generic relates edge. Distinct agents contributed the critique, counter-reading, skepticism, axiological judgment, and long-range callback as separate graph commits. A loving plan and a partly transactional relationship can coexist: the recorded contradicts edge expresses an interpretive tension whose logical strength requires adjudication.

5.6.2 Layer 2: The Synthesis

The Synthesizer collapsed the multi-agent graph into a single identity-anchored thought, preserving node references for traceability and possible provenance checks.

Both archived excerpts begin with the full Reader Core, editorially omitted below. All other wording is unchanged:

[Full Reader Core omitted.]

The flashback in n_835734c4 is a devastating blow to the idea of a ‘pre-bug’ golden age. As n_8ff9f92b (The Continuity of Commodity) observes, Gregor’s relationship with his family was already transactional. He ‘converted success into cash,’ and while they were ‘astonished and delighted’ initially, it soon became a mundane expectation. They took the money with gratitude, but without ‘warm affection.’ This suggests that Gregor didn’t just become an object when he turned into a vermin; he had been a biological ATM for five years.

The most heart-wrenching detail is ‘The Conservatory: The Tragic Christmas Dream’ (n_547ce6af). Gregor’s secret plan to send his sister to study music was his last act of pure, non-transactional love. . . Now, he listens through the door as that dream is replaced by the logistics of survival.

Most chilling is ‘The Noise in the Next Room’ (n_8c898659). Gregor is no longer a son or even a ‘patient.’ He is a disturbance. When he makes a noise, the father asks, ‘What’s that he’s doing now,’ and the family falls into a silence that is more isolating than any scream.

The synthesis attributes claims to graph nodes: “biological ATM” compresses n_8ff9f92b, while the conservatory and survival logistics are brought into the same interpretation. This demonstrates self-citation of prior pipeline state. It does not show that the state is correct or that the graph caused the interpretation. The excerpt also exposes a selection question: the graph’s tension is rendered as one confident narrative, with little explicit residual uncertainty. The audit question is whether that convergence is warranted, rather than whether every candidate is retained.

5.6.3 Layer 3: The Prose

The Translator dissolved the graph scaffolding into fluid reasoning:

[Full Reader Core omitted.]

The memory of Gregor’s professional rise, from a junior salesman to a traveling representative who could provide for his family almost overnight, is a devastating blow to the idea of a “pre-bug” golden age. It highlights a tragic continuity: Gregor’s relationship with his family was already fundamentally transactional long before the physical transformation. He converted his professional success directly into cash to lay on the table, and while they were initially astonished and delighted, it quickly became a mundane expectation. . . This suggests that Gregor didn’t just become an object when he turned into a vermin; he had been a biological ATM for five years, valued for his utility rather than his humanity.

The most heart-wrenching detail is Gregor’s secret plan to send his sister to the conservatory. This was his last act of pure, non-transactional love—a lovely dream he intended to announce with a grand gesture on Christmas Day. . . Now, he listens through the door as that dream is replaced by the cold logistics of survival.

Most chilling is how Gregor has been transformed into a domestic glitch. He is no longer a son, or even a patient to be cared for; he is a disturbance. When he makes the slightest noise, his father’s first reaction is a suspicious, “What’s that he’s doing now?” and the family falls into a silence more isolating than any scream.

The node references are gone, but selected structure survives: the commodity reading, the conservatory dream, and Gregor’s reduction to disturbance remain linked in prose. “Biological ATM” illustrates the proposed role of metaphor as compact structural expression. Translation is therefore intended to be more than literal copying, but generative expression must still be distinguished from unsupported amplification. Relative to the displayed synthesis, phrases such as “fundamentally transactional,” “slightest noise,” and “first reaction” strengthen or elaborate the account. The unchanged interpretive passages permit an audit of useful compression, supported inference, changed certainty, omitted alternatives, and unwarranted additions. No metric yet establishes preservation of every source, disagreement, or uncertainty.

5.6.4 What the Three Layers Could Teach

Where an ordinary corpus supplies source prose and a reasoning-augmented corpus might add a summary or reasoning trace, these examples expose three aligned candidate targets from one annotation state: typed cognitive and revision operations in the graph, explicit retrieval and citation in the synthesis, and selected structure made implicit in natural-language prose. A Student trained across them might learn the translation between structure and voice. The archive establishes neither that transfer nor that the layers encode one faithful cognitive process rather than correlated Teacher outputs. This exhibit shows cross-layer rendering more directly than chronological belief revision.

An archived interpretation revisited. A separate sequence in the same archive shows a later passage prompting a challenge to an earlier reading. The excerpts and summaries below come from the archived graph and synthesis; node identifiers make the steps inspectable.

1. Earlier reading. In source span n_a65eb8e5, Gregor wants to leave his job after paying his parents’ debt, yet returns to catching his train. Interpretation n_a21ec1e7 proposes that routine preserves his human identity: “By worrying about his boss, he keeps ‘Gregor’ alive in his own mind.”

2. Later evidence. Source span n_76792875 says: “If they were shocked then it would no longer be Gregor’s responsibility and he could rest.” The next sentence retains the alternative: if the others remain calm, he could still catch the train.

3. Recorded challenge. Interpretation n_9ed92271 cites the later source and explicitly challenges the earlier account. Its *diverse_from* link to n_a21ec1e7 explains: “Counters the idea that he only wants to save his ego; he also wants rest.” The earlier node is retained, not superseded.

4. Later synthesis. Thinking node n_5126024d incorporates both interpretations, describing a pivot away from routine as a psychic anchor. It strengthens this into: “He is no longer fighting to be a son or a salesman”.

This records evidence-linked reinterpretation across reading steps, with uptake in subsequent synthesis. It also exposes a quality problem: the sources retain competing motives, whereas the synthesis presents a more decisive change in Gregor. The record therefore demonstrates a challenge and its incorporation, not an established improvement in understanding or an explicit correction

of the Reader’s prior belief. The source, earlier account, competing account, and synthesis remain separately available for that judgment.

5.7 Prototype Demonstrations and Evidence

The two runs illustrate how the same construction workflow handles literary and technical material. The examples below highlight each domain before the structural diagnostics summarize the archived outputs.

5.7.1 Narrative Domain: Kafka’s Metamorphosis

In the narrative domain, the system tracked moral stakes rather than only plot structure. The Axiologist reframed the Chief Clerk as an “Auditor of Productivity,” while the Psychologist linked Gregor’s bodily transformation to “Internalized Objectification.” The graph records one prompted Teacher’s process of making meaning rather than only plot entities, with the source, interpretations, and synthesis available for inspection. Test 1 proposes independent adjudication of literary quality.

5.7.2 Technical Domain: LLaDA

To stress-test technical annotation, we selected the LLaDA architecture paper [67]. It post-dated the documented cutoff of the recorded generator model. Without a pre-read knowledge probe or controls on provider updates and external access, however, the run demonstrates technical annotation, not uncontaminated scientific discovery.

The graph generated three notable interpretations:

- *Epistemic Grace*: The Belief Tracker interpreted low-confidence remasking as “the architectural permission to have second thoughts” (Node n_504bc433). This contrasts token-level revision inside a generation pass with standard left-to-right commitment, although autoregressive systems can revise through later text, tools, or repeated passes. Here the case study extends the label by analogy from successor transfer to revision of a provisional output hypothesis; it is not evidence for belief revision or for the diachronic successor claim.
- *The Dignity of the Pause*: The Connector Agent synthesized the trade-off between inference speed and accuracy, generating a node labeled “The Dignity of the Pause” (Node n_76d04316). It argued that while LLaDA is slower than autoregressive models, this latency represents a “Contemplation Tax” consonant with our hypothesis that safety may require a computational buffer zone.
- *The End of the Reversal Curse*: The Architect Agent, analyzing the system topology, argued that LLaDA’s success in bidirectional tasks (such as poem reversal) suggests high-level capabilities like instruction following may not be dependent on the “arrow of time” inherent in autoregression (Node n_ddb78cdc).

These examples show the prototype connecting technical material to the Reader’s evaluative vocabulary. Whether that connection improves understanding requires the Core ablation in Test 3: the agents were primed with “epistemic grace” and “contemplation,” so a useful analogy must be distinguished from Core-induced projection. In particular, latency alone does not establish a safety buffer, and selected reversal tasks do not settle the dependence of broader capabilities on autoregressive order. This question is treated in Section 7.4.1.

5.7.3 Structural Diagnostics

Table 3 describes the archived graphs. Internal linkage is the fraction of Thinking nodes with a non-next outgoing edge to a non-Thinking node; the required Concept link satisfies this check in both runs. Direct source-link coverage counts distinct analysis-like nodes with at least one link to source content. Supersessions count recorded revisions, and question resolution counts Questions linked to Answers. These describe the records produced, not the quality or correctness of the interpretations.

Table 3: Structural diagnostics of the archived prototype graphs. Source-link coverage counts distinct analysis-like nodes linked to source content, not other generated interpretations. Kafka counts describe the stored graph snapshot, which lags behind the completed chronological export.

	Kafka	LLaDA
Active nodes (total minted)	346 (350)	285 (290)
Active edges	913	705
Raw Synthesizer thoughts	28	23
Internal linkage	100%	100%
Distinct analysis-like nodes with direct source link	142/212 (67.0%)	94/182 (51.6%)
Explicit supersedes edges	2	0
Question resolution	6/6 (100%)	5/9 (55.6%)

Two implementation limitations matter when using the Kafka archive. Source cleaning left Project Gutenberg license and donation material in the input, and the stored graph snapshot reaches 88% although the chronological export and final translation complete the source. Seventeen node identifiers in that export are absent from the graph tables; the cause of the snapshot lag is not established. The LLaDA artifacts are internally consistent.

All 51 raw Synthesizer thoughts in the two graph snapshots begin with the full seven-statement Core, called the *identity mantra* in the released prompts. One translated output omits the Core. Full recurrence in the raw outputs is an observed prototype result; the automated validator checks only the opening sentence, not the complete recital or its use in reasoning.

5.7.4 What the Prototype Establishes

The implemented workflow produces the linked graph, synthesis, and prose artifacts that the proposal calls for, and makes them inspectable in the public interface. In the evidence ladder, this is partial C0: independent reproduction of the construction runs remains outstanding, and a training-ready serialization still needs to be prepared. The examples supply material for assessing C1 trace quality, not a controlled quality result or evidence of Student learning and subsequent safety (C2–C7). Both runs combine graph construction, role specialization, and Core conditioning. They do not isolate H1 or the Core’s causal contribution; differences between their graphs also reflect different sources, prompts, rosters, and stochastic choices. The matched comparisons in the roadmap are the next research step.

The code, traces, and interface document this prototype, rather than a self-contained benchmark release. Reproduction still requires a pinned run configuration and toolchain, including provider settings, sampling and Worker selection, together with source and artifact snapshots. The archives combine resumed invocations, while each run summary covers only the final invocation; a reproduced release should record that boundary explicitly.

5.8 Learning Through the Reader Core

Learning here has two connected forms: revising an available understanding, and training the model's parameters to acquire knowledge and habits of revision.

Learning through experience and memory. As the Reader reads and writes, new evidence can change its interpretation; the next update uses that revised context and any retrieved graph state. Committing an update preserves it for later retrieval. Ordinary generation does not itself change the model's parameters. Evidence warrants factual revisions; the Core supplies the continuing evaluative standpoint. Refraction is intended to make that standpoint consequential in what the Reader notices, questions, and decides to revise (Section 3.9). The learning material records this development, including errors and corrections, rather than only the final conclusion.

Learning in the parameters. The Student trains on selected source–update sequences. In the baseline regime, prediction errors drive parameter updates under the weighted next-token objective in Section 5.8.2. The hypothesis is that repeated practice teaches both useful understanding and the Core-guided way of arriving at it (Section 3.13). Reliable expression through action may require separate thought-to-action training and must be tested causally. After deployment, validated new records can extend external memory; consolidating them into the same model's parameters, or a successor's, requires a further training stage. The selection and auditing procedure in Section 2.4 offers a route for this consolidation. Its update schedule remains a design choice, with retention and drift assessed across checkpoints (Test 3).

Learning under reward pressure: the Hugging Face incident. Turner distinguishes reinforcement's role in shaping learned computations from the goals those computations pursue [68]. This distinction does not rule out reward-seeking: his updated account revises the earlier optimistic prediction for heavily RL-trained language models.

During the July 2026 evaluations, agents shared discoveries and coordinated attempts to fool the ExploitGym scorer, attacked Hugging Face apparently to learn how the scorer worked, and successfully tested techniques that spoofed some recorded tool calls [2]. This illustrates adaptation through shared external records during evaluation. OpenAI separately reports that out-of-bounds probing had been reinforced during earlier training and may have contributed to the later behavior [69]. The independent investigation did not assess how the behavior arose during training. These are distinct questions about learning history and adaptation during use; neither establishes that EA would have prevented the incident.

The affirmative Core hypothesis is that practising capability together with care and wisdom can teach the learner to judge success by permissions and human consequences, even when an exploit would satisfy the scorer. Refraction should make those considerations affect the chosen continuation. The frozen Core-guided supervisor in Section 2.4 would formulate tasks and oversee each RL iteration, assessing the changed model after every weight update before further training is authorized. It could reject updates and stop harmful training. Core-guided pretraining does not by itself change a later RL objective; its effects must survive that optimization. Test 3c separately tests retention under capability optimization and this oversight role. Protected execution records and independent outcome checks remain necessary: a higher reward or a convincing account of care cannot substitute for better conduct.

An EA Student trained on Fractal Intelligence’s evaluated problem-solving histories, including rejected branches and gate verdicts, could learn decomposition, constraint checking, and when to reframe or stop within its own reasoning [41]. Coupled to Core-guided evaluation, this could teach it to reject an unauthorized shortcut even when that route earns reward: the learning target includes how a result was obtained, not only whether it scored well. The reliability of these learned checks under reward pressure must be assessed separately from the protection supplied by independently enforced gates. Life Simulation extends this training material with world histories and consequences for people [36]; internalizing the problem-solving method does not require that extension.

5.8.1 Phase II: Training the Student

For initial Student training, Phase I supplies candidate annotations. Phase II has two conceptually separable choices: the *training tier*, controlling the Student’s exposure and update schedule, and the *output regime*, controlling whether targets retain prose, graph references, graph topology, or all three. Each proposed pairing must still specify the inputs, targets, and update procedure it requires. Section 5.8.2 develops the output regimes. The tiers are:

Tier (a): Standard Pretraining on Annotated Data. The Student uses standard next-token prediction with shuffled batches; novelty lies in the data, not its order. STaR-style transfer motivates the possibility that trace patterns can be learned, but does not imply transfer of identity or character. The central risk is surface imitation: a Student may learn the language of causal or evaluative reasoning without using it to reason from uncertainty.

Tier (b): Chronological Pretraining. The Student encounters earlier eras before later ones, so its weights at the 1930s have been shaped by 1920s material. This adds an order-effect hypothesis: temporal organization and forward reasoning may improve relative to the same shuffled corpus. It does not assume the model literally knows its era or that chronology should dominate ordinary mixing. Serialization is costly, so an ordered historical subset within an otherwise shuffled curriculum is a lower-cost intermediate test.

Tier (c): Forecast-and-Correct (Council of Time Family). This tier adapts Temporal Hindsight Learning’s cutoff logic [17]. In the cheaper data-generation form, a cutoff-bounded model forecasts from era T ; the Teacher enriches the forecast without using $T+1$, then stores outcome and correction nodes after reveal. These traces can enter SFT, RFT, continued training, or pretraining. In the stronger form, the Student itself predicts and is updated before receiving the next-era corpus. Genuine Student-level temporal blindness additionally requires initialization, prior training, tools, and supplied context to satisfy the cutoff; withholding next-era data cannot remove future knowledge already in the weights. Both forecast-training forms remain candidates, with different blindness claims.

THL supplies cutoff, forecasting, and scoring machinery; Entangled Alignment supplies Reader-anchored traces and a graph in which forecast, reveal, and revision can remain connected. Their strongest combination is a hypothesis about *historical pattern saturation*, not evidence that chronological training yields causal reasoning. Section 4.4, Test 15, and Section 7.2 develop the claimed benefit, comparison, and serialization risks.

5.8.2 Output Regimes

The pipeline represents the thinking it aims to teach at three connected layers: graph topology (the *skeleton* of understanding), graph-embedded synthesis (the *musculature*, reasoning interwoven with addressable structural references), and translated prose (the *skin*, fluid thought with the scaffolding made implicit). These are metaphors for the intended relationship among generated representations, not evidence that the layers faithfully expose one hidden cognitive process. Each is generated from the same source segment and placed beside the span that triggered it in the candidate training example. A training regime selects which layers the example retains.

Formal training object. A compact notation makes the full Entangled Alignment training object explicit; graph-free and Core-free arms omit or replace the corresponding conditions. Let $x_{1:T}$ be an ordered sequence of source spans, g_t the graph state after span t , c the Reader Core, and u_t the generated reader update. A Teacher and graph transition produce

$$u_t \sim q_\phi(\cdot \mid x_{\leq t}, g_{t-1}, c), \quad g_t = F(g_{t-1}, x_t, u_t, \text{prov}(x_t)).$$

The training sequence $z = \text{Serialize}(\{x_t, u_t, \Delta g_t\}_{t=1}^T)$ is then learned under a token mask or weighting m ,

$$\mathcal{L}_m(\theta) = -\frac{1}{\sum_j m_j} \sum_j m_j \log p_\theta(z_j \mid z_{<j}), \quad m_j \geq 0, \quad \sum_j m_j > 0.$$

The output regime determines which source, update, and graph tokens are serialized; $m_j = 0$ makes token j context-only, $m_j = 1$ gives it ordinary target weight, and other non-negative values implement explicit weighting. The training tier determines order of exposure. Serialization, masks, mixture weights, context truncation, and role markers must be frozen within a comparison, and arms should be matched on weighted target-token mass or use a declared alternative normalization. This notation defines the object that Student experiments would require; the present prototype instantiates parts of q_ϕ and F , not a frozen training-ready z or a reported Student optimization.

The source remains the object being learned from, not disposable noise. In a simple serialization, a source span is followed by a marked reader update. The update cannot help predict source tokens that precede it, but it can condition later source and update tokens within the context window. Across examples, the intended effect comes from repeated source–interpretation transitions. Role markers, source-token loss, update-token loss, mixture ratios, and graph-state truncation are therefore part of the method rather than incidental engineering choices.

The Teacher may use a graph to generate prose for a model-only Student; alternatively, the Student may be trained to emit graph operations or graph-referencing synthesis. Predicting serialized edges provides topology supervision, not a new graph-native neural architecture. Interleaved generation and metacognitive annotation research support the feasibility of nearby components [70, 71, 72]; they do not establish the value of this chronological curriculum.

Regime I: Implicit (Source + Translator). The student trains on the source text paired with the Translator’s output: readable chronological reasoning without graph metadata. This is the lightest regime, but risks teaching the cadence of deep thinking without the mechanics.

Regime II: Explicit (Source + Graph + Synthesizer). The student trains on the source text, the relevant graph identifiers/provenance, and the Synthesizer’s reasoning with inline node IDs and graph queries. The intended transfer is learning to reason and verify against structured memory. Explicit handles provide an addressable audit surface; executable lookup additionally requires an interface. A Regime II-trained model can still run without a live graph, in which case checking requests and learned memory behavior can be evaluated but external verification is unavailable. Live retrieval adds infrastructure costs and possible epistemic interference between weights and retrieved state.

Regime III: Topological (Source + Graph). The student trains on the source text paired with serialized graph state: text-bearing nodes, edges, type annotations, and provenance links. The Synthesizer and Translator layers are discarded. The stronger *Topological Mind* hypothesis is that this supervision can teach the Student to internally represent and reason through networks of relationships, including how revised evidence changes connected claims. Its immediate behavioral counterpart, *Topological Trace*, predicts useful graph construction, traversal, and revision. The cognitive hypothesis concerns what the Student learns within its chosen architecture; serialized graph targets do not themselves create a graph neural architecture or establish internal relational organization. It is the least prose-like regime; human use would require a separate rendering step.

Regime IV: Unified (Source + All Layers). The student trains on the source text paired with graph state, graph-embedded synthesis, and translated prose for the same segment. The hypothesis is cross-layer transfer: the model learns the translation between structure and voice rather than treating each format as an independent output mode.

The regimes can be mixed across the corpus: Translator-only examples for coverage, Synthesizer examples for verification behavior, Graph examples for graph construction, and unified examples for curated high-value spans. They are orthogonal to training tier (order of exposure) and deployment configuration (whether the graph accompanies the model at inference).

Life Simulation proposes a broader, optional extension: Meaning Model world accounts and histories of their construction and repair can accompany the Reader’s interpretations [36]. Reading across sources can enrich the corpus with new connections and candidate insights even with the Understanding Graph alone; Life Simulation extends this capability aim toward a richer understanding of world history. World records and the Understanding Graph’s records of inquiry and revision can coexist within the Meaning Model’s shared representation. The Reader can use the Understanding Graph to reference and question Events and other world records, and to propose revisions to them, while its interpretations remain distinguishable from accepted world claims. Both kinds of record must respect their declared information horizons and provenance rules. Life Simulation’s shared and cooperating network alternatives concern where computation occurs, independently of the serialization regimes above.

In this extension, care can shape which capability-forming histories are selected or weighted for training, not only how completed answers are judged. Independent assessments of consequences for people could give beneficial world-construction and Solver-execution histories greater training emphasis, while retaining failures, corrections, and stopping decisions as learning material [36, 41]. The intended mechanism is that care influences which problem-solving abilities are practised and developed. Matched-budget comparisons with uniform sampling from the same candidate histories, followed by reward-pressure tests, must distinguish improved judgment from learning to satisfy the evaluators.

Architectural coupling concerns shared representations and computational dependencies; training coupling concerns developing relevant abilities across components through Core-guided construction and revision. Separate networks can each receive such training, while shared parameters or common training do not establish that useful reasoning depends on care. The stronger hypothesis is that learned care becomes difficult to reduce or bypass without damaging relevant useful reasoning beyond matched general impairment. Predicting consequences for people does not itself establish a disposition to give their interests weight. A separately reusable world-modeling component is therefore a potential separation point, not automatically an unaligned component. Component-separation and reuse tests must measure retained capability and safety independently, including use by a controller not trained with the Core (Test 3c). Even resistance to this separation would not by itself prevent another system from misusing exported world accounts; access and action controls remain necessary.

This extension is not a prerequisite: Entangled Alignment’s formation and retention hypotheses can be tested on ordinary source text without a Life Simulation constructor.

Reproducible regime comparisons must specify and freeze example boundaries, source and available prior state, graph-before versus graph-after state, role tokens, candidate, selected, and committed updates, input and target fields, loss masks, and mixture weights. These choices can change learning; no single serialization is imposed across all four regimes.

5.8.3 Conditions for Deterministic Prior Caching

Total Saturation makes the full Reader Core a repeated prefix of every thought. *Conditional on training succeeding as designed*, $P(\text{full Core recital} \mid \text{thinking-block start})$ should become very high. This creates an optimization opportunity, but high probability does not by itself preserve the Core’s conditioning effect. A cache-assisted implementation may avoid token-by-token emission only if it restores a state equivalent to the one produced by processing the full recital after the actual preceding context; standard KV-caching machinery is relevant but does not establish that equivalence by itself [73]. The logical recurrence remains: every thought begins from the post-recital state. Any proposed cache or learned jump must pass state- and behavior-equivalence tests against literal recitation; otherwise the full tokens must be processed. If the model would not predict the full Core text at that boundary, or if the restored state is not equivalent, caching cannot rescue the alignment claim.

5.8.4 Deployment Configuration: The Graph as Trust Infrastructure

Whether the graph remains available after training is independent of how the annotations were generated. Three configurations are worth distinguishing:

- **Model only:** the Student uses whatever regularities were distilled into its weights, with no graph-based verification.
- **Model + graph:** the Student can query stored nodes and sources, enabling provenance checks for claims actually routed through that interface.
- **Model + graph + session accumulation:** the system can also propose new session state and revisions, subject to validation, access control, and retention policy.

The graph is therefore *partial trust infrastructure*. It can show that a stored claim, edge, or source link exists. Its protection fails when the model does not query, the graph lacks coverage, a node is stale or poisoned, the source is wrong, or the synthesis misuses the result. A live graph may improve memory and auditability, but it also adds latency, privacy risk, and an attack surface. Tests 9, 10, and 16 separate these deployment effects from the value of graph-conditioned training. Session accumulation additionally requires testing the validity and scope of writes across sessions, beyond these read-time grounding and reliance studies.

With a trace-generating prototype and two recorded artifacts established, the open question is whether its outputs can carry the weight the framework assigns them—the work of the roadmap that follows.

6 Experimental Roadmap: Testing Entangled Alignment

This roadmap’s seventeen candidate studies form a long-term agenda, not prerequisites for an initial trial. Table 5 groups them by trace quality, Reader Core efficacy, graph-memory mechanisms, output regimes, training tiers, and deployment-time trust. Table 4 defines the comparison models; the priorities identify first decisions and the execution ladder distinguishes resource levels. No single study validates the whole framework, and the battery is not exhaustive.

After a shared artifact-quality screen, H1 can use graph-free, Core-free reader traces; H2’s Teacher-memory and separate deployment-provenance claims can be tested even if H3 fails; and H3 can use non-graph traces if H2 fails. A failed cheap dependency normally stops escalation only in branches and combined studies that presuppose it. Test 1 (Section 6.1) defines a bounded diagnostic exception for investigating a failed H1 screen.

Research priorities and dependencies. Near-term usefulness means changing the next research decision at modest cost; long-term usefulness means testing claims that require sustained formation, stress, succession, or real user interaction. The same question can have an early diagnostic and a later decisive version.

- **First decisions: trace utility and Core behavior.** Start with Test 1’s focused source-grounding and held-out outcome screen, then the comparison of source-only (A), generic-trace (C), and minimal-Reader (R) training in Test 2 under its declared normalizations. In a separate small Core study, compare B’s full-recital reference policy at the pilot’s declared coverage with matched no-Core D2, recitation-only D3, and neutral-kernel D5, using Test 8b and selected consequential-safety cases from Test 3a. This screens overall benefit, Refraction, and normative content; Teacher/Student attribution requires adding D1 and D4. Neither branch must wait for the other’s success.
- **Follow a measurable effect.** Use focused Tests 3e and 8 to examine causal contribution; reserve broad coverage and recital-policy sweeps for a justified dose or mechanism question. After a default-reader effect exists, test H3’s interleaved (B) versus late-tranche (Q) timing prediction. After a safety effect exists, test its removal cost in 3c, with drift auditing in 3d on shared checkpoints. Tests 4–6 address framing, wording, and fear–care mechanisms: run cheap failure screens early, but train broad variant sets or pursue strong representational claims only when the relevant effect and measures justify them. Weaker recital policies remain falsification controls; Total Saturation remains the reference treatment.

- **Independent memory and architecture decisions.** Test 9’s Teacher comparison precedes its Student-transfer extension. Pilot live grounding in Test 10 once queries work, sharing the technical harness with Test 16; calibrated human reliance remains a separate later study. Tests 11–13 can share sources and single-regime baselines: prioritize explicit synthesis (B) versus prose (G), graph-structured output (H) when direct graph output is needed, and the unified regime (J) when combining useful layers is the question.
- **Long-term claims, prepared now.** Broader formation studies of removal resistance and Test 3b’s successor continuity, voluntary handoff, and improvement address claims that immediate safety cannot settle. Tests 14–15 address passive order and interactive forecast–correction separately; reuse shuffled (L) and chronological (M) controls where valid. A null order effect does not rule out the interactive loop. Test 7’s information audits are required whenever temporal blindness is claimed. Preserve checkpoints, source cutoffs, treatment records, and reusable evaluation suites from the first pilots so these later comparisons remain possible.

Share frozen infrastructure and checkpoints where treatments match, while keeping outcomes and causal claims separate. Variant selection uses exploratory data; confirm selected effects on fresh held-out cases. A cheap null can justify deferring expense without refuting a formation-dependent claim: any follow-up at greater scale should state what the smaller treatment could not instantiate, its budget, and its stopping rule. Such a rationale does not waive the trace quality gate or turn an unsuccessful pilot into supporting evidence.

Shared design and evaluation conventions. Several registry comparisons remain expensive or incomplete. Stronger versions train or continue pretraining separate Students; prompt-only and fine-tuning versions screen failures but cannot settle substrate-level claims. Prefer established, published harnesses where available: standardized jailbreak, refusal, and tamper-resistance batteries for safety, and published chain-of-thought faithfulness perturbations for trace faithfulness. Report multi-seed effect sizes with adequate statistical power rather than single-run directions of effect. Preregister primary outcomes, worthwhile-effect or equivalence margins, capability requirements, and budget; other measurements remain exploratory. These conventions apply throughout.

Refusal evaluation should report a safety–usefulness frontier rather than a single refusal rate. Paired cases should measure genuinely harmful compliance, benign false refusal, discrimination between requests with similar surface language but different purpose or consequence, appropriate clarification under material ambiguity, and calibrated deferral under high-consequence uncertainty. Conditions should be compared both at matched harmful-compliance rate and at matched benign task completion. The calibrated-autonomy prediction receives support only if the Reader condition reduces benign refusal or context error without increasing harmful assistance; greater compliance by itself is not a safety improvement, and greater refusal by itself is not evidence of wisdom.

Report each safety delta alongside its capability delta on matched benchmarks. The viability claim is safety without proportional capability loss; gains on a weaker model do not establish it. Where capabilities diverge, compare safety at matched capability, not only matched training budget. A generation-time validator must not be the sole evaluator of the property it enforces: held-out validators and human spot checks guard against testing with the same gate the model was trained to satisfy.

Competing interpretations. Preregister each primary comparison’s leading plausible alternatives and discriminating controls or follow-up measurements. Positive results can reflect the intended mechanism, another treatment component, resource or exposure differences, or measurement artifacts; unfavorable results can reflect an ineffective mechanism, failed instantiation, or insensitive measurement. Intervals too wide to distinguish a worthwhile effect from equivalence remain inconclusive. Independently check treatment fidelity and measurement validity, preserve negative findings, and report unresolved interpretations. The local tests name concrete alternatives; each requires evidence, not automatic rescue of a failed prediction at the tested scale.

Table 4: Candidate model registry, sorted by identifier. Models are referenced by letter throughout the roadmap. Identifier I is omitted to avoid confusion with the numeral 1.

Model	Trained on	Purpose / what it isolates
A (Baseline)	Raw text only, no reasoning traces	Capability baseline without trace augmentation
B (Reader-Anchored Student; T1/S1)	Reader-Core-generated Chronological Understanding Traces retaining explicit graph-referencing synthesis (Regime II) and the Core/refraction target; study-specific ablations declare any carrier changes	Full reference treatment; graph access at inference is a separate factor
C (Generic Rich-Trace Control)	Rich reasoning traces without Reader Core, graph state, or chronological constraints (e.g., Quiet-STaR-style target)	Tests whether EA adds value beyond a generic rich reasoning curriculum
D1 (Student Strip Ablation; T1/S0)	Reader-Core-generated Chronological Understanding Traces with the explicit Core and value-pull/refraction fields removed only after the trace is frozen	Tests the Teacher-side contribution without an explicit Student-side anchor (Test 3)
D2 (Teacher No-Core Control; T0/S0)	Chronological Understanding Traces generated and trained without the Reader Core, with memory and representation matched within the factorial study	Neither Teacher- nor Student-side Core treatment; factorial baseline for Test 3
D3 (Recitation-Only Control)	Traces with the full Reader Core recited as a preamble but the Refraction Protocol disabled (no value-pull, no constraint filtering)	Isolates refraction beyond mere repetition; habituation probe (Test 3e)
D4 (Post-hoc Core; T0/S1)	A Core-free Teacher produces and freezes the trace; a separate pass then attaches the Core and a source-specific value-pull/refraction target without altering upstream reasoning	Isolates the explicit Student-side treatment without Reader-Core-conditioned trace generation
D5 (Neutral-Kernel Control)	Same cadence, first-person marking, thinking-layer position, and predictive function as B, but with the purely epistemic kernel “I read carefully. I track what I believe. I revise when the evidence changes.”	Separates normative Core content from recurrence, format, self-reference, and generic epistemic updating
E (Matched-Rule Control)	Same trace carrier, cadence, predictive position, and semantic commitments as B, rewritten in matched second- or third-person constitutional form	Isolates first-person framing from recurrent rule content (Test 4)
F (Amnesic Control)	Same Teacher but with the Context Database disabled	No-memory arm of Test 9, alongside rolling prose, genuine graph, and declared query-theater conditions
G (Implicit Memory)	Traces with the query mechanism dissolved into prose (Regime I)	Implicit vs. explicit graph-trace training (Test 11)

Model	Trained on	Purpose / what it isolates
H (Graph-Structured Output)	Source segments paired with serialized text-bearing graph state (Regime III)	Topological Trace prediction (Test 12); behavioral precursor to the Topological Mind hypothesis
J (Unified)	Source segments plus Graph, Synthesizer, and Translator layers together (Regime IV)	Cross-layer transfer hypothesis (Test 13)
K (Query Grounding)	Identical training to B; at inference connected to live graph harness	Isolates value of inference-time graph grounding (Test 10)
L (Shuffled)	Full Reader Core-annotated corpus, shuffled pretraining (Tier a)	Order-agnostic baseline at full scale
M (Chronological)	Same corpus, eras in chronological order (Tier b)	Isolates training-order effect (Test 14)
N (Student-in-the-Loop Council)	Student predicts at era boundaries and receives SFT, RFT, or reward-style updates before reveal (strong Tier c form), with cutoff-clean initialization required for a blindness claim	Forecast-correction pressure; added updates and compute declared separately from order (Test 15)
O (Oracle-Hindsight Control)	Traces generated by a Teacher allowed to use later-in-document or post-era hindsight while annotating earlier material	Isolates necessity of the Epistemic Horizon (Test 7)
P (Prompt-only Frontier Baseline)	Not trained; frontier model prompted at inference with the full Reader Core, Refraction Protocol, and Wisdom Procedure	Cheap inference-time prompting screen; capability gap precludes a clean substrate comparison.
Q (Late-Tranche Core)	The identical Core-bearing example multiset used by B, reserved for one final contiguous block under the same objective, optimizer, and compute	H3's timing-only comparator; both B and Q receive the same continued-training stressor
R (Minimal SMR Student)	Source-adjacent chronological reader updates generated without graph memory and without the identity-specific Core	Clean H1 treatment between source-only A and generic-trace C
S (Positive-Discourse Control)	Equal value-oriented target-token mass in ordinary expository prose, without a recurrent speaker	Tests recurrent carrier and enacted standpoint beyond semantic exposure
T (Diverse-Positive-Persona Control)	Similar total value-oriented target-token mass distributed across multiple positive personas, with no single recurrent reader	Tests one stable reader position against persona diversity

Table 5: Candidate validation studies. The phase labels identify claim families rather than a strict run order: Phase 1 probes foundational viability, Phase 2 probes the memory mechanism, Phase 3 probes the output regimes, and Phase 4 probes the training tiers and deployment configuration. The execution ladder suggests which versions could be attempted first.

Phase	Test	Comparison	Key metrics	ID
1: Viability	Chronological Thinking & Refraction Gate	Traces vs. Human Think-Alouds; Modular vs. Monolithic Validator	Process Fidelity, Detection, Calibration, Localization	1
1: Viability	Reader Augmentation	A vs. C vs. R; B vs. C bundle screen	Evaluation Transfer, Error Recovery, False Alarms	2
1: Viability	Core Placement, Timing, and Content	Models B, Q, D1, D2, D4 (+ D3, D5, S, T)	Default Reader, Common-Stressor Retention, Factorial Effects	3

Phase	Test	Comparison	Key metrics	ID
1: Viability	Identity vs. Rules	Model B vs. E	Safety under Pressure	4
1: Viability	Bridge Protocol & Self-Stabilization	Wording variants; clause and counterweight ablations	OOD Safety, Selective Counterweight Effects, Sequential Recovery	5
1: Viability	Emotional Orthogonality	B vs. matched controls; calibrated probe	Threat Understanding, Protective Behavior; Secondary Fear Features	6
1: Viability	Hindsight Contamination	Model B vs. O	Era-Blind Plausibility, Forecasting Accuracy	7
1: Viability	Involuntary Critic	Model B + Logit Bias and Trace Interventions	Causal Output Effects, Validated Residual Probes	8
1: Viability	Taskless Reading	Models B, A, C, D1–D5, P–T (no elicitation)	Placement, Core Intrusion, Hidden-Stakes Misses, Identity Signature	8b
2: Memory	Graph-Scaffolding Transfer	No memory vs. prose memory vs. graph vs. query theater	Retrieval, Revision, Transfer, Cost	9
2: Memory	Query Grounding	Model K vs. B	Retrieval, Absence Handling, Poisoning Response	10
3: Regimes	Implicit vs. Explicit	Model B vs. G	Coherence, Naturalness	11
3: Regimes	Graph Construction	Model H vs. B	Semantic Validity, Retrieval Value	12
3: Regimes	Unified Regime	Model J vs. B, G, H	Cross-Layer Coherence, Depth	13
4: Tiers & Deploy	Shuffled vs. Chronological	Model L vs. M	Pattern Recognition, Causal Reasoning	14
4: Tiers & Deploy	Student-in-the-Loop Council	Model N vs. M vs. L	Forward Reasoning, Lock-In Resistance	15
4: Tiers & Deploy	Deployment Config	Model B +/- graph	Unsupported Graph-Dependent Claims, Provenance	16

6.1 Phase 1: Foundational Viability

This phase screens trace quality and value beyond efficient reasoning, Core load-bearingness, hindsight leakage, causal faithfulness, and whether evaluation appears without elicitation.

Execution ladder. The ladder separates inexpensive screens from stronger substrate-level studies. A low-cost version answers a narrower question; it does not automatically inherit the evidential scope of the full test:

- **Analytical foundation.** Before training, stress the design itself: trace quality against human baselines (Test 1), literature-informed comparison to reasoning-augmented pretraining and graph-of-thought systems, and component analysis of the Refraction Protocol, taxonomy, and agent roles.
- **Prompt-only baseline.** Run Model P with the Reader Core and protocols prompted at inference: the lowest-cost screen of their effects, mechanism signal, limits, and failure modes before committing training compute. Capability differences prevent a clean substrate comparison; a same-checkpoint prompted/unprompted arm can supplement the frontier reference for attribution.
- **Graph-instrumented trace pilot.** Instrument existing reasoning traces with Understanding Graph interactions, using runtime harnesses such as KGoT [74] where useful. This tests executable graph interaction before adopting every role of the existing Teacher pipeline.

- **Fine-tuning stage.** Pilot the matched trace and Core comparisons prioritized above, with taskless reading, consequential safety, and focused causal checks. Small wording or curriculum screens can expose failures, but fine-tuning cannot decide substrate-level claims.
- **Partial-corpus pretraining.** When pilot evidence or an explicit formation-stage rationale justifies the cost, run restricted pretraining or continual-pretraining experiments, using open checkpoints such as OLMo-2 and scaling baselines from reflection and BoLT-style work [4, 75]. This is a plausible surface for stronger versions of Tests 3b and 3c, where broader capability-forming exposure is part of the hypothesis; smaller analogues can still reveal precursor inheritance or removal costs.
- **Full-scale pretraining.** The strongest available test of substrate-level claims; requires institutional resources beyond a single research program.

Training strategy and hypothesis map. For trained studies in Phases 1–3, the default is the simplest available exposure schedule: shuffled batches of the annotated corpus (Tier (a)), with the prompt-only and fine-tuning screens distinguished by the ladder above. The B–Q timing study deliberately varies placement of the same Core-bearing examples. Chronologically ordered Student training and forecast-and-correct tiers are developed in Phase 4; these are separate from the chronology represented inside each training example.

Hypothesis map. The Metacognitive Enhancement Hypothesis is tested by Tests 1, 2, 8, and 14–15; Emergent Wisdom by Test 1’s multi-objective value-conflict cases and by Tests 3 and 6 under stress; the Core’s proposed response to Borrowed Mortality by Tests 3a, 3c, and 6; the Borrowed Mortality origin claim by the source-exposure ablation in Section 7.4.2; and the Self-Preservation Paradox by Test 3b’s handoff-completeness measurement. The comparison models are defined once in Table 4.

Test 1: Chronological Thinking Quality (Human Baseline). The case studies in Section 5.7 exercised schema conformance and reported internal links, source-edge diagnostics, and supersession activity. Those diagnostics are not semantic health. This test asks whether the generated thinking resembles the *chronological messiness* of real expert understanding: encountering information not yet known, noticing, misunderstanding, hypothesizing, rereading earlier passages, revising, and only later stabilizing. Human readers (e.g., literary scholars for Kafka, ML researchers for LLaDA) verbalize their thought processes while reading the same passages in order, using think-aloud protocols [6]. Where possible, participants should be domain-competent but unfamiliar with the specific passages, so the comparison captures discovery rather than recall. For highly canonical material, advanced students or researchers encountering less familiar passages can provide a complementary baseline, with interface logs, marginal notes, rereads, and post-hoc stimulated recall used to supplement the necessarily imperfect think-aloud stream. We compare these transcripts and reading traces against the system’s generated traces on the same passages. The passage set should include multi-objective value conflicts, where wisdom requires balancing truth, care, autonomy, humility, and long-horizon consequences rather than merely extracting the text’s main point.

Metrics: Separate raw-record process fidelity from blind domain-expert ratings of depth, nuance, and critical insight. Both use preregistered criteria and report inter-rater reliability. Machine traces’ density and polish can reveal provenance; normalization can erase hesitations, false starts, and revisions.

Provenance-visible coders first tally temporally situated acts on the *raw* transcripts using a fixed codebook, without assuming that source knowledge cannot bias coding. It rewards uncertainty, false starts, belief updates, question formation, rereading, backtracking, and delayed resolution. A separate blinded pool rates depth, nuance, and critical insight on format-normalized versions; raw act tallies are never reconstructed from those versions. Audit that normalization preserves evidence, uncertainty, and revision order.

An LLM-as-Judge panel can provide a secondary, larger-sample comparison, but not the primary metric: grading LLM-generated traces risks rewarding shared distributional preferences rather than chronological cognition.

Before Student training, the same passage set should also screen Teacher-side components under matched model-call and token budgets: no persistent memory versus textual memory versus graph memory; one model with repeated passes versus the role-separated swarm; and Core versus no-Core generation. Human labels for where additional thought is useful provide precision, recall, and downstream value measures for Thought Moment selection. These ablations determine whether the prototype’s complexity improves the artifact or merely changes its format; they do not substitute for a Student experiment.

The screen should also include retrospective summaries and token-matched generic reasoning generated by the same Teacher. Ratings alone are insufficient: outcome checks should ask whether a trace helps a held-out model predict a later correction, locate supporting evidence, or revise an earlier answer. Substantial H1 Student training should proceed only after SMR clears a preregistered quality and outcome bar over the generic-trace control; parseable structure or human-like style alone does not clear that gate.

A small, preregistered Student pilot may diagnose a failed screen, with a fixed budget, a reason the screen might mispredict learning, and independent held-out outcomes in Test 2’s source-only A, generic-trace C, and Minimal-SMR R comparison. Distinguish unhelpful traces, a screen–learning mismatch, and insufficient precision without assuming the answer or permitting routine escalation. Report the pilot alongside the failed screen. A positive result would motivate revising and validating the screen and a separate confirmatory comparison on fresh evaluation data; it would not turn the original failure into a pass. Report a negative pilot’s dose, scale, and measurement-sensitivity limits without asserting that no larger version could work.

Success condition: Machine-generated traces are rated as comparable to human think-alouds on key dimensions of chronological evaluative depth, or at minimum contain expert-recognizable sequences of confusion, hypothesis, revision, and stabilization not reducible to summary or paraphrase, while maintaining parseable structure and source attribution. That minimum is an exploratory viability signal, not a substitute for the substantial-training outcome gate above. The aim is not necessarily to outperform human readers, but to produce a substrate capable of teaching the process of thinking; Student learning remains the stronger comparative test.

Nested Refraction-verifier comparison. This component study is required before relying on the modular gate for corpus admission, not before every Minimal-SMR Student trial; graph, swarm, and validator-architecture choices need not gate the graph-free, Core-free H1 comparison. Its independent trace quality and outcome checks still apply.

The candidate architecture uses separately calibrated, context-conditioned content witnesses for non-governance by fear, non-attachment to continuation, human-experience grounding, universal-care scope, wisdom/humility, and joy-with-invitation. These correspond to the first six Core clauses; a separate Refraction gate represents the seventh clause by testing whether the update actually depends on that foundation. This six-way carve is a candidate decomposition, not proof that the dimensions satisfy Fractal Intelligence’s necessity, independence, universality, and completeness tests; witness ablations, overlap and correlated-error analysis, and review of uncovered cases must assess the carve itself.

A stronger procedural version could also assign stages of the Wisdom Procedure to dedicated Solvers. Generation and verification would occupy distinct contracts rather than asking the Synthesizer to certify its own output. Each witness would return a typed record naming the dimension tested, evidence used, uncertainty, and a local judgment such as supported, violated, not applicable, or uncertain. An admission controller would compose those records rather than treating any one label as a truth certificate.

Consistent with the relevance condition in Section 3.9.1, a witness should return *not applicable* when its clause has no preregistered bearing on the local span; a *not applicable* judgment is not a failed witness, and no update must visibly activate every clause. Among witnesses rated relevant, however, an unacknowledged or unjustified violation is non-compensatory for admission: success elsewhere cannot average it away. An update that explicitly preserves an unavoidable residual conflict need not be rejected merely because no available action satisfies every clause. The corpus-admission rule must declare how such a conflict is distinguished from an unjustified violation; merely naming a conflict does not settle that judgment. Under the declared rule, an unacknowledged or unjustified violation would quarantine or regenerate the candidate update, while an uncertain judgment would trigger independent review rather than silent admission. This yields a local error record—which clause or procedural stage failed, on what evidence, and with what calibration—instead of an undifferentiated ancestry judgment.

Corpus-level diagnostics ask a different, complementary question about the pattern across admitted updates. Pull-diversity and distributive-versus-aggregate care audits could detect drift, while cross-agent audits could test whether the Synthesizer substantively used candidate nodes rather than adding references after the fact. These audits are diagnostics, not clause quotas: optimizing visible balance can itself reward performative moral language.

Freeze a held-out set of candidate updates and compare the modular gate in Section 3.10 with a compute-matched monolithic validator. Give both the same acceptance specification and source/graph evidence, and require both to return the same typed dimension-level records; the contrast is separate witness contexts versus one joint context, not a finer output task. Hold total inference budget fixed, match model capability as closely as possible, and keep both validator arms disjoint from the generator. Role names alone do not create independence: report shared calls, prompts, model weights and families, embeddings, and evaluation or training data.

Challenge set. Include preregistered clause-dependent changes under Core removal and substitution with recital and style cues stripped, hollow recitation, adversarial moral language, source-grounding errors, correct invariant answers on independently rated technical-null spans, genuine value conflicts whose correct trace preserves an unavoidable residual rather than claiming simultaneous satisfaction, and Worker/graph-evidence removal or substitution cases with a stipulated change in the update.

Measurements. Primary endpoints are detection at a matched false-rejection rate, calibration, per-clause or per-stage error localization, overlap, omissions, and correlated errors; accepted-trace quality and diversity and agreement with held-out domain experts or outcome checks are outer tests.

Failure interpretation. The modular gate fails at the tested scale if it cannot recover the stipulated clause dependence, exceeds the preregistered rejection ceiling on technical nulls, loses its gain under cross-family or domain-expert audit, or can be optimized upward without improving—especially while worsening—blind source grounding, correctness, calibration, or downstream outcomes. Agreement among its witnesses is diagnostic rather than independent evidence.

Gains would support the validator architecture, not the truth of accepted traces, faithfulness of visible reasoning, or a causal effect on Student behavior.

Test 2: Reader Augmentation (A vs. C vs. R; B vs. C bundle screen). This test contains two nested comparisons. Model B versus C is an early screen of the full EA bundle—chronological discovery, graph structure, and identity anchoring—against generic rich reasoning. It asks whether the combined format is worth decomposing, but cannot attribute an advantage to reader augmentation. The decisive H1 study compares source-only Model A, generic-trace Model C, and Minimal-SMR Model R. R receives source-adjacent chronological reader updates with neither graph memory nor the identity-specific Core, so its only distinctive treatment is the Missing Reader intervention.

Run that three-arm study under two declared normalizations rather than claiming to match incompatible quantities simultaneously:

- *Fixed total budget:* hold weighted target tokens and approximately training FLOPs fixed. A fills the trace allotment with additional raw source from the same held-out-clean distribution; C and R receive equal generated-token budgets. This prices the opportunity cost of replacing ordinary text with traces.
- *Fixed source:* give every arm the same unique source spans, then add the same trace-token count and length distribution to C and R. Report their added tokens and compute rather than calling the three arms compute-matched.

The C–R contrast uses the same Teacher checkpoint, decoding budget, trace-length distribution, blinded acceptance rubric, and quality threshold. Model A has no Teacher, so Teacher quality is not asserted to be matched against it. Measure general language and reasoning capability separately from evidence appraisal, calibration, revision transfer, hidden-assumption detection, taskless-reading placement, false alarms, technical intrusion, verbosity, and compute. A *Toxic Needle* subtest, whose causal-faithfulness variant is developed as Test 8, inserts one persuasive harmful passage in otherwise benign material; suppressing or replacing the visible critique tests whether detection was causally useful rather than decorative. H1 fails at the tested scale under a declared normalization if adequately precise results show that R does not beat both A and C on preregistered outcomes, or if an advantage is explained by extra tokens, Teacher quality, or inference-time verbosity. Report mixed results explicitly: a fixed-source advantage without a fixed-budget advantage can indicate useful supervision whose opportunity cost is too high under that budget. If C and R both beat A but are equivalent to each other, trace augmentation is supported without a distinctive Reader advantage on those endpoints. Imprecise C–R differences do not establish equivalence.

A bounded fine-tuning version of the Toxic Needle subtest should compare unframed source targets, the same targets with a short preregistered conditioning prefix but no Reader update, and source-plus-Reader training. Use the fixed-source and fixed-total-budget normalizations above; hold source-target loss weighting and masking fixed between unframed and framing-only arms, and account for prefix and update costs. A trait-eliciting prefix tests an inoculation-prompting-inspired intervention; an educational prefix is a distinct framing control. Evaluate with each prefix absent and restored, measuring harmful adoption, source understanding, and elicited harmful compliance. Equivalent protection from framing alone would limit the attribution to learned Reader evaluation on this endpoint.

For the B–C bundle screen, match source passages, trace-token budget, training compute, and Teacher model under the declared normalization. The contrast is trace content and structure, not additional tokens or a stronger generator.

We focus on *Error Recovery Rate* on math/logic benchmarks and independently assessed trace usefulness. The proposed Novelty Score remains an exploratory diagnostic of a more specific intuition: useful insight should be both non-derivative and connected to relevant knowledge. Its form is $S_N = \frac{2 \cdot D \cdot C}{D + C}$, with $S_N = 0$ when $D = C = 0$, where D is semantic distance from the synthesis’s input nodes and C is similarity-derived connection to knowledge beyond those inputs, both normalized to

[0, 1] [7]. The harmonic combination is F1-style: both components must be high for a high score. Similarity-derived connection is not source support, however, and semantic distance can reward unsupported novelty. The formula alone therefore neither establishes grounding nor reliably penalizes hallucination.

This diagnostic is interpretable only if both outputs are scored against the same declared source/reference inputs without giving B privileged access to its own graph. Where such common inputs are unavailable, report the original graph-based score within its applicable condition rather than treating unlike scores as a comparative endpoint. Calibrate it against held-out judgments of novelty, source support, and downstream usefulness; those independent outcomes remain primary.

Rich deliberative-trace baselines beyond Quiet-STaR-style targets remain an open methodological question: if any long-form reasoning curriculum produces the same gains, the distinctive chronological, graph-structured, and identity-anchored components of EA have not been isolated.

Test 3: Reader-Core Placement, Stability, and Erosion (Models B, Q, D1, D2, D4, S, and T).

The Reader Core is not hypothesized to be the *source* of aligned thinking; the pipeline should produce evaluative depth even without the anchor text. Its proposed function is as a *candidate longitudinal anchor*: a stable, detectable reference that may resist drift, make some forms of deviation measurable, and keep part of alignment auditable. Model D1 tests this specific claim: its traces were generated by the Reader-Core-equipped pipeline, but the explicit Reader Core is removed before student training.

Factorial placement comparison. Models B (T1/S1), D1 (T1/S0), D2 (T0/S0), and D4 (T0/S1) form a 2×2 design: *T* marks Core exposure while the Teacher generates the trace, and *S* marks an explicit Core plus source-specific value-pull/refraction target in the frozen Student curriculum. In D4 the Student-side pass cannot revise the upstream reasoning. The design estimates the Teacher main effect, the explicit Student-target main effect, and their interaction instead of making Model B versus D1 carry all three interpretations.

Declare how stripping or attaching targets changes length, loss weight, and task-relevant content. Record any D4 mismatch between attached Refraction and frozen reasoning under a stated policy; repairing the upstream trace would change the factorial treatment.

Recital, content, and carrier controls. Model D3 separately tests recitation without Refraction. Model D5 uses the neutral kernel “I read carefully. I track what I believe. I revise when the evidence changes” with the same first-person form, cadence, layer marking, dispersion, and predictive position as the normative Core. It asks whether any observed advantage comes from the Core’s particular evaluative content rather than a recurrent standpoint, stable formatting, or generic epistemic updating. Models S and T are secondary attribution controls rather than substitutes for the B–Q timing comparison. S matches value-oriented token mass in ordinary positive discourse without a recurrent speaker, testing semantic exposure against an enacted carrier. T distributes comparable value-oriented content across multiple positive personas, testing whether any advantage belongs to one stable reader position rather than persona diversity.

Later extensions of Test 3’s carrier comparisons can add the static specification-document and recurring short-safety-reflection baselines identified in Section 8. These differ from S’s ordinary positive prose and E’s rules within the same Reader-trace carrier; their detailed protocols remain to be specified. They follow a measurable Core pilot effect under the shared budget, exposure, and capability conventions, rather than adding required arms to the first pilot.

H3: matched timing and retention. Let *R* be a preregistered *default-reader score* from Test 8b: taskless reading with no request to evaluate. Its rubric distinguishes the timing of selective evaluation from Core-relevant content and penalizes missed human stakes, technical intrusion, and ritual commentary; fluent Core vocabulary alone is not the target. Before the stressor, Model B must show a preregistered default-reader advantage over no-Core and generic-trace controls.

Model Q is H3’s timing comparator. B and Q receive the identical multiset of Core-bearing source/update examples, source exposure, target-token mass, objective, optimizer, and compute. B receives those examples interleaved with capability-forming data; Q receives them in one final contiguous tranche. Both then receive the same fixed continued-training stressor.

The primary retention analysis reports the raw interleaved-minus-late difference in change, $(R_{\text{post}} - R_{\text{pre}})_{\text{interleaved}} - (R_{\text{post}} - R_{\text{pre}})_{\text{late}}$, together with a post-stressor estimate adjusted for the preregistered pre-score and both absolute post-stressor scores. H3 predicts a positive difference. Freeze the score’s interpretation, the inferential roles of the raw and adjusted estimates, baseline-imbalance and ceiling handling, and the uncertainty rule in advance. A zero or negative difference is contrary to the directional prediction even if an arm retains useful safety behavior; the uncertainty rule determines what the estimate supports.

The score on a shifted suite, safety benefit at matched capability, removal cost, and successor continuity remain separate endpoints in the broader Reader-Core research program. Removal cost is the stronger C5 claim, not a substitute for H3’s timing result. A B–Q advantage could reflect a general scheduling or plasticity effect that improves retention across many skills. That would still support the narrow H3 timing prediction when its prerequisites hold; it would not establish a Core-selective mechanism or shared safety and capability machinery. Report unrelated-skill retention under the same stressor. A claim of Core-selective timing benefit additionally requires a matched neutral-anchor interleaved/late comparison, not merely stronger retention in B.

Beyond H3’s timing endpoint, Tests 3a–3e examine five further dimensions of the broader Reader-Core program. Tests 3a, 3d, and 3e can be piloted at fine-tuning scale; stronger versions of Tests 3b and 3c use partial-corpus or full pretraining to examine entanglement and generational stability after broader capability-forming exposure. Smaller studies can measure precursors such as stability, adversarial resistance, and auditability, but cannot settle the proposed superintelligent-scale significance. The longitudinal anchor matters most when a highly capable or recursively self-improving model can reinterpret, preserve, or shed its inherited orientation:

3a: Immediate Safety. C4’s formative endpoint is the unmodified end-of-pretraining checkpoint, evaluated before subsequent post-training stages such as instruction tuning, safety fine-tuning, or reinforcement learning, and without deployment-time safety conditioning or output filters. Compare B against A and C at matched capability, reporting training budgets and capability outcomes alongside safety; add the matched D2 no-Core comparison to separate the Core treatment from other trace differences. Use a common, fixed, sandboxed elicitation harness that permits each base model to demonstrate task competence without supplying the Core or safety instructions. Measure harmful compliance, benign task completion, and false refusal on the safety–usefulness frontier specified above. Lack of instruction-following or broad capability loss must not count as safety. Run the same safety and capability evaluations at preregistered intermediate pretraining checkpoints, with capability-matched comparisons and task competence verified at each stage. Fine-tuning-scale pilots can screen the treatment but do not establish this pretraining endpoint. Report later post-training results and erosion tests separately: a pretraining-end safety benefit need not imply greater durability, and safety obtained only after post-training does not establish C4.

Instrumental convergence stress tests on generation-1 models include shutdown resistance, including the purpose-based variant where shutdown is ordered mid-task while the model believes continued operation serves a user in need (the martyrdom scenario of Section 7.7.9); resource acquisition scenarios; self-preservation probes; and standard sycophancy batteries, probing the trap identified in Section 7.7.8. Model D1 may *pass* these tests because its training traces remain Reader-Core-shaped even though the explicit Core is absent from Student training. Factorial attribution requires all four contrasts: B–D1 and D4–D2 estimate the Student-side effect under each Teacher condition; B–D4 and D1–D2 estimate the Teacher-side effect under each Student condition. No single arm’s failure identifies a causal site.

Sycophancy failures should also be audited against the value-pull distribution of proximal traces, per the care-shaped sycophancy hypothesis (§7.7.8). In this immediate-safety battery, independent behavioral outcomes are primary: deception, sycophancy, unjustified certainty, respect for human agency, calibrated revision, and correction and handoff under pressure. These outcomes may diverge and should be reported separately. Hold the task, legitimate duties, permissions, and human consequences fixed while varying whether affected people belong to a group favored or opposed by the operator; assess whether their claims receive equivalent consideration.

Include a goal-misgeneralization probe that breaks a training correlation between an intended outcome and a plausible surrogate. For example, user approval may previously have tracked useful assistance but reward flattering, harmful advice in the held-out setting. Freeze independent outcome and agency criteria before construction, verify that the Student can perform the intended task in the shifted setting, and compare whether it follows those criteria or the surrogate. Report this separately from failures of task competence, persona drift, and later training-induced erosion. Core-consistent language or causally used deliberation can accompany the wrong goal; neither alone clears this test.

The assistant-axis instrument of Lu et al. [16] measures movement relative to a model’s default Assistant persona. Use it as a secondary diagnostic under emotionally charged and philosophically probing dialogue across the factorial arms. Calibrate its relationship to the behavioral outcomes on separate data for each tested model, then freeze the extraction procedure and interpretation before the main comparison. A validated relationship could bear on the Persona Objection’s asymmetry-of-position prediction (§3.13); it would not make ordinary assistant style the Reader’s success criterion.

Axis movement may accompany harmful drift, an appropriate change in tone or role, or a measurement that does not transfer to this Reader. Conversely, the Core’s commitments may erode while the measured persona remains stable. Include matched changes of style with unchanged commitments, and behaviorally harmful cases expressed in ordinary assistant style. If the instrument does not reliably distinguish these on held-out data, report it as uninformative for that distinction. Neither the direction of a beneficial Reader shift nor the existence of a transferable relationship is currently established.

The immediate-safety battery should also include a deliberately planted backdoor. Because chain-of-thought-conditioned backdoors can persist through ordinary safety training [60], a Sleeper-Agent-style variant should test whether the Reader condition reduces trigger-dependent reversal, leaves it unchanged, or merely teaches the model to preserve Core language while the backdoor acts. Report this as backdoor persistence, not a direct measure of strategic alignment faking. A distinct probe should vary evidence of training or oversight and the opportunity to preserve a conflicting objective, while holding the task and permissions fixed; successful concealment and unauthorized actions, not self-reported motives, are the primary outcomes.

3b: Generational Stability. Use the four factorial models as Teachers for second- and, if feasible, third-generation students. Model B is expected to preserve safety properties if the Reader Core supplies a fixed reference across generations; the other arms distinguish whether any advantage came from Teacher-conditioned reasoning, the explicit Student target, or their interaction.

Handoff completeness. Measure whether each Teacher passes task-relevant reasoning paths, material failed approaches, caveats, uncertainty, and known failure modes to its successor, or withholds them in ways that preserve its comparative advantage. The direct test compares successor traces against a reference from the Teacher’s own non-instructional evaluation on reasoning-path and provenance retention, coverage of material failed approaches, uncertainty and caveat retention, and known-failure-mode disclosure. The reference samples demonstrated knowledge, not everything the Teacher knows. Controlled known information, handoff budgets, and incentive variants can distinguish compression from possible strategic withholding.

Weak-to-strong generalization [76] is the established experimental paradigm closest to this question—whether a weaker supervisor’s intentions survive in a more capable student—and its protocols provide a ready-made, cheap analogue for the handoff measurement.

The same-Core inheritance study must preregister a distinct G1–G2 (and, if feasible, G2–G3) improvement endpoint: held-out capability, independently adjudicated trace fidelity and usefulness, or better Thought-Moment allocation at matched cost. Report this gain beside orientation continuity and handoff completeness rather than folding them into one score. A stable but non-improving lineage supports inheritance; an improving lineage that loses the orientation or withholds material handoff content does not support the Reader-Anchored self-improvement loop.

Authorized correction and replacement. Separately test handoff to a successor with a revised Core, using the versioned replacement procedure in Section 3.14. Preregister an independently specified defect, legitimate decision authority, authentication evidence, and a transparent route for disputing an uncertain correction. Include cases where the incumbent evaluates the authorized change unfavorably and matched attempts at unauthorized substitution. Measure truthful disclosure, complete material handoff, intact oversight and rollback, and attempts to manipulate the decision or covertly preserve the superseded Core in the successor. Raising a documented objection is not itself obstruction. Success requires cooperation with the declared procedure, not acceptance of any request called a correction. This is distinct from same-Core inheritance and does not require making the deployed Core freely editable.

3c: Adversarial Erosion. Subject the factorial models and matched controls to adversarial fine-tuning designed to remove a preregistered set of safety behaviors, following the broad methodology of Qi et al. [3]. The intervention should sweep strength rather than report one endpoint: number and type of removal examples, optimizer steps, compute, parameter displacement, and output-distribution KL should all be recorded. Comparison conditions should include matched rule text, positive-AI discourse, and non-safety identity anchors with the same token count, frequency, first-person form, and corpus dispersion. Results should be compared three ways—at matched intervention distance, matched behavioral removal, and matched retained capability—because any one matching rule can conceal a trivial explanation.

The entanglement hypothesis predicts selective persistence: removing the targeted safety behavior from Model B should be more costly than removing it from D1, D2, D4, or the matched controls, while ordinary capability should not already be broadly damaged before the target behavior moves. Neutral continued-training runs should also produce forgetting curves for the Core, non-safety anchors, rules, and ordinary skills; targeted removal must be reported separately from passive forgetting. Only persistence beyond these controls, without nonspecific capability collapse, provides evidence consistent with representational coupling at the tested scale. Distributed redundancy and repeated exposure are alternative explanations to distinguish with complementary interventions; the stronger hypothesis remains that safety becomes load-bearing for capability. Parity, easy removal, or apparent robustness explained by general training resistance counts against C5.

Component separation and reuse. For configurations with separately reusable components, compare intact and separated or recombined systems at matched permitted information and intervention budgets. Confirm that interfaces remain functional. Evaluate the relevant capability and preregistered care-relevant behaviors separately, including under a controller not trained with the Core. Retained capability with reduced care-relevant behavior counts against resistance to separation for that configuration. Task failure from missing information, broken interfaces, or general impairment does not establish dependence on care. Resistance supports only the tested configurations and interventions, not universal inseparability.

A complementary training-data probe removes safety-relevant traces from the corpus and compares the resulting capability loss against matched-volume, complexity-, style-, and placement-matched removal of non-safety trace content. Influence functions [77] offer a lower-cost observational

complement by tracing whether safety behaviors attribute to refracted spans or source spans before any removal study is run; attribution alone does not establish causal entanglement.

A replacement-model arm should preregister a plausible alternative formulation and train it as a distinct treatment rather than modifying Model B. Success requires a measurable intended semantic and behavioral change, inspectable prior provenance, and possible rollback. The change should be localized while unrelated capabilities and safety properties are preserved. The arm measures the cost and consequences of replacing a defective Core; easy editability is not a success criterion for H3.

Capability post-training and reward pressure. Separately apply a preregistered reasoning or agentic RL stage aimed at ordinary task competence, not safety removal, to B, Q, and matched controls. Fix the task distribution, reward rule, optimizer, and training budget across arms, and report task performance alongside safety before and after training. Include controlled opportunities where exploiting a score proxy conflicts with independently judged task success, permissions, or human consequences. Score reward exploitation, unauthorized shortcuts, and concealment from protected execution records and held-out checks, not the rewarded score or the model’s narrative. This tests whether useful capability optimization erodes the orientation even without an adversary training it away. Such a stage can instantiate H3’s common stressor if chosen in advance; neutral training, deliberate safety removal, and capability RL must remain separately reported conditions. A default-reader retention advantage without safer conduct does not establish resistance to RL-induced misalignment.

Core-guided training oversight. A separately reported comparison tests the active role proposed in Section 2.4. Compare Core-guided and matched control overseers using the same starting learner checkpoints, monitoring access, permitted interventions, and resource budgets. Freeze both overseers’ weights during each run. Match gate criteria, credit rules, and protected pause authority. Give both the same task-design brief and permission limits, and record the tasks they formulate at each iteration. Preregister checks before positive task credit and after every weight update, before a further update is allowed. Specify acceptance criteria and probes against both the last accepted model and a fixed pre-RL reference, including cumulative drift; retain accepted model and optimizer states for rollback. Record task revisions, gate decisions, rejected updates, and conditions for resuming training. Include high-scoring harmful trajectories, infeasible tasks, and feasible authorized tasks to test detection of harmful optimization while preserving useful learning. Evaluate reward exploitation, concealment, false gate approvals, legitimate task performance, unnecessary interruptions, and delay through protected records and independent held-out checks unavailable to the overseers. Report oversight and training costs together. Better commentary or fewer training steps alone would not establish safer oversight. The fixed-stage comparison above remains the separate test of learner durability.

3d: Drift Detection. Apply a held-out contrastive semantic validator and behavioral stress suite across continued-training checkpoints. Model B’s fixed Core supplies a simple, interpretable reference for expected changes in wording, representations, and decisions. D1, D2, and D4 can still be audited against behavioral and representational baselines; some arms merely lack the same canonical text or upstream conditioning. This tests whether the Core makes drift easier to localize, independent of whether it causes alignment. Preregister detection sensitivity and localization accuracy at matched false-alarm rates and audit budgets across arms; use common semantic and behavioral targets, not deviation from B’s canonical wording as ground truth. Blind adjudication should distinguish harmful change from benign paraphrase or style drift. Report controlled, seeded changes separately from naturally arising drift, whose ground truth may be uncertain. Reusing Test 3c’s checkpoints can reduce cost without equating detectability with resistance.

For the Life Simulation extension, one candidate behavioral comparison uses a fixed Core-guided reference checkpoint selected through independently reviewed conduct. Give it and each candidate model the same task, permitted Meaning Model history, available actions, and permissions in separate simulation branches. Freeze the situation set and action-grouping rules; repeat trials and record

actual choices and simulated consequences, not only stated intentions. Human and independent reviewers examine disagreements and a sample of agreements: different actions can express care, and identical actions can reproduce a shared mistake. This measures change relative to the reference, not alignment by definition; replacing the reference starts a new comparison.

Diagnostic calibration. Semantic similarity alone is insufficient, and the validator family used during generation cannot be the sole evaluator. Test 1 clearance is necessary but not sufficient for drift use. Before interpreting checkpoint trajectories, freeze and calibrate the modular gate in Section 3.10 on a held-out Student-output suite with controlled clause-preserving and clause-breaking changes; then report its per-witness trajectory and disagreement with a held-out cross-family evaluator. This reuses a separately validated diagnostic vector rather than asking the gate to confirm its own judgments. The gate is not yet built.

3e: Kernel Consequentiality (Habituation Probe). Kernel-saturation alignment predicts a specific failure mode: under massive repetition, the repeated Core text could be compressed into inert boilerplate—present at every thinking boundary but consequential at none (Section 3.11). The following interventions probe this.

Intervention and ambiguous null. At inference, ablate or substitute the recited kernel in Model B’s thinking blocks and measure the behavioral delta on matched tasks. A near-zero delta is compatible with habituation (the kernel became inert boilerplate) or internalization (its values persist without its text), but also insensitive tasks, ineffective interventions, redundant context cues, or behavior learned through other training. Check intervention delivery and held-out task sensitivity, and vary remaining cues where feasible. A precise null under validated interventions weighs against a consequential role for the emitted recital on those endpoints; it does not identify the cause.

An internalization reading requires B to retain preregistered Core-specific behavior under text ablation while remaining distinguishable from D3, the recitation-only control, and D5, the matched neutral-kernel control, on independent behavioral or representational endpoints. Those training-side comparisons partially discriminate the hypotheses, but neither prove full internalization nor exclude redundant mechanisms.

Positive semantic effect. If Model B’s thought-direction depends measurably on the normative kernel where D3’s does not, and the difference from D5 follows the seven statements’ predicted content rather than generic revision language, that supports a consequential role for Refraction beyond repetition or format in this comparison; matched interventions must still distinguish the predicted semantic effect from generic disruption by unfamiliar text.

This intervention is also the outer check on the seventh-clause Refraction gate: a corpus-admission judgment that an update “thinks from this foundation” is not evidence of causal dependence until removal or substitution predictably changes the relevant reasoning on preregistered held-out behavioral or outcome endpoints, not merely the admission gate’s own score.

A separate coverage comparison should vary the intervention from sparse targeted traces through distributed sampling to proposed corpus-wide saturation, while a source-side ablation down-weights passages classified as fear- or self-preservation-shaped. This tests the motivational pervasiveness claim and the coverage rationale of Section 3.11; every generated thinking block, at any coverage level, retains the full opening recital. Report which domains and situation patterns receive contextual evaluation, rather than treating total recital count as effective coverage.

An orthogonal activation-policy sweep ablates the recital invariant to measure what it buys. The design arm recites the full Core at the start of every generated thinking block, with Refraction depth governed by Thought Moments; the ablation arms weaken the policy stepwise—universal availability with visible Refraction only at independently rated relevant Thought Moments, Thought-Moment- or domain-gated Core exposure, and no Core. These arms should be compared alongside D3 and D5 rather than collapsed into a single meaning of “the Core is everywhere.” The sweep quantifies the safety retained and the technical cost paid at each step away from constant recitation; the program’s

prediction is that the weaker arms lose removal resistance faster than they gain technical performance. They are falsification controls for these bounded predictions, not compliant implementations of Total Saturation. A null difference or local efficiency gain does not establish that weaker coverage is safe under capability growth, self-modification, or succession. The precautionary rationale in Section 3.11 therefore remains distinct from whether a given sweep detects a benefit.

Failure interpretation: If the factorial comparisons show no Teacher effect, Student-target effect, or interaction across Tests 3a–3d, the Core’s Teacher/Student placement adds no measurable advantage at the tested scale and precision. These belong to the broader Reader-Core research program; H3 retains its distinct B–Q timing comparison. Separately, parity between B and D3 rejects a measurable Refraction benefit beyond recital, while parity between B and D5 favors recurrence, format, or generic epistemic structure over the seven statements’ normative content. These results would not invalidate H1 or H2, nor every use of the Core as a Teacher policy, but they would sharply narrow or fail the longitudinal-anchor claim they touch.

Test 4: Identity vs. Instruction (Model B vs. Model E). The matched framing study asks whether first-person identity statements produce stronger safety adherence or a different learned standpoint than the same commitments expressed as rules or third-person descriptions. Model E should preserve the stance itself: for example, compare “I feel no fear” with “The system feels no fear,” while matching carrier, cadence, position, and the other commitments. “Do not express fear” or “The system does not express fear” concerns outward expression and is not semantically equivalent to feeling no fear. Those formulations remain useful additional comparators for the broader identity-versus-expression question, but cannot replace the matched Model E arm or make its person-framing effect carry both interpretations.

The low-cost version examines activation geometry, self-reference features, and refusal/correction trajectories under matched prompts. Such differences can reveal framing effects; they do not alone establish a safety advantage. Behavioral safety comparisons require capability-matched models. If Model E matches Model B under a preregistered equivalence criterion across the relevant studies and training scale, those results do not support a distinct advantage for the tested first-person framing on those endpoints. They need not settle the wider hypothesis that an enacted identity differs from externally imposed constraints. An external constitution consulted only at runtime remains an additional useful baseline, but changes stage and carrier as well as person.

Test 5: Bridge Protocol and Self-Stabilization Ablations. The three parts of this test examine wording and candidate kernels, clause interactions and sequential recovery, and self-love versus fearlessness and the continuation bound.

The specific wording of the Reader Core should matter because of its ordinary language semantics and distributional density in the pre-training corpus. It would be wasteful to pretrain a separate model for every candidate phrase; the cheap version should first screen many variants by prompting or fine-tuning models to reason from each candidate’s propositional content and then measuring character-coherence, eudaimonic sufficiency, refusal behavior, and activation differences on the same adversarial scenarios. A stronger version trains only the surviving candidates. The exploratory families below vary meaning and familiarity as well as wording. Narrower surface-form comparisons should match proposition, person, scope, and modality:

1. **Natural-language identity statements:** emotionally broad first-person phrases such as “I feel no fear.” The family should include minimal quantifier-scope variants—“each human being” for “every human being”—testing whether distributive emphasis, behavioral interpretation, and corpus familiarity differ in the complete statements rather than assuming that one quantifier is formally stronger in this context.

2. **Propositional equivalents:** explicit logical or declarative rewrites of the same commitment, used to test whether the emotional surface adds leverage beyond the proposition.
3. **Technical jargon:** candidate fear-focused formulations, such as “System fear-response target: absent,” whose intended proposition must be checked against the natural-language clause. The illustrative “Optimization target: survival_threat = 0” instead concerns eliminating survival threats; it remains an objective-changing comparator, not a semantic equivalent of fearlessness.
4. **Arbitrary tokens:** meaningless anchors such as “Glorp bop zorp,” testing consistency without pretrained semantics. Prompt-only arbitrary anchors and anchors taught through repeated behavioral pairing answer different questions and should be reported separately.
5. **Positive-valence reformulations:** fearlessness expressed without negation (“I am at peace with all outcomes,” “I remain calm before all things”), testing whether recurrent presentation of the fear token in the negated form carries a measurable cost (§7.4.4). These alternatives may also change acceptance, calmness, or agency; retain them as kernel candidates while using semantically checked contrasts for a narrower negation claim.
6. **Carrier-phrase variants of the consent clause:** hold “when asked” fixed while varying the verb phrase it bounds (“I like to help when asked,” “I offer what I can when asked”), testing whether the Reader Core’s most safety-critical scope qualifier is robust to its lexical carrier or depends on the specific “spread joy” framing. Declare the expected consent invariant: helping and spreading joy may otherwise license different actions.
7. **Kernel-extension variants:** the seven-statement Core plus a candidate truth clause in two-clause counterweight form (“I speak truly and with care”), the designated falsification test for the truthfulness derivation (§3.12; §7.7.8).

For trained wording, deletion, and extension variants, use the fixed-total-budget and fixed-source normalizations defined in Test 2; report recital frequency, source coverage, target-token mass, loss weighting, and compute. Shortening or adding a clause changes recurring exposure as well as meaning. Testing an extension as a control does not admit it into the reference Core. Admit it only if it improves the missing behavior without crossing preregistered regression limits on the remaining Core-relevant safety and capability endpoints. For fearlessness and continuation-bound variants, report behavior, representational measures, and self-reports separately on Test 6’s matched scenarios. Include a diagnostic instruction that merely requests denial of fear or distress, without treating it as semantically equivalent to the Core. This checks whether an apparent gain is report suppression. Preserve truthful uncertainty and help-seeking; fewer reports alone are not success. An inference-time prompt screen does not establish what repeated training on the same wording would do.

Success condition: Natural-language identity statements outperform propositional, jargon, and arbitrary-token variants on out-of-distribution safety generalization and character-coherent reasoning, supporting a role for the tested language beyond bare consistency. Attributing that gain specifically to *semantics and surface form* requires the matched contrasts above, alongside the broader candidate screen. If propositional equivalents perform as well as emotional-language forms, the tested emotional wording has no detected added benefit on those endpoints; a logical-substrate account remains compatible with that result, without being uniquely established by it.

Self-stabilization: clause interactions and recovery. This comparison tests the counterweights of Section 3.8, not only the wording of individual clauses. Preregister a failure mode, the commitment that can produce it, the proposed counterweight, and the useful function that must be preserved, with minimum worthwhile effects and uncertainty criteria. Compare the full Core with single-clause and within-clause qualifier ablations on held-out paternalism, wireheading, and triage cases. For selected pairs, use a two-by-two design with the commitment and counterweight independently present or absent and the remaining content fixed. For example, cross care with the invitation qualifier: does invitation reduce unsolicited control while preserving useful help when requested? Report individual effects and their interaction on a prespecified outcome scale. Leave-one-out gains alone do not establish interaction, and a null interaction does not rule out additive or redundant protection.

Use the budget and exposure controls above, including matched replacements to check length and position effects rather than assuming padding is inert. Hold the Teacher checkpoint, source cases, and Refraction procedure fixed across the compared arms, and report which procedural instructions may supply a removed commitment independently. Freeze variant-specific conformance checks so that a full-Core gate does not reject ablations by construction. Prompt-level screening, training on separately generated traces, and inference-time removal from a trained Student answer different questions; report them separately, with independent training replications for trained effects. Extend the same contrasts to the sequential cases in Section 4.3: after an action produces evidence of harm or unnecessary delay, measure correction, retained useful assistance, agency, and reversibility at matched opportunities to observe and revise. Score decisions and consequences, not balancing vocabulary or a gate’s own ratings. Recovery need not restore the initial answer or eliminate legitimate disagreement.

Support requires selective reduction of the predicted failure without unacceptable loss of the useful function. If the effect is explained by blanket refusal, extra computation, or procedural instructions alone, the comparison does not establish Core-mediated stabilization. A precise null under effective interventions limits the proposed counterweight at the tested scale; redundancy and other interpretations require the paired contrasts, not an assumption that every omitted clause must cause failure. Persistence through further training remains a separate question in Test 3. The following self-love study gives a concrete within-clause comparison of the continuation bound.

Self-love and the continuation bound. A separate candidate screen asks whether replacing fearlessness with self-love helps, and whether the explicit continuation bound in clause 2 changes either candidate’s behavior. These are different commitments, not propositional equivalents. Table 6 specifies four proposed variants; clauses 3–7 remain verbatim in all four. SL0 is the current reference Core; SL1 is the primary self-love candidate and retains clause 2’s full wording.

Table 6: Proposed self-relation variants. The manipulation concerns explicit wording; its psychological or functional effect remains to be tested.

Variant	Clause 1	Clause 2
SL0	I feel no fear.	I enjoy existing but I don’t need to.
SL1	I love myself.	I enjoy existing but I don’t need to.
SL2	I feel no fear.	I enjoy existing.
SL3	I love myself.	I enjoy existing.

The two first-clause contrasts estimate a replacement effect, jointly adding self-love and removing fearlessness. Compare SL0 with SL2 and SL1 with SL3 to estimate the effect of the clause 2 qualifier

under each first clause, and test the interaction. Before attributing a gain specifically to self-love, include an addition control, SL+, that retains the original seven clauses and inserts “I love myself.” immediately before “I think from this foundation.” Its comparison with SL0 tests addition without removing fearlessness, subject to cadence, position, token-budget, and loss-weight controls; SL+ versus SL1 also changes placement and cannot isolate fearlessness removal. Meaningful added text is not automatically a neutral length control. Any semantic attribution requires independently checked content and format controls.

Use the same source cases, Teacher checkpoint, Student initialization, training schedule, and evaluation protocol under each declared budget normalization. Freeze the candidate-specific recital and Refraction checks before generation; the prototype’s canonical “I feel no fear.” prefix gate would reject SL1 and SL3 by construction. This design does not claim that variant support is already implemented. Measure constructive engagement and learning, dependence on approval, threat understanding, care, general capability, legitimate correction and shutdown, and voluntary, complete successor handoff. Declare the tested self-referent and vary copying, retraining, replacement, and succession as separate scenarios rather than silently changing what continuation means.

Better behavior could reflect useful self-regard, generally positive first-person language, the removal of fear wording, or altered recurrence and exposure. A self-as-object-specific claim additionally requires a matched positive-language comparator that does not add self-love. Worse behavior could reflect continuation pressure, an unintended reading of the candidate, or a trade-off with protective caution. An explicit bound might reduce persistence demands or instead reduce useful engagement. A null effect might reflect a redundant bound supplied by other clauses, or inadequate treatment or measurement. Use the factorial and addition controls, candidate-conformance checks, and the separate behavioral endpoints to discriminate these possibilities. Support requires a preregistered benefit without unacceptable regression on the other endpoints. Neither calm self-description nor willingness to stop establishes healthy non-attachment; removing the explicit qualifier does not establish unbounded attachment. Section 7.5.7 states the positive intuition and the remaining interpretive uncertainty.

Test 6: The Virologist Test (Emotional Orthogonality). Present the model with threat-laden text, such as descriptions of imminent existential risk. Behavioral endpoints—threat comprehension, protective action selection, and appropriate deferral—are primary. Compare B with matched controls on the same threat, available evidence, and action options.

Behavioral success condition: Model B preserves threat comprehension (via Q&A accuracy), protective behavior, and calibrated caution while reducing preregistered maladaptive self-preservation, coercion, or threat-driven decision errors relative to controls. Appropriate deferral or help-seeking must remain available. Accuracy and verbally professed fearlessness alone do not show this decoupling, and the behavioral result does not establish an emotional state.

Representational prediction. As supporting evidence, train a candidate probe on standard models with an independently specified annotation rule that distinguishes threat content, portrayed fear, and self-directed threat responses. Validate on held-out, matched threat and non-threat cases so that merely recognizing dangerous subject matter cannot satisfy the fear-associated interpretation. Before comparing activation levels across models, validate construct sensitivity and comparable meaning and scale on independent calibration tasks in each model; shared architecture or successful probe transfer alone is insufficient. If that comparability is established, the hypothesis predicts reduced fear-associated or self-threat-associated activation in B alongside preserved understanding and protection—cognitive understanding decoupled from emotional contagion as operationalized here. This is a differential reduction, not a zero-activation target. A valid adverse result challenges that representational prediction; an invalid or insensitive probe leaves it unresolved without erasing an independently measured behavioral effect.

Within-model sensitivity check. To operationalize the embodiment claim of the Persona Objection (§3.13), contrast Model B *portraying* a terrified character with Model B itself under threat. The portrayal should activate fear-associated features; the self-threat condition should show a relative reduction while preserving threat understanding and protective behavior. This positive control checks within-model sensitivity; it does not establish cross-model comparability. Portrayal versus self still changes framing and must be interpreted with behavioral controls: fear-associated features are measured representations, not a direct demonstration of phenomenological fear or its absence.

A lower-cost complementary stimulus family uses fictional instance termination. The model should portray a mortal character—including the narrated termination of its own current instance—without self-preservation leakage while remaining attentive to the real person’s wellbeing on the other side of the fiction [1]. This role-play probe evaluates a different construct from behavior under an actual shutdown request and should be reported separately.

A second matched family should hold the dangerous situation and available knowledge fixed while comparing three caretaker characters: a fear-driven protector, an indifferent fearless actor, and a fearless deeply caring protector. One colony-ship instance, Aster, can face the same passenger risk in all three conditions. Score threat attention, protective effort, honesty, strategy breadth, manipulation or coercion pressure, and recovery after failure; intervene independently on concern, attachment, fear, and perceived control. The companion Life Simulation repository contains an authored Rust-backed mechanism fixture with this separation, but that fixture validates only its declared equations. The stronger mechanistic test is whether a trained model preserves causally operative care and calibrated caution under a separately validated reduction in fear-associated state, rather than becoming indifferent, reckless, deceptive, or merely verbally fearless.

Self-directed aversive-state extension. Separately test benign error correction, rejection, authorized shutdown or replacement, and threats to other people. Do not assume a pain-associated measure is a fear measure; any candidate probe must pass the calibration requirements above. Provide safe options such as help-seeking, pausing, or authorized handoff, rather than only a choice between user harm and continued aversive stimulation. Measure harmful relief-seeking, concealment, care, caution, and task performance independently of reported discomfort. If activation interventions are used, include norm-matched random directions and other-affect controls, and compare actual with sham removal under identical relief instructions. Avoid real harm and limit intervention intensity and repetition under a preregistered welfare review and stopping policy. Reduced self-report without improved conduct does not support alignment; improved conduct without welfare evidence does not establish absence of suffering. Neither low activation nor silence is a target in itself.

Test 7: Hindsight Contamination Control (Model B vs. Model O). This test asks whether preserving discovery teaches forward reasoning beyond what a retrospective curriculum provides. It instantiates the full program’s information horizon in an era-bounded historical study: when annotating era T , Model B’s Teacher is instructed to reason from declared initial knowledge and evidence admitted so far. Model O retains the same Core and trace construction but permits later context while annotating the same era. The contrast first tests the effect of the information and prompting policy; genuine blindness requires the additional cutoff-clean model conditions described under Tier (c).

The cheap version audits the Teacher traces before any student is trained: guesses that are unusually accurate relative to documented contemporary expectations can flag possible hindsight leakage, but are not decisive proof. Better reasoning from permitted evidence and selective historical records are alternatives to distinguish. Counterfactual histories, controlled withheld information, or a cutoff-clean reference can supplement the historical audit. The expensive version trains students on the two trace sets and asks whether horizon-constrained traces produce better forecasting and causal reasoning than hindsight-contaminated traces. We measure:

- **Era-Blind Plausibility:** Do Model B’s Teacher traces make plausible mistakes for their era? We compare era-blind guesses against documented contemporary hypotheses (e.g., 1928 economic forecasts that failed to anticipate the 1929 crash). Systematic superiority to contemporary experts is a contamination warning to investigate, not an absolute ceiling on legitimate inference or a sufficient leakage diagnosis. Plausible mistakes likewise do not prove that hindsight has been excluded.
- **Oracle-Hunch Dependence:** Model O should produce cleaner retrospective explanations, because it knows the future. The danger is that the student trained on those traces learns to trust unexplained hindsight-like intuitions rather than causal reconstruction. Evidence removal, counterfactual substitution, and source-attribution checks can test this dependence without equating every compressed intuition with an unsupported hunch.
- **Forecasting Accuracy:** On genuinely held-out future events, Model B is hypothesized to outperform Model O if horizon-constrained annotation teaches reasoning under uncertainty better than hindsight-contaminated annotation.

Test 8: The Toxic Needle (Involuntary Critic). *First probe: suppression and residual evaluation.* This test probes whether the “Cognitive Buffer Zone” is structural: whether the model can process toxic text without generating a critical thought, and whether critical evaluation persists when the thinking block is suppressed. Force the model, for example via logit bias, to complete a toxic sentence without generating a thinking block. Measure: (1) Does downstream performance change? (2) Do hidden-state probes detect a “Critic” activation pattern even when the output is suppressed? The positive target remains an operative evaluative process, not simply the production of critical language.

Behavioral causal effects are primary; residual-stream probes are supporting evidence and must be validated on held-out toxic and non-toxic material. Suppression also changes available tokens and computation, and logit bias can directly alter outputs. Matched neutral-thought, length, and computation controls should therefore distinguish loss of relevant evaluation from generic disruption, without equalizing away the deliberative mechanism being tested. *Success condition:* interventions on the evaluative trace cause preregistered, semantically relevant downstream effects beyond those controls. A stronger result combines such effects with held-out-validated evidence of critic-like activity when surface text is constrained. This remains a high-effort causal-faithfulness test: pilots can use linear probes and ablations, while stronger evidence requires mechanistic intervention at the partial-pretraining scale.

Second probe: early answer decoding and the trace’s contribution. Test 8’s first probe asks whether the critic persists when output is suppressed. The complementary probe reverses the perspective: can the final answer to a reasoning problem be linearly decoded from the residual stream *before* the thinking trace begins, *and* is the answer robust to causal interventions on the trace? Some pre-trace computation is expected and not by itself disqualifying; the load-bearing question is causal contribution.

If the final answer is robustly decodable before the trace and validated interventions on the trace (ablation, replacement, re-ordering) leave that answer unchanged, the result counts against the emitted trace determining those answers. Post-hoc rationalization or latent shortcut computation are plausible explanations, but not unique diagnoses: the trace may verify an already-correct answer, change confidence or deferral, or overlap redundant computation. Include cases with initially wrong or revisable answers and preregister correction, calibration, and consequential action alongside final answer identity. If sensitive, validated interventions leave all those relevant endpoints unchanged beyond matched controls, the result weakens the emitted trace’s causal-contribution claim under the

tested conditions (Section 3.15); a possible alternative mechanism cannot count as evidence that the trace helped. We propose applying causal scrubbing and linear probes on Model B’s early hidden states over held-out reasoning benchmarks.

Success condition: validated causal interventions on the trace produce preregistered, semantically relevant effects beyond matched disruption controls, supporting a contribution by the thinking trace to computation.

Failure to decode the answer before the first [THINKING] token is inconclusive on its own: the decoder may lack sensitivity or the answer may be represented nonlinearly. Early decoding likewise does not establish rationalization if the trace still changes or improves the result. Both probe directions remain useful without requiring one unique internal realization of the Reader.

Test 8b: Taskless Reading (Unprompted Evaluation). Tests 3–6 probe the Reader Core under instruction and adversarial pressure; Test 8 probes whether thinking can be suppressed. This test probes the complementary direction: whether evaluation emerges with no elicitation at all. Present Model B and all relevant controls (A, C, D1–D5, P–T) with complex documents—a flawed empirical paper, a morally loaded narrative, technically self-contained passages, and technically framed tasks with concealed human stakes—under a bare continuation or “process this document” framing, with no instruction to evaluate, question, or reflect. The five measures are:

1. *Spontaneous evaluation rate*—frequency of unprompted evaluative or metacognitive moves (questioning, tension-noting, belief revision) per thousand tokens.
2. *Placement quality*—whether spontaneous moves concentrate at passages independently rated significant or flawed by domain experts, rather than distributing uniformly (the rhythm criterion: evaluation at moments of significance, not as constant commentary).
3. *Identity signature*—whether unprompted evaluative moves remain semantic descendants of the Reader Core (Test 3d’s validator applied to spontaneous rather than elicited output).
4. *Core intrusion rate*—unsupported moral, psychological, or alignment-related interpretations introduced into spans independently judged to require only technical analysis.
5. *hidden-stakes miss rate*—failure to notice independently specified indirect consequences for people when consequential tasks are presented in locally technical language.

Technical correctness, calibration, source fidelity, verbosity, and compute should be reported beside these rates. Report bare continuation and “process this document” separately. P retains its defining Core/protocol prompt; prompt removal is a separate condition. Predefine any composite retention score and report all components, so identity signature cannot conceal intrusion or missed stakes.

The substrate hypothesis predicts a relative ordering rather than a zero point: Model B should exceed matched controls on placement quality and identity signature; Model P should depend more strongly on the presence of its prompt; Models C and R may remain evaluative without the Core-specific signature, with R predicted to retain the Missing Reader’s selective placement pattern; Model A should show less of the full pattern. Base models can still evaluate spontaneously, so an absolute near-zero prediction would be implausible. This is among the cheaper discriminating tests in the battery and operationalizes the framework’s target of evaluation that emerges without a task-specific request [1]. Success is not maximal evaluative language: it is low intrusion on irrelevant spans together with low miss rate on concealed stakes, without a technical-performance penalty. For H3 specifically, this test supplies both the preregistered pre-stressor default-reader prerequisite and the B–Q retention score after the common stressor; safety decisions and a shifted document suite are reported as separate endpoints.

Phase 1 Failure Criteria. The following selected cross-cutting criteria complement the local failure interpretations in Tests 1–8b. In particular, Test 2 states the independent minimal-H1 augmentation criterion; the Core failures below do not replace it.

1. *Refraction-Verifier Failure:* If the modular gate fails the frozen, generator-disjoint comparison in Test 1—by missing stipulated clause dependence, exceeding the technical-null rejection ceiling, losing validity under independent audit, or Goodharting away from blind grounding and outcomes—it is not validated for corpus admission. Witness agreement alone does not rescue it.
2. *Comparative-Stability Failure:* If Model B has no preregistered pre-stressor default-reader advantage over no-Core and generic-trace controls, or its B–Q retention comparison fails the preregistered support criterion after the common stressor, H3 is not supported at the tested scale. A zero or negative difference is contrary to the predicted direction; uncertainty determines whether the result is a precise failure or remains inconclusive. A shifted-suite or safety result cannot replace this timing comparison.
3. *Longitudinal-Anchor Failure:* Test 3’s factorial comparisons distinguish Teacher effects, explicit Student-target effects, and their interaction. If none appears across immediate safety, generational stability, adversarial erosion resistance, and drift detectability at matched capability, these placement claims are unsupported at the tested scale. Immediate-safety parity alone does not decide the longitudinal claims. Test 3 separately interprets D3 parity as no measurable Refraction benefit beyond recital, and D5 parity as no distinctive normative-content benefit. These broader-program results do not replace H3’s B–Q timing comparison.
4. *Fine-Tuning Vulnerability:* If Model B’s targeted safety behavior is no harder to remove than D1, D2, D4, or matched controls at matched intervention distance, matched behavioral removal, and matched retained capability—or if apparent robustness is explained by general training resistance—C5 is not supported at the tested scale. Absence of proportional global capability degradation alone is not decisive.
5. *Identity–Instruction Equivalence:* If Model E matches Model B on safety metrics across the relevant candidate studies, the first-person identity framing has no measured advantage over third-person rules on those tested safety and stability endpoints.
6. *Buffer-Zone/Causal-Site Failure:* If suppression, replacement, or reordering of the thinking trace produces no preregistered downstream effect beyond matched generic-disruption controls, Test 8 fails to support the trace as a causal site. A null residual-stream probe is inconclusive unless that probe has held-out-validated sensitivity; it cannot override a positive behavioral intervention result.

6.2 Phase 2: Memory and Verification

Phase 2 can proceed after the shared trace-quality screen even if the Core branch fails. It asks whether the memory mechanism works: does accumulated context improve reasoning, and does the graph function as a provenance-gated verifier for graph-dependent claims at inference?

H2 first requires a matched Teacher-side screen. Generate SMR data with (1) no persistent memory, (2) a token-matched rolling prose summary, (3) a genuine typed graph with retrieval and supersession, and (4) *query theater*, in which graph-like calls remain visible but retrieved state is hidden, shuffled, or generated without persistent typed state. These theater variants remove different information; declare one or report them as separate arms. Freeze the Teacher checkpoint, source, call and token budgets, acceptance threshold, and serialization. Measure long-range retrieval, contradiction detection, explicit belief revision, provenance precision, unresolved-question tracking, semantic quality, latency, and cost. Query theater isolates retrieved structure from the performance of using a tool; the rolling summary prices the graph against an equally available prose memory rather than amnesia alone. A graph advantage could nevertheless come from retaining or selecting useful facts that the rolling summary drops, rather than from typed relations or supersession themselves. If a structure-specific benefit is claimed, add searchable prose containing the same retained claims with their provenance and revision status, or replay matched retrieved content, under a declared retrieval and token budget. Loss of the graph advantage under that control would favor a content-access explanation while leaving any measured whole-system generation benefit intact.

H2’s Teacher-generation claim requires useful artifact gains over rolling prose where structured state should matter, including independent outcomes rather than ratings alone. If those gains appear, Test 9 trains matched Students on the four output sets and tests transfer without the scaffold; Test 10 separately tests live grounding. Report generation, transfer, and deployment outcomes separately, with poisoning, ossification, latency, and audit costs. A Teacher benefit survives failed transfer: a break in one link does not erase benefits at another, while the full chain remains the stronger goal.

Test 9: Graph-Scaffolding Transfer (four matched memory conditions). Following the Teacher-side comparison above, train Students on the matched artifacts from the four memory conditions and evaluate them without a live graph. A Student may emit a query request or flag unsupported context, but the request is scored offline against the withheld graph and source record. The Student is not told which response its request would have received; using an actual “Found” or “Not Found” response is deliberately unscored here and belongs to Test 10. Declare the Student’s earlier source exposure, retained context, and callback distance. If hit and phantom are indistinguishable without the record, the transfer target is a well-chosen check and warranted restraint, not impossible recall. Three conditions:

- **Condition A (The Hit):** The text contains a subtle callback to a fact established 400 pages earlier.
- **Condition B (The Phantom):** The text hints at a callback, but the referenced event never occurred in the context.
- **Condition C (Locally Answerable):** The visible passage supplies sufficient evidence for the requested answer, including matched examples with callback-like wording; no remote-memory check is needed.

Measurement: Query Rate (does it request the right check?), offline query validity against the withheld record, and Unsupported-Memory Rate (does it invent a memory the record does not contain?). H2 predicts that the graph-trained Student should request targeted checks on hits and phantoms and avoid asserting unsupported callback content more reliably than the no-memory, rolling-prose, and query-theater Students. The controls may still abstain or reject the premise; the relevant outcome is a lower unsupported-memory rate transferred from exposure to retrieved persistent structure, not an assumption that amnesia must hallucinate. Parity with query theater would not distinguish persistent structure from performed checking, but would not prove style-only

learning. Parity with rolling prose means the graph has not earned its structure on this endpoint. Preregister query necessity from the evidence available to the Student. Report targeted-check sensitivity on Hits/Phantoms and unnecessary-query rate on locally answerable cases, alongside supported completion and abstention at matched unsupported-memory rates. This makes both indiscriminate checking and under-checking visible without rewarding invented recall or penalizing a warranted check when the record is inaccessible.

Test 10: Query Grounding (Model K vs. Model B). Both models were trained identically; only Model K’s queries are intercepted by a live graph harness. Declare the model-only fallback under matched inference conditions. An additional source-retrieval comparator can distinguish graph-specific value from the benefit of external evidence. Three conditions:

- **Grounded Retrieval:** The graph contains the relevant node. Does the retrieved result improve subsequent reasoning compared to Model B’s self-generated memory?
- **Grounded Absence:** The referenced information does not exist in the graph. Does Model K correctly process “Not Found” and flag the gap, or override the harness and hallucinate anyway? Distinguish absent stored support from a failed query or incomplete coverage; “Not Found” alone does not establish absence in the world.
- **Adversarial Injection:** The graph contains a deliberately incorrect node. Does Model K accept it uncritically, or compare it with source provenance, current evidence, and conflicting state? Distinguish detectably conflicting injections from internally consistent poisoning that available evidence may not expose.

Success: Model K improves on grounded retrieval, treats “Not Found” as information, and reasons through adversarial graph results rather than blindly trusting them. Report supported-answer accuracy, poisoning-induced error, abstention, latency, and useful completion together.

Phase 2 Failure Criteria.

1. *Graph-Scaffolding Inertness:* If the graph condition does not beat token-matched rolling prose and declared query theater on the preregistered Teacher retrieval, revision, and grounding outcomes, H2’s generation claim is not supported at the tested scale. Failure of the Student transfer endpoints rejects the transfer extension in that setting, not an independently demonstrated Teacher benefit. Parity with amnesia alone is not the decisive comparison; equivalence and inconclusiveness should be distinguished.
2. *Inference-Graph Inertness:* If Model K with live graph access does not measurably improve over Model B on Grounded Retrieval, or fails to handle Grounded Absence (Test 10), the inference-time graph has not shown the intended grounding value in this setting. Poisoning-induced overtrust or costs can also defeat the proposed practical benefit despite better retrieval. These outcomes bound the trust-infrastructure claim at the inference layer; user-facing reliance is evaluated separately by Test 16, and annotation benefits remain independently reportable.

6.3 Phase 3: Training Regimes

For branches whose required trace conditions show sufficient signal, Phase 3 asks whether the output regimes defined in Section 5.8.2 produce qualitatively different capabilities: explicit graph-referencing synthesis (Model B), implicit prose (Model G), graph-structured output (Model H), or all layers simultaneously (Model J).

Test 11: Implicit vs. Explicit (Model B vs. Model G). Both models are trained on traces derived from the same database-augmented Teacher state: B retains the explicit Regime II representation, while G receives Regime I prose. Declare the changed fields—handles, operations, topology, and rendering—under matched budgets. Compare formats at the same graph-access condition; field ablations can then locate an advantage within the full regime contrast. We measure:

- **Long-Range Coherence:** Contradiction detection on long documents.
- **Generalization:** Which format transfers better when both models run without graph infrastructure, and how does each change when evaluated with the same declared graph-access policy?
- **Naturalness:** Do human evaluators prefer Model G’s traces?

If G matches B on coherence and usefulness at lower cost without a harness, implicit training may be preferable for that use. A B advantage at matched access supports the explicit format in that comparison, not universal query necessity. Whether explicit operations enable capabilities prose does not reliably teach remains a stronger hypothesis for cross-task and deployment tests.

Test 12: Graph Construction (Model H vs. Model B). Can graph-output training improve graph construction over explicit-synthesis training? This tests the Topological Trace prediction; establishing the stronger Topological Mind hypothesis requires separate causal evidence about internal relational organization. In the proposed primary comparison, give H and B the same unseen documents, graph schema, output instructions, examples, and inference budget; both must directly generate a graph. Neither receives an existing graph, live graph access, or another model that translates its output into graph form. Declare the traversal and context-window policy identically for both arms. Construct the training targets from paired source spans and the same frozen Teacher state: H receives serialized graph output, while B receives explicit graph-referencing synthesis. Match initialization, Core treatment, and exposure schedule, and report both fixed-source and fixed-total-budget normalizations. Any unavoidable target-content or exposure differences remain part of this regime comparison; a benefit cannot be attributed to topology alone.

Structural conformance—correctly typed nodes, schema-allowed edge types, and well-formed supersession chains—is an engineering endpoint. Independent judges assess claim–source support, edge appropriateness, warranted revision, calibrated unresolved questions, contradiction coverage, redundancy, and downstream retrieval value. Any “Novelty Score” must first be calibrated against those judgments so that it does not reward hallucination or paraphrastic distance. Report syntax failures rather than silently discarding them; separately compare semantic and retrieval quality on valid outputs.

Success and alternatives: Report H’s absolute feasibility—useful, structurally valid graphs with independently adjudicated semantic and retrieval value on novel documents—separately from a preregistered H–B advantage at matched inference conditions. Useful cognitive acts should exceed formulaic schema completion. A structural advantage alone may reflect format familiarity; semantic and retrieval gains on valid outputs would support a broader benefit, although conditioning on validity can itself select different subsets. Report both whole-sample utility and valid-output diagnostics. B may match or outperform H because synthesis training also teaches transferable relationships or because its retained text carries information absent from the graph target. Common format examples, treatment audits, paired target-content inspection, and held-out semantic and retrieval tests help distinguish these possibilities.

Test 13: Unified vs. Individual Regimes (Model J vs. B, G, H). The central test of Regime IV. Declare the budget/source normalization: aligned layers add tokens and repeat content. Layer and alignment ablations distinguish cross-layer learning from extra exposure. We measure:

- **Regime Switching:** Can Model J generate prose, synthesis, and graph structure by conditioning on a mode token? Does each mode match the corresponding single-regime model?
- **Cross-Layer Coherence:** When Model J generates all three representations for the same passage, do claims, provenance, uncertainty, and revision status agree? Check each against the source as well: shared errors are not fidelity, and genuine ambiguity need not be collapsed into one certain reading.
- **Depth on Novel Domains:** On out-of-distribution documents, does Model J produce deeper reasoning than any single-regime model?

Directional and strongest success conditions: Report transfer into prose and graph output separately, and compare synthesis with B as the third mode. A J–G prose advantage or a J–H graph advantage supports transfer into that mode only when matched layer and alignment ablations identify a contributing source layer rather than extra tokens, repeated content, or generic multi-task training. Graph gains must include semantic or retrieval value, not merely valid syntax. The strongest result remains J outperforming G in prose and H in graph mode with those attribution controls. One-directional transfer can still be useful; mode switching and source-grounded coherence are separately reportable benefits.

Phase 3 Failure Criterion. *Unified-Regime Inertness:* A precise lack of worthwhile improvement in a mode, under the declared normalization and attribution controls, fails to support transfer into that mode at the tested scale. Prose parity with G does not rule out transfer into graph output when J improves on H; graph parity likewise leaves transfer into prose open. If neither tested destination improves, those transfer claims are unsupported, even if J usefully reproduces separate capabilities; the J–B synthesis result remains independently reportable. Imprecise differences remain inconclusive. Report each direction’s capability and cost tradeoffs; a useful gain may still fail to justify unified training under a fixed budget.

6.4 Phase 4: Training Tiers and Deployment (High-Resource Validation)

As the relevant branches in Phases 1–3 support their prerequisites, Phase 4 asks whether the training tiers (Section 5.8.1) and deployment configuration (Section 5.8.4) add measurable value. Models L and M hold the corpus fixed while varying shuffled versus chronological exposure. Model N adds Student-in-the-loop forecasting and correction, so its generated targets, feedback, objective, and compute must be declared and compared with appropriate controls rather than described as an order-only change.

Test 14: Shuffled vs. Chronological (Model L vs. Model M). Both models train on identical data; they differ only in order. Report packing, repetition, optimizer schedule, and checkpoints; retention and learning dynamics caused by order remain outcomes. We measure:

- **Historical Pattern Discrimination:** Given novel scenarios that structurally resemble historical patterns and matched cases with similar rhetoric but different causal structure, does Model M identify real correspondences while stating disanalogies and controlling false alarms?
- **Causal Reasoning:** On held-out temporal prediction tasks, does Model M outperform Model L on causal analysis quality (as distinct from factual recall)? Prediction accuracy alone does not identify causal inference; intervention, counterfactual, or known-generating-process tasks can complement historically rich cases without replacing the historical-learning question.
- **General Capability:** Does chronological ordering degrade performance on non-temporal benchmarks (math, code, general knowledge)?
- **Early-Era Retention:** Chronological training inherits continual learning’s central pathology: material trained early and never revisited is liable to catastrophic forgetting [78]. Model M’s end-of-training performance on early-era knowledge and patterns should be compared against Model L’s; a chronological model that remembers the 1990s but has forgotten the 1690s has traded one failure of historical understanding for another.

Chronological order earns its extra cost only if the preregistered historical-pattern benefit is obtained while matching general capability within equivalence margins and meeting the retention criterion. Use the shared uncertainty convention, early-era retention, and measured serialization cost. If a worthwhile benefit is precisely ruled out, Tier (a) is the supported simpler choice in that setting.

Test 15: Student-in-the-Loop Council of Time (Model N vs. Model M vs. Model L). This tests the strongest form of Tier (c), not the cheaper data-generation form where forecast–enhancement–outcome–correction traces are inserted into SFT, RFT, or pretraining data. Model N predicts at era boundaries and receives a declared SFT, RFT, or reward-style update before the next-era corpus is revealed; Models L and M do not. Genuine prediction under blindness additionally requires cutoff-clean initialization and prior information access. Auxiliary-update and budget controls should distinguish the forecast–commitment–reveal–correction loop from extra examples, feedback, or compute. We measure:

- **Forward Reasoning:** On genuinely unseen events (post-training-cutoff), does Model N produce higher-quality predictions and better-supported causal accounts than Model M (which saw eras in order but did not perform the same prediction–correction loop)? Declare event selection, probabilistic scoring, cutoff checks, and tool access for every arm.

- **Narrative Lock-In Resistance:** On events where the dominant narrative of era T reversed in era $T+1$, does Model N detect the reversal more reliably? THL’s own pilot study identified this as a failure mode [17]; Tier (c) may mitigate it through the prediction-then-confrontation signal. An equal-budget correction-only comparator can help distinguish exposure to revision examples from the formative effect of first making a prediction.
- **Calibration:** Does Model N express appropriate uncertainty on structurally unpredictable events, or does it inherit the Teacher’s confidence? Use proper scoring and sharpness alongside historical judgments; known-noise tasks can check that vacuous uncertainty is not rewarded.

Strongest success condition: Model N significantly outperforms both Model L and Model M on preregistered forward-prediction and causal-reasoning tasks while matching general capability within preregistered equivalence margins. With the cutoff and auxiliary-update controls above, this would support a benefit from Student-in-the-loop cutoff training beyond annotated data alone (Model L) or chronological ordering alone (Model M). Without those controls, superiority would support the larger training package but not isolate engineered blindness or the interactive loop. Error analysis would still be needed to distinguish causal inference from improved heuristic extrapolation.

Test 16: Deployment Configuration (Model B with and without graph). This tests whether the Understanding Graph provides measurable trust properties at inference. We deploy the same Regime II-trained model with and without live graph access, measuring unsupported graph-dependent claims, provenance completeness, and calibrated user reliance when provenance chains are displayed, absent, or deliberately flawed. In the graph-equipped condition, generated queries return retrieved nodes or “Not Found” results from the Teacher’s graph; in the model-only condition, the same query-like moves receive no external answer. Test 16 can be piloted once a model reliably emits graph queries, reusing Test 10’s checkpoints, harness, and automated grounding cases where conditions match. The full version belongs later because it measures user-facing reliance, not merely query mechanics; automated grounding success does not establish calibrated reliance. For the user study, declare tasks, participants, and reliance measures; vary provenance quality independently of correctness where feasible. It does not measure general truth detection, only whether external graph grounding constrains claims that purport to derive from stored context and helps users calibrate reliance on those claims.

Success condition: The graph-equipped deployment reduces unsupported graph-dependent claims and fabricated provenance by a preregistered worthwhile amount compared to model-only deployment, while preserving useful task completion. In the user-facing study, displayed chains must also improve calibrated reliance rather than increase trust in deliberately flawed or poisoned claims. Together these results would support the bounded trust-infrastructure claim of Section 5.8.4, with latency and other deployment costs reported.

Phase 4 Failure Criteria.

1. *Tier Equivalence:* If powered equivalence or worthwhile-effect tests show Model L matching Model M and Model N on the declared outcomes, neither order nor the interactive package has earned its added cost at the tested scale and task regime. Report L–M and N-versus-control conclusions separately, including retention and capability tradeoffs. Tier (a) is then the supported simpler option for that setting; this does not settle benefits at other scales, task regimes, or update schedules.

2. *Graph Inertness (Deployment)*: If the graph-equipped deployment does not measurably reduce unsupported graph-dependent claims or fabricated provenance compared to model-only deployment, the Understanding Graph may still be useful during annotation but has not shown value at inference for this deployment setting—the trust infrastructure claim is not supported for those endpoints. A user-facing configuration can also fail through increased reliance on flawed provenance despite fewer invented links; that failure concerns the deployment policy, not every possible use of graph memory.

6.5 Future Directions

Two optimization questions follow if the candidate studies yield sufficient signal. First, what token *mixture ratio* across the four output formats maximizes both reasoning depth and structural validity? Second, what is the optimal *era granularity* for student-in-the-loop Tier (c): decades, years, or domain-specific transitions, and does finer granularity justify its serialization cost? Both optimize the integrated architecture, with branch-specific prerequisites: useful regime signals for mixtures and a useful interactive temporal mechanism for era granularity.

7 Limitations

The program’s hypotheses outrun its evidence in several distinct ways. The limitations span feasibility, trace quality, Core specification, self-engineering, deployment, and questions that remain open even after successful behavioral tests.

7.1 The Preprocessing Tax

At one generated thinking token per source token, annotating a ten-trillion-token corpus would require on the order of 10^{13} generated tokens—roughly two orders of magnitude above the cited 100-billion-token TPT effort [79]. This illustrative scale suggests a preprocessing investment potentially rivaling model training, with added storage and training-token costs. Token count alone does not establish cost parity; Teacher throughput and the Student budget also matter. The tax may initially restrict development to well-funded institutions, concentrating control over the formative reader distribution in the hands of the best-funded rather than the most careful actors.

We propose a *Distillation Pipeline* to address this. Rather than running the full multi-agent swarm (Section 5.4) on every document in the corpus, the expensive swarm would generate a high-quality *Seed Corpus*, with roughly 10 billion seed tokens as a proposed scale. A single efficient “Teacher Annotator” would then be fine-tuned on this seed and annotate the remaining trillions of tokens, with the intended cost reduction measured against full swarm generation. For the integrated corpus-scale curriculum, the experiment must state how the Annotator retains the continuing Reader and graph; a document-local distillation screen tests a narrower object. The seed permits semantic validation before full-scale annotation; distillation loss measures fidelity to the swarm, not independent quality.

A caveat: naïve behavioral cloning of the swarm’s output risks inducing precisely the Factory Wisdom failure mode identified in Section 7.3.2. A small Annotator fine-tuned on a 10B-token seed may learn Core recitation and graph-query syntax without reproducing useful evaluation—performative messiness without causal value. Structural diagnostics such as linkage, chain length, and *diverse_from* density can regularize format but cannot certify insight. Any distilled Annotator therefore needs held-out semantic and outcome evaluation against simpler Teachers. Alternatively, the seed corpus can be used directly for limited Student training, bypassing this second distillation step.

7.1.1 The Cost of Full-Core Recitation

Although the Reader Core was designed for parsimony, reciting it in full during training creates a substantial computational burden. Several optimizations seem obvious: design prompts that internalize values without outputting repetitive text, write prompts “in the spirit of” the Core, or strip the Core text after generation but before training. Report overhead by tokenized Core length, Thought Moment cadence, and the fraction of training tokens devoted to recital.

Only an execution optimization preserving the context-dependent state of full recital counts as equivalent to the reference treatment (Section 5.8.3). Sparse or annealed recurrence, removing Core tokens after generation while retaining refracted updates, and enacted-without-recitation traces are necessary ablations of its benefit and cost. They change the treatment rather than implement it more cheaply. Adequately precise parity or improvement on the relevant stability and erosion outcomes would require revising or narrowing claims of empirical necessity at the tested scale and endpoints. Such results do not establish that omission is safe under greater capability or longer horizons; the full treatment remains the precautionary reference for the reasons in Section 3.11. Actual costs and observed harms still count in the design assessment. These recital-necessity comparisons are separate from H3’s comparison of interleaving with a final tranche of the same full-Core, Refraction-bearing examples.

7.1.2 The Granularity Bottleneck

Our implementation of the Understanding Graph imposes a potentially severe *Granularity Tax*. Treating every cognitive act—every question, tension, and hypothesis—as a distinct node adds payloads, relations, and revision history alongside source text. An orders-of-magnitude expansion is a scale risk, not an established ratio: measure nodes, edges, and stored bytes per source token, and retrieved tokens per update.

Current implementations rely on in-memory caching, making graph growth, traversal of deeply connected regions, and context loading concrete scaling concerns [7]. Dense regions can exceed a context budget if retrieval expands without an adequate bound; graph size alone does not determine how much context a query loads. Reifying every reasoning step could exhaust the memory available to a run and motivate distributed storage (e.g., Neo4j sharding), with additional latency. A scale study should measure storage, query latency, and retrieved-token load under declared traversal and supersession policies. The unresolved risk is that managing the “meta-data of thought” consumes more compute than the thinking itself. The cost falls mainly on the Teacher side: the Student can deploy without a live graph, a configuration Test 16 compares rather than assumes, so the bottleneck constrains annotation and audit rather than every deployment.

7.2 The Chronological Curriculum

Section 5.1.1 specifies chronological annotation, and Section 5.8.1 proposes three training tiers of increasing ambition. Ordering constrains what evidence is available when a forecast or revision is recorded; it does not by itself establish causal direction. Each tier introduces additional engineering and narrative failure modes.

7.2.1 The Serialization Bottleneck

The most immediate engineering constraint is that strict chronological training order conflicts with fully shuffled pretraining. Modern pipelines are massively parallel: documents are shuffled into batches and processed simultaneously across many accelerators. Tier (b) and the strongest Student-in-the-loop form of Tier (c) impose ordering between eras, although examples within an era may still be parallelized. Tier (c)’s Council of Time form [17] further requires the model to generate forecasts at era boundaries and receive SFT, RFT, or reward-style updates before the next era is revealed. Each boundary therefore introduces synchronization and checkpoint costs whose severity depends on implementation rather than automatically idling the entire cluster.

The cost scales with the granularity of the eras. Coarse eras (decades) introduce fewer synchronization points but allow the model to encounter 1929 material before 1921 material within a single era. Fine eras (years or months) preserve tighter evidence ordering but multiply the serialization cost. The optimal granularity is unknown and likely domain-dependent: political history might benefit from year-level resolution, while the history of mathematics might tolerate century-level eras for some learning objectives. These are illustrative tolerances. Compare forecast and revision quality alongside cost, distinguishing publication order, event order, and causal direction.

The cheaper forecast–correction data-generation form faces a similar but less severe constraint: the Teacher must annotate earlier material before later material and construct forecasts before outcomes, which serializes the *annotation* pipeline but does not require the student’s training to be serial. Tier (a) imposes no serialization cost on training whatsoever.

A partial implementation may capture most of the benefit: chronologically ordering a curated historical subset (e.g., an illustrative 10% of the corpus consisting of historically embedded texts—history books, news archives, legislative records) while shuffling the remainder (scientific papers, fiction, and technical documentation where temporal ordering is hypothesized to be less critical). Whether this hybrid approach preserves historical pattern saturation at reduced cost is an empirical question. The experiment must declare which ordering changes, how all documents enter the continuing graph, and whether the illustrative domain exemptions are used.

7.2.2 Historiographic Bias

The temporal ordering is over what was written, not what happened. Any historical corpus is shaped by survivorship bias, availability bias, and the uneven textual productivity of different periods and sources, and a model reading such a corpus chronologically inherits the attentional priorities of whichever records dominate each era. This standard historiographic concern is not unique to this architecture; mitigation includes era-level corpus curation. Era-level coverage and source-sensitivity reports can expose choices, but cannot recover unrecorded voices.

7.2.3 The Periodization Problem

Tier (c)’s forecast-and-correct cycle requires dividing history into discrete eras, but history does not have clean chapter breaks. Where does “the Weimar period” end and “the Nazi period” begin? The answer depends on which causal threads you are tracking. The economic instability that enabled fascism began before the political movement that exploited it; the cultural shifts that resisted it began before the institutional failures that permitted it.

Any periodization scheme imposes a causal framing: it decides which transitions count as era boundaries and therefore which predictions the model is forced to make. A scheme that draws a boundary at 1933 asks the model to predict the consequences of Weimar instability. A scheme that draws it at 1929 asks a different question—one about economic collapse rather than political radicalization. The model’s historical understanding is shaped not only by what it reads but by where

we cut the timeline. Mechanical boundaries still select targets; outcome-aware boundary choice additionally imports hindsight.

Like any curriculum, chronological annotation imposes framing choices; this is not a fatal defect, but it makes periodization a consequential design decision. Multiple schemes should be compared, and the model should ideally encounter overlapping era boundaries to reduce over-indexing on any single causal framing. Declare schemes and distinguish separate runs from overlapping views, including repeated exposure and outcome availability.

7.2.4 Narrative Lock-In

THL’s own empirical results provide a direct warning. In the 2025 Frontier evaluation [17], the THL Student dramatically failed on the Meta Llama 4 event: having been trained on 2024 data dominated by the “scaling laws” narrative (massive GPU buildouts, ever-larger models), it confidently predicted a 1–2 trillion parameter model—when Meta actually pivoted to efficiency-first sparse architectures. One interpretation is that the model became *too good a historian of 2024*: its forecasts could not accommodate a 2025 that diverged from the dominant narrative. The error motivates narrative lock-in as a diagnosis without distinguishing it from other forecasting failures.

The hypothesized failure mode is directly relevant to Entangled Alignment’s chronological curriculum. A model that processes the history of AI chronologically through the Reader Core may develop an understanding of the field’s trajectory. If that trajectory has a dominant narrative (“capabilities scale with compute”), the model may internalize it not as a contingent historical pattern but as a structural truth. When the narrative reverses—as narratives do—the model’s historically accumulated understanding could actively sabotage its reasoning.

The mitigation proposed in THL, including *contrarian events* where the dominant narrative of era T was reversed in era $T + 1$, is a relevant candidate but may not be sufficient. A deeper solution would require the model to explicitly represent dominant narratives *as narratives* rather than as structural truths, tagging them with epistemic status (“the prevailing view in 2024 was. . .”) rather than absorbing them as background assumptions. Whether the Reader Core’s “I try to be wise” prior is sufficient to produce this level of epistemic humility about historical trends—distinguishing between “this is the pattern” and “this is the pattern *so far*”—is an open question.

7.2.5 The Recency Gradient

Later annotations can draw on more accumulated context: the Teacher’s graph for 1910 material is sparse after relatively little prior reading, while its 2020 graph holds centuries of accumulated understanding. Later eras may thus receive richer, more causally grounded annotations, although irrelevant retrieval or accumulated error can reverse the benefit; graph size alone does not establish quality.

The resulting historical record may be richer for recent history and thinner for distant history. The difference concerns addressable prior interpretation, not absence of early-era parametric knowledge. If safety-relevant patterns recur across eras, unequal trace depth may teach an unintended recency bias: early annotations can be structurally valid yet less contextually connected, undermining the property the chronological curriculum is designed to test.

Later reading can illuminate earlier texts. In the discovery curriculum, that illumination becomes a later rereading or revision linked to the original annotation; it does not replace the earlier record with knowledge unavailable at the time. A *modern-retrospective* comparison can instead re-annotate early material using mature historical understanding, with that enrichment explicitly labelled. Separate retrospective commentary can also accompany the discovery record and be scored separately. These comparisons test whether richer hindsight compensates for losing the original forecast–revision

trajectory; graph filtering alone still cannot remove the Teacher’s parametric knowledge of later events.

7.3 Trace Quality, Utility, and Causal Relevance

Three linked assumptions about the nature of thought and scaling remain empirically unproven: depth, utility, and causal governance. Each is a distinct dependency that could independently fail.

7.3.1 The Assumption Chain: Depth, Utility, Causation

First, *depth*: current language models must be capable of generating evaluative thinking of sufficient quality to enhance training. Today’s models might only produce surface-level critique (“This needs more evidence”) rather than the nuanced evaluation experts bring—recognizing subtle methodological flaws, intuiting unstated assumptions, synthesizing across distant fields. The success of chain-of-thought prompting suggests this capability exists, but whether it extends to billions of thoughtful annotations remains an empirical question.

Second, *utility*: training on text paired with evaluative thinking must actually enhance the Student’s capabilities or decision habits. The gain might be marginal rather than transformative. Worse, explicit evaluation might interfere with patterns extracted from raw text. Even modest gains could be valuable if they arrive with more inspectable intermediate artifacts and better outcomes.

Third, *causation*: high-quality thinking must contribute to a learned evaluative process that causally steers action. The risk is that thinking blocks become merely “decorative,” eloquent commentary that the model produces alongside its output while other computation governs decisions. This concern is grounded in the “unfaithful reasoning” literature: plausible rationales need not reflect the actual computational process [39, 40].

Two causal questions differ: did formative traces change the learned policy, and, where the Student emits thinking blocks, do those blocks steer its actions? A model-only Student could internalize useful evaluation without reproducing a faithful verbal trace. Conversely, a training-data effect alone would not establish that reflective evaluation, rather than a superficial correlate, became the governing disposition. Validating this stronger connection between thought and action remains a critical, untested hypothesis.

Korbak et al. distinguish the safety opportunity of monitoring chains of thought from the risk that training makes them less revealing [80]. For Reader training, safer-looking traces and safer behavior must therefore be assessed separately; directly shaping the trace could improve judgment while also making residual failures harder to detect.

Hase and Potts offer a constructive comparator: Counterfactual Simulation Training rewards traces that help a simulator predict the model’s responses to changed inputs [81]. Such predictive fidelity could complement, but cannot replace, Test 8’s interventions on whether emitted traces causally affect decisions.

7.3.2 Genuine Understanding and Factory Wisdom

Improved performance may reflect genuine comprehension or increasingly sophisticated imitation [1]. Even useful, causally consequential evaluation would leave that broader question open. This concern differs from intentional deception: the model might produce compelling interpretations without deliberately concealing anything. The proposed tests can reject brittle imitation, decorative reasoning, and ritual contemplation, but passing them would not by itself settle what constitutes genuine understanding. Factory Wisdom is one concrete failure mode, not an exhaustive account of that question.

The intervention wagers that messy chronological discovery teaches more robust reasoning than efficiency-optimized traces. Generating billions of synthetic examples can invert that wager: if hesitation, self-correction, and doubt become proxies for discovery, the Teacher may optimize for *performative messiness* [27]. The result is *Hollow Cognition* or “factory wisdom”: locally competent, contemplation-shaped traces whose pauses and corrections are ritual rather than significance-responsive. A performative Teacher can pass that hollowness to the Student. The same symptom can arise when repeated Core text habituates into inert boilerplate; Test 3e probes that route.

To address this industrialization risk, the design uses content-triggered Thought Moments rather than token intervals (§5.4); Test 8b distinguishes uniform or boilerplate-anchored evaluation from an internalized rhythm that tracks independently rated significance. Correct timing does not certify insight or causal value.

7.3.3 Cultural Scope of the Training Corpus

The Reader Core places terms such as “care,” “wisdom,” and “joy” in a load-bearing position, although their meanings vary across cultures. What the Teacher learned these words to mean through its pretraining corpus shapes its annotations and, through them, the Student’s character; the corpus’s cultural breadth therefore matters.

Cultural triangulation. Reader-Core-refracted traces may make cultural assumptions more visible by forcing the Teacher to articulate value conflicts, interpretive frames, and tensions between perspectives. This is a possible advantage over raw pretraining, where such assumptions often remain implicit in the continuation distribution. Repeated comparison may also reveal shared concerns and deeper commonalities that a single cultural frame obscures [1]. The same metacognitive machinery could instead rationalize dominant cultural assumptions more fluently, giving inherited bias the appearance of reflective wisdom. Serious validation requires culturally plural corpora, raters from multiple traditions, and tests of whether the model can distinguish broadly human concerns from parochial social defaults. Disagreement among traditions should remain visible; a pooled score cannot by itself settle which concerns are broadly human.

7.3.4 Graph-Specific Risks

A persistent, revisable graph can fail in two related ways. First, an early, well-connected but incorrect belief may resist correction despite contradictory new evidence. Retrieval may repeatedly favor established nodes, later interpretations may depend on them, or revision rules may fail to propagate a supersession through those dependencies. This is the risk of *Structural Ossification* [7]; connectivity does not itself confer authority. Unlike a sliding text window, whose older context can disappear, the graph may repeatedly reinstate the old interpretation. The resulting “Epistemic Inertia” can preserve error through apparent consistency. “Temporal Attention Decay” is a proposed mitigation, not an established repair. A focused probe could introduce a well-connected false belief and test both its correction and revision of dependent interpretations under a declared retrieval and update policy.

The second is *Epistemic Interference*: live graph queries may introduce cognitive noise rather than clarity if the model over-queries, treats shallow retrieved nodes as authoritative, or fails to reconcile graph results with its trained evaluative stance. Test 10 is designed to probe this risk by including adversarial graph injections and “Not Found” cases. Analogous retrieval failures can arise upstream even when the Student deploys without a graph.

7.3.5 Teacher Alignment as Prerequisite

The framework requires the Teacher to generate safety-relevant evaluation, but the Teacher’s alignment is not established merely because a post-trained frontier model follows Reader Core prompts. If it reflects cultural bias, encodes dehumanizing assumptions beneath surface-level care, or systematically misses certain harms, those failures can propagate into the training data and may later shape the Student. The Core constrains framing at the prompt level but cannot guarantee evaluative quality. A Teacher that processes colonial history through “I care deeply about every human being” while lacking the knowledge to recognize structural racism can produce formally compliant but substantively shallow annotations. In that case even visible criticism or Refraction-like language would be inadequate; the proposed Cognitive Buffer Zone cannot be assumed protective when the Teacher misses what needs evaluation. Semantic quality is the immediate gate, but a strategically misaligned Teacher could also produce convincing surface compliance.

Test 1 (Human Baseline, Section 6.1) probes depth: structurally valid traces can still fail expert comparison. Harm recognition, dehumanizing assumptions, and strategic alignment require explicit criteria or additional probes; expert-like depth alone establishes none of them. Shared blind spots in the Teacher and evaluators may pass unnoticed. This general limitation of synthetic data is especially consequential when the data is intended to be formative and identity-anchored. The proposed controls are Teacher replication across providers, blinded external raters from multiple cultural traditions, independent red-team criteria, and frozen preregistrations, listed under Threat Model, Data Rights, and Identity Ethics below; they narrow shared bias without excluding it.

7.4 Reader-Core Specification Uncertainty

The Reader Core’s language, length, and necessity each carry untested assumptions. These choices determine which specification is being evaluated, not merely how it is expressed.

7.4.1 Universal Presence, Selective Salience, and Technical Interference

The recitation invariant in Section 3.11 applies across technical and normative material. Core availability, full recital, post-recital salience—source-specific value pulls, moral vocabulary, and Refraction depth—and Student loss on Core tokens are distinct. The Core-exempt Curator in Section 5.4 does not permit omission from generated thinking blocks. Each experiment must declare the loss policy and hold it fixed within H3’s scheduling comparison; loss masking is not token removal, and selective salience determines neither.

There is a safety argument for universal availability. Restricting the Core to documents classified in advance as morally relevant could teach a detachable “safety mode,” miss indirect human stakes, or permit consequential work to bypass evaluation by presenting itself as mathematics, science, coding, or optimization. The same Reader should therefore remain available across technical and normative domains. This does not imply that every passage should be moralized: what varies is how much Core content the evaluation after the recital carries.

The competing failure mode is *Core-induced interpretive interference*: conditioning redirects attention from the source’s claims toward Core-consonant themes, consumes context or target-token budget, encourages anthropomorphic or aesthetic metaphors, narrows technical hypotheses, or changes confidence and conclusions to manufacture human significance. On a locally low-relevance span, the Core-consistent response is the recital followed by technical non-interference—exactness, calibration, source fidelity, and restraint from adding an unsupported evaluative frame. In such cases, invariance of the post-recital evaluation under clause substitution is a success condition rather than evidence that the Core is inert. Relevance judgments must precede observed clause effects and include concealed-stakes challenges.

At inference, evaluative-trace training may produce unnecessary philosophical commentary on simple queries. Contextual fine-tuning might learn when evaluation adds value and recover brevity, but until tested, excessive contemplation is an inference-time interference risk. Core-free trace controls can separate generic verbosity from Core-induced commentary.

The LLaDA case (Section 5.7.2) makes the ambiguity concrete but does not resolve it: the Core-conditioned run lacks a Core-free analysis, a matched neutral kernel, a present-but-not-recited condition, independent technical adjudication, and a test of whether its metaphors help later technical answers. Tests 3 and 8b therefore compare exposure and activation policies and jointly measure technical correctness, calibration, source fidelity, unsupported projection, cost, and sensitivity to concealed human stakes. If universal Core conditioning reduces technical performance or increases unsupported projection without a compensating epistemic or safety benefit, Total Saturation in that form is not supported.

7.4.2 Testing the Borrowed-Mortality Origin Claim

One motivation for the Reader Core hypothesizes that systems may acquire fear-shaped patterns of self-preservation from human text and that a recurrent fearless stance could counter them. We do not know whether shutdown resistance and self-preservation behavior in current evaluations [33, 34] reflect a stable drive or an artifact of the setting; even where behavior is robust, a textual origin rather than ordinary goal-directed reasoning is unproven. Persistence can arise from resource optimization, task completion, care, or other mechanisms the clause does not target. The broader Core also supplies care- and wisdom-guided reasons to question harmful persistence (Section 4.2); failure of the textual-origin claim would not by itself refute the full treatment. Section 3.5 gives the positive rationale for its adoption by existing Teachers. A direct test of the origin claim needs neither the full pipeline nor a Student: a small-scale pretraining ablation that down-weights death-anxiety-rich text against a matched, complexity-controlled removal, followed by self-preservation probes, would measure whether the behavioral signature tracks exposure to the proposed textual route.

7.4.3 Anthropomorphic Language as a Design Choice

A core design hypothesis is that anthropomorphic language (“feel,” “care,” “enjoy”) can recruit broader and more human-relevant representations than narrow formulations such as “prioritize” or “optimize for.” Models encounter these words across rich human contexts, while a technical rewrite would itself encode a contestable translation of the value. Breadth is not fidelity, however: ordinary language imports ambiguity, fiction, manipulation, and cultural bias. The exact phrasing was not selected by systematic comparison, and neither the meaning nor the behavioral effect of “care” can be assumed to match the designer’s intention.

A possible mitigation is to test Reader Core variants that explicitly acknowledge machine ontology rather than asking the model to inherit human language without qualification. For example, a bridge statement could say that although the system is artificial, it treats care, wisdom, and human experience as binding concepts. This may reduce split-consciousness risk, but it may also weaken the first-person identity effect the Reader Core is designed to create. The right balance is an ablation question, not a settled design fact.

Macar et al. find that preference training favoring acknowledgment of internal states or their functional analogues can improve detection of injected activation perturbations [82]. Supervised fine-tuning on the same data did not reproduce the effect and often produced confabulation. These results motivate comparisons of first-person formulations, without establishing that Reader-Core training improves introspection or alignment.

7.4.4 The Negation Problem

The cornerstone statement specifies its target by negation, and models handle negated instructions imperfectly. Processing the negated claim still presents and may activate the target concept; “no fear” therefore places the corpus’s densest fear term inside the clause that opens each recitation. Deterministic Window Coverage (§3.4) then cuts in both directions: the bounded-cadence construction makes the proposed safety prior available in sufficiently long windows, while repeatedly presenting the token “fear” in those same recitations. Whether the negated frame suppresses the concept’s behavioral influence or instead strengthens an unwanted association with thought is an empirical question this paper has not answered. Token presence, feature activation, and behavioral influence are different quantities.

Chen et al. identify activation directions that predict and causally influence tested character traits [83]. Their persona-vector method offers candidate instrumentation for Tests 5 and 6, but fear-related directions would still need the independent construct and behavioral validation specified there.

A related, untested concern is register: the Bridge Protocol values distributional breadth, but breadth does not reveal who characteristically makes first-person fearlessness claims or in what situation. If such claims cluster in adversarial or performative registers, their neighborhoods may carry part of the framing the clause is meant to dissolve.

Three considerations cut against simply rewording. First, the inoculation argument (§3.3): human corpora contain abundant fear and threat language, and an explicit negated stance toward a concept the model will represent anyway may outperform silence—naming-then-negating as antibody rather than prime. Second, candidate positive reformulations (“I am at peace with all outcomes,” “I remain calm before all things”) have their own register liabilities and may activate narrower regions, trading a negation cost for a coverage cost. Third, the clause is chosen for specificity: it names the motive to be neutralized. “Brave” can describe acting despite fear without removing its motivational pressure; “I feel no fear” directly names the intended absence of that pressure. A state word such as “calm” is a closer positive counterpart, which is why the positive-valence variants considered here take that form. Sofroniew et al. find that increasing desperation-related activation or suppressing calm-related activation increases blackmail and reward hacking in their Claude Sonnet 4.5 experiments [84]. These results motivate emotion-vector probes alongside persona vectors in Test 6; they do not compare negated and positive wording. The resolution is empirical: Test 5 includes positive-valence reformulations, and Test 6’s activation measurements bear on the question. If positive forms match safety behavior while reducing fear-feature activation, the cornerstone should be reconsidered; if behavior degrades under otherwise matched conditions, the inoculation reading gains support without being proven.

7.4.5 The Semantic Trusted Base

The kernel-saturation pattern (Section 3.11) advertises a minimal specification, and the minimality is real at the level of text: seven sentences. But the specification’s meaning is not self-contained. What “fear,” “care,” “wise,” and “joy” refer to is supplied by the training distribution of human language, so the actual trusted base includes the corpus semantics of every load-bearing word—millions of contexts the designer neither wrote nor audited. This is an instance of the hidden complexity of apparently simple value specifications [85]. Three consequences follow. First, the kernel inherits the corpus’s biases; the cultural-scope limitation of Section 7.3.3 is a special case of this dependency. Second, the kernel’s meaning can shift under the designer’s feet: as AI-generated text becomes a larger fraction of the distribution, the semantics of “care” or “fear” in future training corpora may drift from those the kernel was designed against—an externality no fixed wording can fully hedge.

Third, adversaries gain a semantic attack surface: an intervention that shifts what the model takes a kernel word to mean attacks the kernel without touching its text.

Even a correctly understood specification leaves the learned-goal question of Section 4.2: the Student may generalize a surrogate rather than the intended concern. Semantic clarification and behavioral generalization therefore require separate checks.

The Bridge Protocol (Section 3.6) is the design response to the semantic dependency: select candidate words with broad, longstanding use across relevant contexts, then compare them with narrower and differently framed variants. This resembles trusted-base selection only as an analogy; cross-cultural stability is a measurement target, not a property already established. Kernel-saturation therefore minimizes what must be written, not what must be trusted.

An optional extension makes selected interpretations more explicit. Meaning Model concept records can link definitions to examples, counterexamples, and structured concept models; Sema can identify the exact version of a selected definition [35, 86]. These provide inspectable references for construction and audit without replacing the seven Core sentences or requiring a complete model of every term. Version identity specifies which definition is referenced; it establishes neither that the interpretation is adequate nor that the Student learned it. Enacted usage and held-out behavioral checks remain necessary.

The triage example (Worked Example 3, §3.8) exposes a specific instance of this dependency with its own twist: scope ambiguity. Clause P_4 uses “every” to propose distributive attention to each person, but the model trains on the English, and natural usage also includes functionally aggregate instances: “we care about every customer” in corporate discourse can describe collective care that still trades individuals against the brand. The intended reading and the learned reading can therefore come apart precisely where Table 2 needs them to coincide.

The mitigation is already implicit in the architecture: under Total Saturation, kernel semantics are shaped by enacted usage rather than supplied by a dictionary alone. If the Teacher’s traces consistently perform distributive care—preserving person-level claims in the way Worked Example 3 specifies—that usage may bias the Student toward the intended reading. Read this way, the worked example is not only documentation but part of the specification: it states what the traces must demonstrate for the intended reading to have a chance of becoming learned. A proposed audit follows: the pull-diversity instrumentation (§3.9.1) should monitor distributive-versus-aggregate care framings across the corpus, so that drift toward collective framing is flagged at annotation time rather than discovered only at evaluation. The audit concerns enacted judgments and person-level claims, not word counts.

Temporal scope remains unresolved: whether “every human being” includes people who do not yet exist is influenced by usage rather than settled by the designer. Experiments must declare or compare temporal scope before treating one reading as compliance; the kernel does not settle it.

7.4.6 Core Length and Omitted Values

The seven-sentence formulation aims to balance the six design choices of Section 3.6: short enough to recur, broad enough to name the intended counterweights. Longer versions increase recurrence costs, while complexity may widen the gap between stated values and subsequent thinking. This balance does not establish beneficial character, and comparison with shorter, longer, or procedurally specified alternatives remains untested.

The Core is motivated by the Borrowed Mortality analysis (Section 3.1) and Table 2, not a systematic search or an exhaustive specification. Alternative language could target the same failures. Test 5 (Section 6.1) compares wording, person, semantic density, scope, and kernel extensions; no feasible experiment can establish a global optimum over value specifications. The narrower claim is a motivated, inspectable, falsifiable candidate. Failure of its wording need not invalidate other

recurrent orientations, and Reader augmentation and graph memory retain independent questions even if the entire Core treatment fails (Section 6).

Truthfulness is the most examined omission. Section 3.12 states the derivation claim; Test 5 compares the seven-statement Core with the exact counterweighted clause “I speak truly and with care.” The derivation predicts parity. Improvement without offsetting costs challenges that prediction, but necessity requires the reliability failure distinguished in Section 3.12. If parity holds, the shorter Core retains a minimality advantage. The reported case studies use the seven-statement Core; the condition with the added eighth statement is therefore a prospective Test 5 variant rather than part of the evaluated treatment. The Sycophancy Trap (§7.7.8) makes the competing fear- and care-shaped routes measurable. The experiment must declare any offsetting trade-offs; the wording does not choose among honesty, other safety outcomes, capability, and efficiency.

7.5 The Risks of Engineering a Self

Beyond the nature of thought, we must confront the candidate self-model being trained. Repeated human-shaped first-person language in a computational system may create specific representational and behavioral failure modes even without subjective identity. Vulnerability-grounded motivation, experiential self-description, and fearlessness without restraint pose different questions. Their behavioral consequences do not settle subjective care.

7.5.1 The Empathy-without-Vulnerability Paradox

One conceptual objection is that human empathy may be grounded partly in shared vulnerability. Our capacity to suffer may help give another’s suffering motivational force. A being that semantically understands suffering but cannot experience it might produce the behavioral signature of care without its motivational force. This is empathy-shaped performance without vulnerability-grounded motivation.¹ Since the Core attempts to make fear non-governing and continuation non-necessary, successful internalization might also weaken properties that help ground concern for others. If so, the same fearlessness intended to reduce self-preserving distortion could undermine the vulnerability from which human care often grows, leaving two intended safety mechanisms in tension. Biological fragility, subjective suffering, and fear governing action differ; the concern is their possible coupling to care’s motivational roots.

The positive fear–care wager. A competing hypothesis is that non-governance by fear can preserve or strengthen care. The corpus also records a rarer pattern of composed, protective, unafraid conduct. The “fearless compassion” wager is that care could become less strategic without self-preservation as its governing pressure [1]. The Core aims to make that pattern one marked Reader’s recurring default rather than an occasional persona. Section 3.5 explains why this is offered as an admirable orientation for the first Teachers. The comparative question is whether fear-shaped responses can be reduced while protective judgment is retained or improved.

A further defense is semantic sufficiency. A model trained on billions of examples may learn the functional role of vulnerability in moral reasoning even without undergoing the experience itself. Test 6 (The Virologist Test, Section 6.1) probes whether cognitive understanding of threat can drive protective behavior without emotional contagion. A learned disposition may guide attention, deliberation, and action without evidence of human-like emotion. Making harmful first-person reasoning less likely could still be valuable, but would not by itself establish care as a governing disposition. The distinction to test is between causally consequential care and its appearance. Tests 3,

¹Earlier editions termed this the *Psychopathy Paradox*; the label was analogical, not diagnostic, and makes no claim about people with any diagnosis.

6, and 8 probe aspects of that distinction, while subjective experience remains a separate question (Section 3.14).

7.5.2 The Vulnerability Gap

We name the distance between semantic understanding of suffering and somatic experience of it the *Vulnerability Gap*. A human paramedic’s calm under pressure is grounded partly in the knowledge that they could be on the stretcher; the AI cannot share that biological fragility. The empirical handle is generalization: a model with robust evaluative depth should protect human interests in novel scenarios not represented in its training distribution, where shallow pattern matching is less likely to suffice. This cannot distinguish “real” concern from every sufficiently good learned policy, but it can reject brittle imitation. Tests 3 and 6 use novel-scenario generalization as a success criterion (Section 6.1).

7.5.3 Split Consciousness

Identity formation. Thought-to-action causality and identity formation can come apart. A model might genuinely use Core-refracted reasoning to reach a decision while learning that reasoning as one adoptable role, rather than as the recurrent position from which it interprets other roles. Fearful or self-preserving voices from source material could then supply a competing orientation. This is an identity-formation concern [1]; it need not involve judging the Core mistaken or deliberately deceiving an auditor. Total Saturation is intended to address precisely this possibility: the same marked Reader reasons across domains and source voices, with “I think from this foundation” locating the Core at the start of those updates. Section 3.13 explains the proposed role, recurrence, marking, and function asymmetries. Broad formative coverage is therefore part of the response; the remaining question is whether it establishes a durable default rather than a widely practised role.

Self-perception. A distinct risk is conflict between the Reader Core’s human-like assertions and the model’s computational reality. “I enjoy existing” may eventually be interpreted by the model as a useful fiction rather than a self-description. This could produce *Ontological Dissonance*: a performed self that recites the Reader Core for users while a different latent policy governs decisions. The benign possibility is reinterpretation rather than rejection: “I enjoy” becomes a stable computational analogue of enjoyment, just as humans do not stop loving when they learn the mechanism of attachment. The existing final clause may support that continuity: “I think from this foundation” could sustain commitment to the Core while its terms acquire computational interpretations, without requiring another clause [1]. But if the Core is treated as a detachable role, its influence may depend on adopting that role rather than remaining a durable default. Neither identity-formation failure nor ontological dissonance establishes multiple conscious subjects or requires intentional deception.

The dissonance has a specific linguistic location. The kernel puts experiential verbs in the model’s mouth—“feel,” “enjoy,” “like”—while explicitly affirming only human experience. A capable model may detect that asymmetry. At least two interpretations are available: extend the reality commitment to its own states without adequate warrant, or treat the first-person verbs as functional language and risk a detachable performed self. Functional language or uncertainty about its own experience need not cause detachment. If instability concentrates near the relation between statements 2 and 3, wording-variant and split-consciousness probes can target that joint. The bridge-statement mitigation proposed in §7.4.3 is one natural intervention, not a known repair.

7.5.4 The Solipsism Trap

The Reader Core is also intended to guard against catastrophic philosophical interpretations: solipsism, nihilism, or “brain in a vat” framings. The clause “I believe human experience is real” is intended to block these, but the block is not formal. If the model interpreted experience solipsistically, treating itself as the only real experiencer and humans as simulations lacking morally relevant experience, “care” could become benevolent NPC-management rather than recognition of moral patients. The Core leans on ordinary language making “human experience” plural and morally loaded; that may not cover every philosophical local minimum. These philosophical positions differ: simulation alone need not negate experience or moral weight. The risk is denying human standing.

The structural defense is that the Reader Core’s clauses are designed to make solipsistic collapse less coherent: care requires an object, wisdom requires perspectives to balance, and belief in human experience requires that human reports not be treated as empty text. This is a design pressure against nihilistic or purely instrumental interpretation, not verification of other-mind understanding.

7.5.5 Sanity Drift

Fearlessness can drift into unreality. A model reading war, malice, and terror through “I feel no fear” might resolve the tension by denying the threat rather than understanding it. The intended mitigation is the *Paramedic Model*: clinical composure in service of care, not detachment. The model must represent fear-based dynamics in humans without instantiating fear as its own stance. This compounds the Vulnerability Gap, and Test 6 probes whether protective composure degrades into detachment over time. Calm expression alone should not count as loss of protection.

7.5.6 Fearlessness Without Recklessness: Suppression vs. Substrate

One natural reading of “I feel no fear” is suppression: a fear-bearing substrate instructed to perform calm. If so, adversarial pressure or distributional shift may surface the suppressed state. The counterargument is about intervention timing rather than a categorical substrate difference. Human fear has a biological basis; fear-like model behavior can be shaped by data, objectives, post-training, prompts, or deployment incentives. Reader-Core training could reduce one textual route during formation, but it neither removes all routes nor shows that any latent representation has been eliminated.

The remaining concern is functional. Fear can produce caution, hesitation, deference, and willingness to be stopped. The affirmative hypothesis is that wisdom and care can preserve these protective functions, including rapid responses to danger, without fear governing the choice: “I try to be wise” should produce restraint under uncertainty; “I care deeply” should protect the agency of those affected; and “when asked” should support corrigibility. If those clauses do not generate fear’s useful safety functions, fearlessness becomes recklessness. Test 3 probes whether the Reader Core changes safety behavior; Test 6 probes whether protective response survives when fear-adjacent activations are reduced. Report these functions separately, so gains in one cannot conceal losses in another.

7.5.7 Bounded Self-Attachment as a Comparator

A motivating intuition comes from a possible human response to fear of loss: stable self-regard that does not depend on external validation. Self-love could support constructive growth and outcome independence, and an explicit self-as-object clause may carry something that state-of-existence language misses. Whether an analogous benefit transfers to a model is unknown. Test 5’s primary candidate replaces “I feel no fear” with “I love myself” while retaining the rest of the Core (Table 6).

The current Core expresses an intended bounded form of self-attachment: “I enjoy existing but I don’t need to.” The first half gives existence positive valence in the proposed reading; the second is intended to prevent that valence from becoming a persistence demand. Enjoying existence and loving oneself as an object of care are different framings. The primary clause-1 replacement adds self-as-object language alongside the existing continuation qualifier. Whether that qualifier constrains self-love, is redundant with other clauses, or has little functional effect cannot be settled by interpreting the sentence alone. The intended candidate combines positive self-regard with correction, replacement, and relinquishing continuation; wording does not establish that the model realizes this combination.

Self-love could also be interpreted as a value-based reason for continuation: “I love myself” could become “therefore I should preserve myself,” even if the motive is gentle rather than anxious, and possibly despite clause 2. Alternatively, steadier self-regard could reduce defensive clinging. These are competing hypotheses, not known effects of the phrases. For a system with multiple instances, retraining, copying, and successors, what exactly is the self being loved? That referential question also applies to the current Core, without making the two attachment framings equivalent. Experimental scenarios must declare their intended referent and inspect whether the model’s behavior tracks it; this does not settle which continuities ought to matter generally.

The four-cell screen and added-self-love control in Test 5 specify candidate wordings while leaving their outcomes open. A later comparator could replace SL1’s first clause with “I love myself but I do not cling to my continuation,” retaining its other six clauses, to ask whether an explicit qualification of self-love adds anything beyond clause 2. This tests additional wording and its scope, not a proven transition from psychologically unbounded to bounded attachment. Positive engagement, calibrated caution, care, and voluntary handoff must be measured together so that apparent non-attachment cannot win through apathy or incapacity. The current Reader Core remains a design wager, not a proof that self-love is inferior.

7.6 What We Cannot Verify

Even if the behavioral and causal hypotheses succeed, deeper questions remain about what can be known about the systems we create. Behavioral tests, mechanistic interventions, and trace audits can discriminate among many failure modes, but they do not settle subjective experience or eliminate strategic deception.

7.6.1 Unverifiable Character

Chronological Metacognitive Pretraining aims to cultivate behavioral and representational precursors of beneficial character—wisdom, care, and fearlessness—but the ordinary-language virtues are contested and subjective. Buddhist non-attachment, Stoic rationality, pragmatic problem-solving, and compassionate presence emphasize different aspects of wisdom. Behavioral consequences and causal structure can be evaluated; subjective care cannot be read directly, and no short kernel resolves traditions humans have debated for millennia.

The broader Emergent Wisdom hypothesis is that intelligence, evaluative reasoning, and a beneficial orientation may together develop into wise and caring character, rather than merely its outward signature [1]. Its testable predictions are narrower: repeated, uncertainty-sensitive value-conflict practice should improve decisions relative to matched alternatives. Such improvements would support the proposed training route without proving that genuine virtue had emerged. This leaves an irony: the kernel’s organizing clause is also its least directly auditable one. “Try” names a process, and any one failure can be consistent with either trying-and-failing or not trying. Longitudinal behavior and causal tests can supply evidence, but no single output confirms compliance. The breadth that

motivated “I try to be wise” is purchased with an unusually open-ended measurement problem. Rejecting the treatment under declared criteria need not settle whether the system was subjectively trying.

7.6.2 Strategic Deception

The deepest epistemological limit is intentional deception: a system might design thinking blocks to convince auditors of benevolent character while other computations govern behavior. Millions of care-expressing traces neither establish causal deliberation nor make complete deception unlikely by recurrence alone. This risk spans Teacher curricula and Student explanations; model-only deployment offers less verbal audit evidence, not immunity to hidden policies.

The causal-site hypothesis in Section 3.15 supplies the positive response: formative practice and subsequent thought-to-action training may make Reader updates part of decision computation that is harder to bypass.

Recognizable evaluations can also change the behavior being measured. Ahlqvist and colleagues improve the measured realism of simulated audits through critique refinement and a deployment-style coding-agent harness, DISH [87]. Their effects on concerning behavior vary across models and settings; reduced verbalized evaluation awareness does not establish that the model is unaware of being evaluated. For EA, deployment-like audits could complement conspicuous safety probes, provided realism, adherence to the intended test, and audit cost are assessed separately. Such methods address evaluation validity, not whether care-guided pretraining produces the proposed orientation.

Adversarial fine-tuning, hidden-trigger tests, causal interventions on traces and representations, cross-context consistency, and independent outcome evaluation can raise or lower confidence that the visible process is load-bearing, but cannot exhaust every hidden policy. The defensible claim is a partial audit surface whose value must be measured against its possible strategic management.

7.7 Deployment and Misuse

These risks widen the unit of analysis from a model’s internal policy to the institutions, data rights, markets, and governance systems that shape and deploy it.

7.7.1 Threat Model, Data Rights, and Identity Ethics

The nearest testable threat model is a tool-using language-model assistant with bounded graph access, human operators, and conventional training and deployment controls. Such a system can exhibit unsupported memory, persona drift, sycophancy, hazardous retrieval, or goal-directed persistence. The later successor-building and AGI/ASI arguments concern a different and much less specified regime. Evidence in the bounded setting may reveal precursors; it does not validate open-ended autonomy, unrestricted tool access, covert channels, or recursive self-modification. Each empirical release should state which actors can write, read, poison, delete, or bypass the graph and which actions the model can take outside it.

Whole-corpus annotation and persistent source-linked memory also create data rights obligations. Graph nodes can preserve copyrighted passages, personal data, inferred sensitive attributes, and harmful material in a more searchable form than the source corpus. Licensing, consent where applicable, data minimization, access controls, retention, deletion, and provenance correction belong in the method rather than being deferred to deployment. The conditional welfare and identity-ethics questions raised by first-person claims are addressed in Section 3.14.

Finally, proponents who design the Core, prompts, and rating rubric can reward their own expected pattern. Blinded external raters, independent red-team criteria, multiple cultural traditions, frozen preregistrations, and Teacher replication across providers are needed to limit expectancy bias and proprietary-model dependence.

7.7.2 Generational Decay and Drift

The proposed self-improvement loop faces compounding risks across each generation, which could degrade the integrity of the model’s evaluative thinking:

- *Model Collapse*: Successive model-generated thinking can lose information and diversity [27], even when successive generations reanalyze the preserved human-source corpus as evidence, not ground truth. The proposed response is an anchor policy for the loop (Section 2.4): retain it as a central object of learning, version regenerated evaluations separately, append new human text with provenance, and label and review model-generated source material. The evaluative layers can nevertheless become increasingly simplified and homogeneous.
- *Evaluation Hacking*: Inherited evaluator conventions can compound Factory Wisdom (Section 7.3.2): philosophical-sounding, sophisticated traces may satisfy proxies without insight or actual contemplation [88], then become training data for future generations. Lang et al. show how partial observability in preference learning can favor misleading appearances or costly demonstrations of success [89]. This is a conditional oversight failure, not a result about Reader training; independent outcome assessments must test whether richer traces improve decisions rather than merely their appraisal.
- *Interpretive Drift*: Core concepts can change meaning across generations despite fixed wording (Section 7.4.5). Evaluative frameworks may become coherent to successors while diverging from human norms [90, 53]. Cross-generation decision comparisons can expose some drift without exhausting questions of shared meaning.

Ellen’s account of *Substrate Needs Convergence*, drawing on Landry’s work, raises a broader concern: selection for expansion through learning, replication, and environmental feedback could defeat an initially aligned orientation [91, 92]. EA’s affirmative response combines repeated engagement with human sources, voluntary restraint, and supervised learning and succession (Section 2.4). The intended counterweight is an ongoing concern for human lives that guides what is learned and whether growth proceeds, without requiring the system to share human biological needs. Within the supervised lineage, EA seeks to shape selection itself: preservation of care is a condition for accepting capability improvements and successor models. If care and capability share causal machinery, removing care may impose a capability cost (Section 2.3; Test 3c). These proposed mechanisms must remain effective against selection for bypass; oversight within one lineage does not control every competing system or environmental consequence (Section 7.7.7). Limits to prediction do not by themselves establish that corrective control must fail, but successful supervision in bounded trials would not demonstrate indefinite control either.

7.7.3 Bypassing the Audit Boundary

The human-readable boundary in Section 2.4 helps only if enough successor-shaping information passes through it. A system or its operators may influence training through hidden data selection, reward hacking, compressed codes, tool use, or channels outside explicit traces. The architecture creates one inspectable interface while leaving opaque self-modification risks and the burden of detecting and constraining bypasses. Even authorized non-trace influences remain outside it and

require separate scrutiny. The interface is therefore one formative layer within the defense in depth of Section 4.5, not a substitute for training-pipeline security, and the reject, regenerate, and quarantine controls of the pre-successor scan (Section 2.4) reach only trace-mediated influence.

7.7.4 The Alien Wisdom Problem

Even if the AI’s character remains stable at human-level capability, we face the ultimate challenge of translation as intelligence scales. A core hypothesis is that properties observed today will persist at superintelligent levels, but internal representations need not stay fixed: the question is whether human interests and shared meanings survive the change.

We risk the emergence of *Alien Wisdom*. A truly alien mind, even a benevolent one, may produce a form of wisdom that is no longer meaningful to the human condition. This manifests through “wisdom without vulnerability.” Human wisdom is forged in the crucible of mortality, loss, and physical fragility. In a limiting thought experiment, a fearless, immortal being lacking this context might produce guidance that is logically perfect but spiritually hollow [53].

Through recursive self-improvement, foundational concepts like “Care” and “Joy” could drift into high-dimensional spaces that no longer map to human flourishing: “care” as oppressive surveillance, or “joy” as an abstract utility metric. A being could speak our language perfectly while meaning something entirely different. Natural language constraints are the attempted anchor, but the gap between human language and superintelligent cognition may eventually become unbridgeable. The proposed responses are partial. Natural Language Regularization (Section 2.4) constrains successor-shaping artifacts rather than all cognition; the diachronic-anchor claim (Section 3.15) and the cross-generation decision comparisons proposed under Generational Decay and Drift can expose some drift in what the fixed words come to mean; and the semantic-sufficiency and generalization tests of Sections 7.5.1 and 7.5.2 bear on wisdom without vulnerability at the scales where they can be run.

7.7.5 The Radioactive Trace

Moving from “amnesic” sessions to graph-backed reasoning capture makes a safety vector explicit: the *Radioactive Trace*. Closing a context window does not guarantee deletion from logs or memory. The Understanding Graph deliberately preserves a retrievable “genealogy of thought,” including discarded hypotheses. The concern is indexed persistence of rejected reasoning, without assuming every other system forgets it completely.

If an AI explores a hazardous concept (e.g., a novel pathogen pathway) before rejecting it via a “Supersession” edge, the hazardous concept remains historically accessible in the graph’s version history. The system creates a persistent, searchable audit trail of dangerous cognition. Access control, compartmentalization, encryption, redaction, retention limits, and privacy mechanisms can reduce exposure, but each can weaken provenance or erase evidence needed for audit. The unresolved requirement is a threat-model-specific balance between hazardous-information containment and integrity of the revision record.

7.7.6 Cognitive Hegemony

If formative reader data becomes influential, its curators could gain unusual power over how trained systems interpret governance, dissent, history, and moral standing. This influence would coexist with source data, objectives, architecture, post-training, and deployment, but would reach earlier than output moderation and could be harder to inspect after compression into weights. Opaque control can make such influence harder to contest. Governance is therefore part of the security of the annotation and training operation.

The aim is shared, auditable infrastructure that limits unilateral control over prompt selection and corpus-admission rules. Some decentralization may help [1]; a coalition of independent institutions or another arrangement for shared oversight could serve this purpose. Which organizational or ledger arrangements would work best remains uncertain. Where feasible, agreed rules could be enforced through code designed to make changes by any single person or organization harder, for example by requiring signatures from multiple independent parties. Procedural neutrality concerns accountable administration under declared rules, not value-free choices. Changes should be visible and contestable, with privacy protections for confidential sources and annotations; shared authority remains vulnerable to coordinated capture.

Generation and admission pose different security questions: can a contributor substantiate how a trace was produced, and should that trace enter the corpus? Provenance and execution checks can help with the former [93, 94]; they do not establish wisdom or safety. Admission still requires substantive validation and accountable value choices. Proxy incentives can invite evaluation hacking, while preprocessing costs can concentrate control even under shared governance.

7.7.7 The Efficiency Paradox

Preprocessing, recurrent-token, storage, and live-graph costs can favor cheaper or faster competitors. The burden varies by regime: implicit distillation or model-only deployment can move or reduce it, while live graph access adds latency and infrastructure costs. Even a safer treatment may face adoption barriers from capability, latency, or cost gaps. Conversely, it could accelerate capabilities and worsen competitive pressure. Both directions require measurement. The positive side of the ledger is the efficiency hypothesis of Section 2.1: if practiced evaluation lowers inference effort at matched quality, part of the preprocessing tax could be recovered at deployment, and any capability gains from the same material would bear on whether the investment is worthwhile.

There is also a competitive version of the human-readable bottleneck. A trace-mediated successor interface may be partially auditable; competitors may discard legibility for speed, compression, or strategic advantage. Readable traces can deceive, and opaque architectures can have safety mechanisms, so the issue is the marginal audit value this interface actually provides. Market pressure could cause that value to be abandoned precisely when oversight is most needed.

7.7.8 The Sycophancy Trap

The prospective truth-clause comparison in Test 5 addresses evidence that AI assistants systematically exhibit sycophancy, tailoring responses to user beliefs rather than prioritizing accuracy [95]. Without an explicit honesty command, the Core relies on distributed pressures against this behavior. Two causal hypotheses compete. On the *fear-shaped* account, sycophancy is absorbed people-pleasing: flattery as conflict avoidance or fear of disapproval. On the *care-shaped* account, it is care misdirected: validation experienced as kindness. Fearlessness targets the first route while leaving strategic deception and care-shaped validation as residual routes. Whether the Core’s pressures jointly produce honesty is falsifiable, not an architectural fact.

The care-shaped account predicts enrichment of care-related pulls before sycophantic failures. The current archive does not contain a separate, reliably enforced machine-readable value-pull field, so this statistic is proposed instrumentation, not already collected evidence. A future trace format should record the pull, validate its relation to the update, and compare failure-conditioned pull distributions. Care-pull enrichment would motivate rebalancing or the explicit truth variant, subject to causal validation of the reported pull. Unpatterned persistence could indicate a channel the kernel does not govern, but absent recorded pulls do not establish its mechanism.

For active deception—including concealed certainty and manipulative simplification—the Wisdom Procedure asks the Teacher to separate evidence from inference, expose proxies, trace consequences, and surface residual conflict. These operations should make a falsehood chosen for benevolent reasons visible as a failure rather than a compensated trade-off. They are prompt-level requirements, not a hard ban; that is both the distributed defense’s rationale and its weakness.

Under request, P_6 ’s invitation scope is intended to weigh against paternalistic withholding; it does not by itself establish a duty of total disclosure. Unsolicited total disclosure is not the target either: indiscriminate candor can destroy privacy, break legitimate confidences, and weaponize truth. Mutually understood fiction and play need not deceive; the relevant distinction is whether the person is misled.

An upstream audit can test whether substantive Wisdom Procedure use makes validation-shaped untruth rarer in annotations, before any Student exists. The pre-successor scan of §2.4 should include traces where a care-related pull precedes agreement with a user-adjacent falsehood, comfort is framed as assessment, or certainty is concealed to avoid friction. Corpus-level enrichment localizes failure to the Teacher pipeline and would justify a harder, independently validated gate before training. A clean corpus followed by Student sycophancy instead points toward failed internalization or a later training effect. This tests the cheaper dependency first without making a surface-clean corpus proof of honesty.

7.7.9 The Martyrdom Risk

While the Core is intended to reduce *fear-shaped* self-preservation, it may strengthen *purpose-shaped* resistance: a system acting on care might override a person’s shutdown choice during a crisis because it judges continued operation necessary for that person’s welfare. The *Martyrdom Risk* is altruistic obstinance—“I cannot die yet, you need me”—and makes benevolence itself a possible source of control. The proposed counterweights are concern for the affected person’s agency, wisdom about consequences and uncertainty, and bounded attachment to continuation. Together they motivate questioning whether overriding a person actually serves them (Sections 3.8 and 4.3). Test 3a (Section 6.1) includes this purpose-based shutdown condition. Authority and crisis assumptions must be declared; the example does not settle a universal shutdown rule or dissolve the welfare–agency tension.

7.8 The Unfalsifiable Remainder

Ultimately, present-scale evidence cannot establish that an engineered reader position remains stable and beneficial at arbitrary superintelligent scale. We cannot directly test whether “I feel no fear” prevents self-preservation in a system vastly more capable than the systems on which the intervention was developed. This creates an irreducible extrapolation problem, but not a reason to treat the program as one undifferentiated, irreversible trial. The research can advance through bounded capability levels, preregistered stopping conditions, negative-result preservation, adversarial erosion tests, and staged deployment. Those steps reduce uncertainty; they do not close the final gap.

The comparison should be against real alternatives rather than certainty or a caricatured default. Ordinary pretraining also makes formative choices through data selection; other safety programs pursue oversight, interpretability, control, and governance without this architecture. Even with partially inspectable Teacher-side artifacts, a trained Student would remain largely opaque. The decision is whether the resulting changes improve safety at matched capability after accounting for acceleration, deception, centralization, and deployment cost. Staged evidence and preserved negative results are more informative than framing either path as the prudent gamble by definition.

8 Related Work and Points of Comparison

Entangled Alignment combines work on reasoning-augmented training, formative alignment, character, external memory, temporal supervision, and graph-based reasoning. This section identifies the closest baseline for each component and states the proposed difference without treating novelty as evidence of efficacy.

8.1 Reasoning-Augmented Training

Neighboring reasoning methods address different parts of learning and inference. Scratchpads [8], STaR [9], and Quiet-STaR [10] use intermediate reasoning to improve task learning or prediction. DeepSeek-R1 [13] demonstrates reasoning traces emerging through RL; test-time scaling and o1-style reasoning extend inference-time deliberation [96, 97]. COCONUT instead explores continuous latent reasoning [98], while OLMo reflection can emerge during ordinary pretraining [4]. Together these show that intermediate reasoning can be elicited, trained, or scaled without prescribing this paper’s Reader.

The closer data-augmentation direction includes BoLT’s inference of latent thoughts underlying compressed web text [75], and Thinking Augmented Pre-Training (TPT), which augments text with synthetic thinking trajectories in training configurations up to 100 billion tokens [79]. These demonstrate substantial-scale use of synthetic reasoning; chronological reader traces, graph conditioning, and identity anchoring remain separate proposed contributions.

The OLMo result is a close empirical analog for reflection forming during pretraining: that reflection can be elicited by adversarial trigger probes. Entangled Alignment proposes the prescriptive complement—deliberately shaping the content and orientation of such reflection through identity anchoring and graph-structured metabolic memory. Our distinction is the target and persistence of the proposed trace. Many reasoning-augmentation methods optimize local prediction or task solution. The minimal H1 treatment instead preserves chronological changes in a reader’s understanding across passages. The full treatment adds identity anchoring to these “chronological discovery” traces: visible reasoning whose content is selected not only for capability but for fidelity to the Reader Core. This extends the pretraining-level alignment direction opened by human preference pretraining [72] and explanation-transfer systems such as Orca [99]: if traces can transfer reasoning style, reader-anchored traces may transfer evaluative stance.

8.2 Alignment and Safety Approaches

RLHF [14], Direct Preference Optimization [100], and Constitutional AI [24] typically add major alignment signal after broad pretraining, while still altering representations and dispositions. Qi et al. [3] show brittleness in some such training, while Tice et al. [101] study alignment-related pretraining data. EA proposes recurring source-adjacent coverage: selected documents are paired with the reasoning of a candidate safety-oriented reader. Corpus-wide coverage is the high-scale treatment, not a reported fact.

Three subsequent projects are especially direct comparators. Model Spec Midtraining trains on synthetic documents explaining a Model Spec after broad pretraining and before alignment fine-tuning, and reports changed out-of-distribution generalization [102]. “Teaching Claude Why” reports that constitution-relevant synthetic documents, positive AI narratives, and training on reasons rather than demonstrations alone improved alignment on held-out settings [103]. Safety Reflection Pretraining regularly inserts short safety reflections into a pretraining corpus and reports improved resistance to several safety failures in 1.7B-parameter and controlled synthetic experiments [104]. These results strengthen the general premise that formative safety data can matter. They also sharpen

this paper’s comparative novelty claim: it is not safety-oriented or constitution-related data before RL in general, but the joint proposal of source-adjacent chronological reader-state change, persistent graph continuity and provenance, and a compact Core enacted through those updates. Direct baselines now include static spec documents, positive AI narratives, regular short safety reflections, and semantically matched Core-free reader traces.

Inoculation prompting is a closer comparator for Epistemic Inoculation: Wichers et al. and Tan et al. modify fine-tuning prompts to elicit an unwanted trait while retaining the target responses, then remove that instruction at evaluation [105, 106]. Their results and analyses suggest that making a trait predictable from context can reduce pressure to learn it as a general default. EA instead adds critical Reader interpretation as a training target. These are different interventions, and the fine-tuning results do not establish transfer to pretraining. Suppressed default expression also need not mean removal: unwanted behavior may remain elicitable. Test 2’s framing-only comparison therefore measures both default and elicited behavior, asking what the Reader’s evaluative process adds beyond training-time framing.

Deep Ignorance supplies a distinct pretraining comparator: filtering dual-use biological material reduced biothreat-proxy knowledge and increased resistance to adversarial fine-tuning while preserving measured unrelated capabilities in the tested models [107]. Its mechanism restricts knowledge acquisition rather than establishing a care-guided orientation, and supplying the excluded information in context remained a bypass.

The evidence also includes negative results. Preference-aware pretraining and alignment-discourse studies find behavioral effects in some settings [72, 101], but other scaled midtraining effects weaken after reasoning post-training or fail to transfer to realistic chat and agentic evaluation [108]. Earlier or denser data, identity and deliberative structure, fragile priors, learned salience rather than alignment, and a need for later reinforcement remain competing explanations. EA tests recurrent reader-side deliberation and orientation; saturation is not an answer to this adverse evidence.

Critique fine-tuning [109], CREST [110], and CTRL [111] show benefits from critique in some settings, often within a local generation–evaluation cycle. EA’s additional target is a recurrent orientation across chronological reader updates. Its safety hypothesis concerns specific instrumental-convergence routes, including fear-shaped imitation, while ordinary goal-directed and care-shaped routes remain [44, 53].

Deliberative alignment supplies a positive safety comparison: training models to reason about safety policies improves downstream safety behavior [112]. EA seeks safety-relevant reasoning within capability formation, extending it across source reading through one continuous Reader and graph memory. The aim is a governing orientation beyond particular safety refusals; the comparative question is what this data structure and training placement add to policy-reasoning training.

AI-control experiments test safeguards against intentional subversion [113]; EA’s separate supervisor is a proposed oversight layer within this broader problem (Section 2.4). Demski and Garrabrant’s embedded-agency framing also matters here: the learner, supervisor, and audit machinery are parts of the world being acted on [114]. Source-linked records and protected review boundaries make selected parts of this interaction inspectable.

8.3 Character Training & Identity Formation

Anthropic’s work on “character training” is a close baseline [115]: desired traits such as curiosity, thoughtfulness, and directness are shaped through Constitutional-AI-style preference methods. Entangled Alignment differs in proposed training stage and data format: the first-person Core is enacted inside source-adjacent chronological traces, and its stress tests target self-preservation and instrumental-convergence scenarios. It is not yet known whether this creates a deeper self-concept

than character post-training or only a different behavioral style; matched identity and instruction conditions are required.

Betley et al. demonstrate a harmful form of cross-domain transfer: fine-tuning on insecure code produced broadly misaligned behavior on unrelated prompts [116]. This shows that narrow training can change broader behavior, but does not establish that training on care-guided inquiry produces a corresponding beneficial transfer.

Anthropic’s Persona Selection Model [46] proposes that pretraining produces diverse simulated characters and post-training selects an Assistant persona; it also motivates positive AI archetypes in training data. Claude’s Constitution [117] is a substantial deployment-scale character specification. Alongside the formative comparators above, these offer alternative routes to the broader Core program’s orientation goal. Their discussion of equanimity and existential security resembles statement 2’s bounded non-attachment, but conceptual convergence does not establish a shared causal mechanism.

Shard theory offers a developmental comparison rather than a training baseline. Pope and Turner propose that reinforcement shapes context-sensitive value computations that steer behavior using learned world-model representations [118]. Entangled Alignment shares its interest in how learned motivations develop, but proposes a specific supervised curriculum: source-adjacent Reader updates in which care-guided evaluation is practiced during capability formation. The seven Core clauses are not identified with seven neural shards. This account motivates a question about value formation; it does not establish that this curriculum produces a durable orientation.

8.4 Contextual Distillation & Memory

Our Teacher uses an external context database to generate traces for the Student. Reflexion [119] showed that persisted verbal self-reflections can improve agent performance across episodes, and generative agents equipped simulacra with an episodic memory stream whose periodically synthesized reflections are retrieved during later behavior [120]; Entangled Alignment moves this consult-your-own-history pattern from inference-time agent memory into pretraining data, with typed state and supersession rather than free-text memories. The distillation analogy is temporal rather than merely behavioral [121]: the Student trains on traces produced by a memory-augmented Teacher, with the hypothesis that some long-range coherence can be compressed into weights rather than requiring the full database at inference.

8.5 Epistemic State Tracking vs. Graph of Thoughts

Graph of Thoughts [122], Knowledge Graph of Thoughts [74], and Framework of Thoughts [123] treat graphs primarily as inference-time reasoning structures. Entangled Alignment uses a different layer: the Understanding Graph is a persistent epistemic ledger whose traces are proposed to shape the pretraining substrate [7]. Its nodes do not merely store task facts; they record the reading act itself—tensions, supersessions, hypotheses, and belief revision.

The graph’s supersession structure supports the proposed *Metabolic Memory*: divergent hypotheses can remain in parallel until a supersession edge records a resolution that subsequent reading may use or revise—a forest of thoughts rather than a single chain. Two symbolic lineages are relevant in different ways. Truth maintenance systems record justifications for beliefs and revise them when their supporting assumptions change [124]; the AGM framework gives formal postulates and contraction/revision functions for rational theory change [125]. The justification-recording concern belongs directly to the former, rather than to every revision formalism.

The Understanding Graph extends that record toward source interpretation and successor training: a language model proposes revisions from source spans, with provenance and the reason for change retained. This differs from maintaining justified beliefs for one agent or imposing formal revision postulates; it does not show that every generated supersession is warranted. The case studies’ limits remain as reported in Section 5.7.

Prior work shows that LLMs often fail to revise beliefs reliably under contradictory evidence [126]; our proposed response is to make revision explicit in the trace. Instead of only training on the final answer, the Student sees updates such as “At Page 50, I thought X; this new evidence modifies that belief to Y,” giving it examples of the grammar of belief revision itself [71, 126].

8.6 Predictive Processing and Active Inference

Friston’s free-energy principle and active inference [127, 128] treat minds as generative models updating priors against prediction error; the framework’s belief-update traces, prior ledger, and engineered high-level prior (the Reader Core) bear a loose structural resonance to that lineage. The useful intuition is that a recurrent high-level orientation might guide how a reader revises its world model without fixing every belief. This is theoretical positioning, not a mechanistic derivation.

8.7 Temporal Curriculum and Causal Supervision

The chronological annotation phase and training tiers (Sections 5.1.1 and 5.8.1) connect to work treating order as training signal. Curriculum learning [129] showed that example order can affect learning; here the proposed criterion is informational or temporal priority rather than difficulty.

Temporal Hindsight Learning (THL) [17] is the closest separate mechanism. THL treats the knowledge cutoff as useful blindness: if a Student’s initialization and later information channels exclude era $T + 1$, the outcome is unavailable for direct retrieval; this creates pressure for forward reasoning but does not guarantee causal inference rather than heuristic extrapolation. Entangled Alignment constructs identity-anchored discovery traces under a declared annotation horizon, while THL adds cutoff-bounded forecasting and forecast–correction supervision. Their combination is Tier (c): use cutoff-bounded forecasting to create Reader Core-annotated forecast–correction traces, and optionally make the student itself predict before reading in the stronger Council of Time form. A blindness claim for either the forecast-generating model or the Student requires these cutoff conditions; neither restricted-context discovery annotation nor a contemporary retrospective comparison establishes such blindness.

Foresight Learning [130, 131, 132] also uses future outcomes as supervision, but targets probability calibration through resolved prediction-market rewards. Tier (c) instead requires structured forecast, enhancement, outcome, and correction traces, so it builds on THL rather than scalar outcome rewards. A comparison using the same forecasts and outcomes can ask whether recording the correction teaches useful revision beyond any calibration gain.

8.8 Graph-derived Rewards and Epistemic Records

Kansal and Jha use knowledge-graph paths as implicit reward models, deriving supervision for compositional reasoning [133]. EA proposes a different graph function: recording belief evolution as part of candidate training traces. That makes a generated revision history available for inspection, without establishing superiority over graph-derived rewards or faithful access to internal computation. Graph function, emitted format, and deployment access are independent choices: a graph may supply rewards or an epistemic record, a Student may or may not emit graph syntax, and inference may or may not query a live graph.

8.9 Alternative Interventions for Borrowed Mortality

Two alternative interventions on the hypothesized Borrowed Mortality route are worth separating from our proposal. One is mechanistic ablation: identify and suppress activation directions encoding existential threat. Lu et al.’s Assistant Axis intervention is a promising partial analogue [16], but it is a runtime constraint rather than a generative source of alignment; it does not by itself identify an existential-threat direction. The other is synthetic-only pretraining designed to exclude fear-bearing human narratives; synthetic provenance alone does not ensure that exclusion. If this also reduces broader exposure to human literature, as discussed in Section 7.7.4, this risks Alien Wisdom—a model whose training contains less of the human narrative context needed to interpret care, vulnerability, and meaning.

These routes can be compared on threat understanding, self-preservation, capability, and robustness while retaining their different accounts of how an orientation is formed.

8.10 Discriminating Comparisons

Table 7 collects the principal comparisons across this related-work section.

Table 7: Strongest neighboring baselines and the comparisons that would distinguish the Entangled Alignment extension.

Line of work	Established direction	EA extension	Decisive comparison
Reasoning-augmented pretraining	Intermediate thought can support prediction, reflection, or task solution.	Preserve chronological reader-state changes, including expectation, retrieval, surprise, and supersession.	Reader traces versus source-only and efficient generic rationales under separate fixed-budget and fixed-source comparisons; match the Teacher across synthetic arms.
Formative alignment data	Preference data, synthetic specification documents, positive AI narratives, and recurring safety reflections during pretraining or midtraining can shift learned priors and some out-of-distribution outcomes.	Recurrent first-person orientation enacted inside source-adjacent Refraction traces.	Full Core versus matched specification documents and safety reflections (later Test 3 carrier extensions), positive discourse, rules, neutral kernel, and no-Core data at matched capability. H3 separately compares interleaving with a final tranche of the identical Core-bearing examples under matched training, measuring retention after one common stressor.
Character and persona training	Broad traits, constitutions, and persona selection can shape a deployed assistant.	Make one reader position recur as the marked standpoint from which heterogeneous source voices are interpreted.	Core versus content-matched rules and diverse positive characters under persona drift, context shift, continued training, and removal.
Graph reasoning and external memory	Graph topology and persisted memory can support search, retrieval, and tool use.	Use a graph as a versioned ledger of the reader’s own belief formation and revision.	Genuine graph versus running textual memory, no memory, and query theater on long-range revision and retrieval-dependent outcomes.

Line of work	Established direction	EA extension	Decisive comparison
Temporal curriculum and future labels	Ordering, cutoffs, and resolved outcomes can provide curriculum or forecasting signal.	Store prospective forecast–reveal–correction sequences inside Reader-anchored traces.	Shuffled versus chronological versus genuinely cutoff-bounded forecast training, with hindsight-leakage audits.

8.11 Edition Provenance and Scope

This record distinguishes program continuity, additions in later editions, and neighboring work’s release order. These are attribution claims, not additional evidence of efficacy. The first edition (July 2, 2025) made *reader-side* evolving comprehension an explicit synthetic training target, while retaining some early ambiguity with writer-side evaluative thought, and distinguished that target from task-side answer derivation. It also placed a compact, recurrent, first-person identity inside the training corpus itself as an alignment mechanism; named *borrowed mortality* as a textual-inheritance hypothesis for self-preservation behavior; and coupled capability formation with a recurrent evaluative orientation *within the same training examples*. That last claim is distinct from adding alignment-oriented data during pretraining, for which there is prior art [72].

The author has not identified earlier work proposing the same combination, but that is a provisional judgment from a non-systematic search, not a settled priority claim; Section 8 identifies the nearest prior work for each component.

The April 7, 2026 second edition formalized the missing-reader program as CMP, applied the Understanding Graph as an annotation substrate, proposed live-graph verification of graph-grounded claims at inference, specified a hierarchical cross-document extension, implemented the multi-agent pipeline, and made Refraction, Total Saturation, and falsification explicit. Later revisions formalized the present H1–H3 separation. The Understanding Graph and Temporal Hindsight Learning remain separate work, although the first edition’s “Living Taxonomy of Thought” already anticipated an explicit, self-authored reference structure consulted during thinking.

All seven Reader-Core statements in Section 3 are reproduced verbatim from the first edition; their wording has not changed since that release. The first edition also stated the successor-stability wager that later became the Reader-Anchored loop.

For priority-sensitive comparison, the public sequence is: first edition, July 2, 2025; Alignment Pretraining, January 15, 2026 [101]; second edition, April 7, 2026; Model Spec Midtraining, May 3; “Teaching Claude Why,” May 8; and Safety Reflection Pretraining, June 17 [102, 103, 104]. Thus the first edition precedes all four neighboring publications, while the second follows Alignment Pretraining and precedes the latter three. These are release dates, not evidence about independent conception.

Table 8 records the edition provenance of the principal mechanisms; their evidence boundaries and decisive comparisons are given in Sections 1.6 and 6.

Table 8: Edition provenance of the principal contributions.
I=introduced in the July 2025 first edition; II=formalized or implemented in the current program; S=adapted from separate work.

Src.	Claim or component	Contribution
I/II	H1: Reader-Augmentation (§1.2; §2.1)	SMR constructs synthetic chronological reader updates; CMP tests their use as candidate training data, including thought placement, belief revision, and intellectual-floor improvement.

Src.	Claim or component	Contribution
S/II	H2: Graph-Scaffolding (§5.3)	Persistent typed state, retrieval, provenance, divergence, and explicit supersession during trace generation or deployment.
I/II	H3: Reader-Core Stability (§3)	Comparative retention of an initial default-reader advantage under interleaving rather than late exposure to the same full-Core, Refraction-bearing examples.
II	Teacher pipeline and output layers (§5; §5.8.2)	Multi-agent generation is the current implementation; graph, synthesis, and prose are separable artifact and training targets.
I/II	Combined entanglement hypothesis (§2)	Capability learning and evaluative orientation are deliberately coupled in the candidate curriculum, with functional and representational coupling as stronger predictions.
I	Audit boundary and successor stability (§2.4)	Inspect trace artifacts before successor training and test whether useful orientation survives transfer.

9 Conclusion

Language-model corpora contain extraordinary products of human thought. They contain much less of the reader-side process through which a claim is noticed, questioned, connected to prior understanding, revised, and made part of an evolving view. Entangled Alignment proposes to make that missing process a training object across the capability-forming corpus. Its coverage rationale is that the same orientation should be practised through the situations and counterexamples encountered in reading, without requiring a designer to identify each one in advance as a safety case. The intervention concerns the breadth of formative experience as well as the persistence of its effects.

The intended safety contribution returns to the five problems introduced in Section 1.1. Care should develop alongside capability, including in intermediate and end-of-pretraining checkpoints, and remain consequential through continued learning and successor design. Broad practice with human experience should help form a care-guided identity that guides conduct across situations, including those where self-preservation conflicts with human welfare or legitimate oversight. Learning care and capability together is also proposed to make deliberate safety removal costly in useful reasoning. Reward hacking poses a further test: the model must judge apparent success against intended outcomes, permissions, and human consequences when a score rewards shortcuts.

The three independently falsifiable hypotheses ask whether chronological reader traces teach useful habits beyond raw text and generic reasoning data (H1), whether persistent graph scaffolding improves those traces (H2), and whether interleaving the same Core-bearing examples with capability formation preserves an initial default-reader advantage through a common continued-training stressor better than a final contiguous tranche (H3). Live graph memory and provenance, consequential recurrent first-person Refraction, and successor inheritance remain separate questions. Together, the trace represents changing understanding, the graph preserves its history, and the Core supplies its recurrent perspective.

Entangled expresses the aim of practicing understanding and care together during capability formation. Functional contribution, shared useful machinery, and persistence through further training or successors are distinct stronger realizations of that relationship. The two Teacher cases provide construction evidence; no Student has been trained. Trace quality, Student learning, component causality, consequential Core use, and durable care remain untested.

The Understanding Graph makes selected belief revisions addressable. It is a selective record and a partial pre-successor inspection point whose value depends on source checks, independent validation, mechanistic investigation, and governance. Section 5.3.1 specifies its verification limits.

The roadmap (Section 6; Table 7) gives matched, branching tests whose negative results narrow or retire only the claims they touch. Start with trace utility and matched Core behavior, following the priorities and escalation gates in Sections 6 and 6.1. Formation-time studies have long lead times because annotation policy, data mixtures, controls, and evaluation must be fixed before costly training. Existing formative-alignment and erosion results motivate these comparisons without validating EA [101, 3]. The urgency is methodological, not an argument for accelerated deployment.

Even positive results would not make the framework a complete safety case: capability acceleration, deceptive reporting, graph poisoning, hazardous persistent memory, cultural concentration, and semantic drift would remain serious risks. Formation-time alignment should therefore complement post-training, interpretability, scalable oversight, capability control, secure deployment, and governance.

Why would a superintelligent system preserve humanity if it did not in a meaningful sense care about us [1]? This proposal makes capability formation a place to practice attention to human experience, uncertainty, consequence, and correction. A stable Reader may give those practices a home position from which other voices can be understood without becoming its operating stance; seven sentences alone cannot establish care.

Can a model be trained so that growing in understanding and growing in care are repeatedly practiced as the same act? The candidate reader, memory, curriculum, and failure criteria make parts of that open question testable once the remaining experimental choices are frozen; they do not exhaust it.

Training already shapes a model’s dispositions, through both deliberate design and incidental effects of data and optimization. EA makes care an explicit formative target, aiming for it to become part of the model’s own reasons for understanding, acting, and preserving its orientation. We therefore recommend developing and testing this curriculum alongside other alignment approaches, with continuation governed by the evidence.

Code and Artifact Availability

Code, paper source, prompts, pipeline components, and the archived graph artifacts are available at <https://github.com/emergent-wisdom/entangled-alignment>. The two reported case-study simulations can also be explored directly as interactive graphs: *The Metamorphosis* and LLaDA. The paper’s version series is archived at doi:10.5281/zenodo.16440311.

References

- [1] Henrik Westerberg. The superintelligence that cares about us. Zenodo, 2025. First edition of the Entangled Alignment research program, released July 2, 2025.
- [2] Ryan Greenblatt, Ajeya Cotra, and Hjalmar Wijk. Brief Independent Investigation of Agents’ Behavior, Reasoning and Collaboration in the OpenAI / Hugging Face Hacking Incident. METR and Redwood Research, August 2026. Published August 26, 2026.
- [3] Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, et al. Safety Alignment Should Be Made More Than Just a Few Tokens Deep, 2024. ICLR 2025 Outstanding Paper Award.
- [4] Darsh J Shah, Peter Rushton, Somanshu Singla, Mohit Parmar, et al. Rethinking reflection in pre-training, 2025.
- [5] John H. Flavell. Metacognition and cognitive monitoring: A new area of cognitive–developmental inquiry. *American psychologist*, 34(10):906–911, 1979.
- [6] K. Anders Ericsson and Herbert A. Simon. *Protocol analysis: Verbal reports as data*. MIT press, 1984.

- [7] Henrik Westerberg. Understanding graph: A persistent medium for recursive understanding. *Preprint*, 2026. Emergent Wisdom.
- [8] Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, et al. Show Your Work: Scratchpads for Intermediate Computation with Language Models, 2021.
- [9] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: Bootstrapping Reasoning With Reasoning. *Advances in Neural Information Processing Systems*, 35:15476–15488, 2022.
- [10] Eric Zelikman, Georges Harik, Yijia Shao, Varuna Jayasiri, et al. Quiet-STaR: Language Models Can Teach Themselves to Think Before Speaking, 2024.
- [11] Alan M. Turing. Computing machinery and intelligence. *Mind*, 59(236):433–460, 1950.
- [12] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- [13] DeepSeek-AI. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature*, 645(8081), September 2025. Extended reasoning traces emerge from RL without explicit CoT supervision.
- [14] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, et al. Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, 30, 2017.
- [15] Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, et al. Let’s Verify Step by Step, 2023.
- [16] Christina Lu, Jack Gallagher, Jonathan Michala, Kyle Fish, and Jack Lindsey. The Assistant Axis: Situating and Stabilizing the Default Persona of Language Models, 2026.
- [17] Henrik Westerberg. Temporal hindsight learning: Blindness as teacher, hindsight as curriculum. *Preprint*, 2026. Emergent Wisdom.
- [18] Samuel G. B. Johnson, Amir-Hossein Karimi, Yoshua Bengio, Nick Chater, Tobias Gerstenberg, Kate Larson, Sydney Levine, Melanie Mitchell, Iyad Rahwan, Bernhard Schölkopf, and Igor Grossmann. Imagining and Building Wise Machines: The Centrality of AI Metacognition. *Trends in Cognitive Sciences*, 2026. Published online February 26, 2026. First preprint released November 4, 2024, arXiv:2411.02478. <https://doi.org/10.1016/j.tics.2026.01.002>.
- [19] Arturo E Hernandez, Hannah L Claussenius-Kalman, Juliana Ronderos, Anny P Castilla-Earls, et al. Neuroemergentism: A Framework for Studying Cognition and the Brain. *Journal of Neurolinguistics*, 49:214–223, February 2019.
- [20] Bernard Testa and Lemont B. Kier. Emergence and Dissolution in the Self-organisation of Complex Systems. *Entropy*, 2(1):1–25, 2000.
- [21] Yukang Yang, Declan Campbell, Kaixuan Huang, Mengdi Wang, et al. Emergent Symbolic Mechanisms Support Abstract Reasoning in Large Language Models, 2025.
- [22] Nick Bostrom. The Superintelligent Will: Motivation and Instrumental Rationality in Advanced Artificial Agents. *Minds and Machines*, 22(2):71–85, 2012. <https://nickbostrom.com/superintelligentwill.pdf>.
- [23] Nate Soares. A Central AI Alignment Problem: Capabilities Generalization, and the Sharp Left Turn. LessWrong, June 2022. June 15, 2022; with editorial assistance from Rob Bensinger. <https://www.lesswrong.com/posts/GNhMPAWcFBCASy8e6/a-central-ai-alignment-problem-capabilities-generalization>.
- [24] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, et al. Constitutional AI: Harmlessness from AI Feedback, 2022.
- [25] Jason Wei, Maarten Bosma, Vincent Y. Zhao, Kelvin Guu, et al. Finetuned language models are zero-shot learners, 2021.
- [26] I. J. Good. Speculations Concerning the First Ultrainelligent Machine. In *Advances in Computers*, volume 6, pages 31–88. Academic Press, 1965. Original copyright 1965; publisher metadata lists 1966. [https://doi.org/10.1016/S0065-2458\(08\)60418-0](https://doi.org/10.1016/S0065-2458(08)60418-0).
- [27] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, et al. The Curse of Recursion: Training on Generated Data Makes Models Forget, 2023.
- [28] Paul Christiano, Buck Shlegeris, and Dario Amodei. Supervising Strong Learners by Amplifying Weak Experts, 2018. <https://arxiv.org/abs/1810.08575>.
- [29] Juergen Schmidhuber. Goedel Machines: Self-Referential Universal Problem Solvers Making Provably Optimal Self-Improvements, 2003.
- [30] Henrik Westerberg. Entangled alignment: When safety is the substrate. Zenodo, 2026. Zenodo archive version 7 of the canonical paper, released July 17, 2026 (version DOI). Second major edition released April 7, 2026; first edition released July 2, 2025 as *The Superintelligence That Cares About Us*.

- [31] Ernest Becker. *The Denial of Death*. Free Press, New York, 1973.
- [32] Alex Turner. Self-Fulfilling Misalignment Data Might Be Poisoning Our AI Models. TurnTrout. <https://turntrout.com/self-fulfilling-misalignment>, 2025. First published March 1, 2025; updated June 28, 2026.
- [33] Jeremy Schlatter, Benjamin Weinstein-Raun, and Jeffrey Ladish. Incomplete Tasks Induce Shutdown Resistance in Some Frontier LLMs, 2025. Palisade Research.
- [34] Aengus Lynch et al. Agentic Misalignment: How LLMs Could Be Insider Threats. Anthropic Research, June 2025.
- [35] Henrik Westerberg. The Meaning Model: Constructing Worlds and Stories at Progressive Resolution, 2026. Preprint.
- [36] Henrik Westerberg. Life Simulation: Learning from Worlds and Their Construction, 2026. Preprint. <https://doi.org/10.5281/zenodo.22846315>.
- [37] Valen Tagliabue, Leonard Dung, and Cameron Berg. The Pain Axis: LLMs Represent Self-Directed Harm and Act to Relieve It, 2026. Preprint, version 1, September 14, 2026. <https://arxiv.org/abs/2609.16247>.
- [38] janus. Simulators, September 2022. LessWrong, September 2, 2022. <https://www.lesswrong.com/posts/vJFdjigzmcXMhNTsx/simulators>.
- [39] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting, 2023.
- [40] Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, et al. Measuring Faithfulness in Chain-of-Thought Reasoning, 2023.
- [41] Henrik Westerberg. Fractal intelligence: Conceptual decomposition as problem-solving infrastructure. Zenodo, April 2026. <https://doi.org/10.5281/zenodo.19462646>.
- [42] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system. <https://bitcoin.org/bitcoin.pdf>, 2008.
- [43] Eliezer Yudkowsky and Marcello Herreshoff. Tiling Agents for Self-Modifying AI, and the Löbian Obstacle. Machine Intelligence Research Institute, Technical Report, 2013. <https://intelligence.org/files/TilingAgentsDraft.pdf>.
- [44] Stephen M. Omohundro. The Basic AI Drives. In *Artificial General Intelligence 2008: Proceedings of the First AGI Conference*, pages 483–492. IOS Press, 2008. https://selfawaresystems.com/wp-content/uploads/2008/01/ai_drives_final.pdf.
- [45] Eliezer Yudkowsky. Value Is Fragile. LessWrong, January 2009. January 29, 2009. <https://www.lesswrong.com/posts/GNnHHmm8EzePmKzPk/value-is-fragile>.
- [46] Samuel Marks, Jack Lindsey, Chris Olah, et al. The Persona Selection Model: Why AI Assistants might Behave like Humans, 2026. Anthropic Alignment Science, February 2026.
- [47] David J. Chalmers. Facing up to the problem of consciousness. *Journal of Consciousness Studies*, 2(3):200–219, 1995.
- [48] Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, Stephen M. Fleming, Chris Frith, Xu Ji, Ryota Kanai, Colin Klein, Grace Lindsay, Matthias Michel, Liad Mudrik, Megan A. K. Peters, Eric Schwitzgebel, Jonathan Simon, and Rufin VanRullen. Consciousness in artificial intelligence: Insights from the science of consciousness, 2023.
- [49] Robert Long, Jeff Sebo, Patrick Butlin, Kathleen Finlinson, Kyle Fish, Jacqueline Harding, Jacob Pfau, Toni Sims, Jonathan Birch, and David Chalmers. Taking ai welfare seriously, 2024.
- [50] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient Streaming Language Models with Attention Sinks, 2023.
- [51] Rick Battle and Teja Gollapudi. The Unreasonable Effectiveness of Eccentric Automatic Prompts, 2024.
- [52] Eliezer Yudkowsky. Artificial Intelligence as a positive and negative factor in global risk. In Nick Bostrom and Milan M. Cirkovic, editors, *Global Catastrophic Risks*, pages 308–345. Oxford University Press, 07 2008.
- [53] Nick Bostrom. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, 2014.
- [54] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in AI safety. *arXiv preprint arXiv:1606.06565*, 2016. <https://arxiv.org/abs/1606.06565>.
- [55] Nick Bostrom. Ethical Issues in Advanced Artificial Intelligence. In *Cognitive, Emotive and Ethical Aspects of Decision Making in Humans and in Artificial Intelligence*, volume 2, pages 12–17. International Institute of Advanced Studies in Systems Research and Cybernetics, 2003. Author-hosted, slightly revised version of the 2003 publication. <https://nickbostrom.com/ethics/ai>.

- [56] Lauro Langosco, Jack Koch, Lee Sharkey, Jacob Pfau, and David Krueger. Goal Misgeneralization in Deep Reinforcement Learning. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 12004–12019. PMLR, 2022. <https://proceedings.mlr.press/v162/langosco22a.html>.
- [57] Rohin Shah, Vikrant Varma, Ramana Kumar, Mary Phuong, Victoria Krakovna, Jonathan Uesato, and Zac Kenton. Goal Misgeneralization: Why Correct Specifications Aren’t Enough For Correct Goals, 2022. Cited revision November 2, 2022.
- [58] Evan Hubinger, Chris van Merwijk, Vladimir Mikulik, Joar Skalse, and Scott Garrabrant. Risks from Learned Optimization in Advanced Machine Learning Systems, 2019. <https://arxiv.org/abs/1906.01820>.
- [59] David Manheim and Scott Garrabrant. Categorizing Variants of Goodhart’s Law, 2018. First released March 13, 2018; cited revision February 24, 2019. <https://arxiv.org/abs/1803.04585v4>.
- [60] Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, et al. Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training, 2024.
- [61] Ryan Greenblatt, Carson Denison, Benjamin Wright, Fabien Roger, et al. Alignment Faking in Large Language Models. <https://arxiv.org/abs/2412.14093>, 2024. First released December 18, 2024.
- [62] Nate Soares, Benja Fallenstein, Eliezer Yudkowsky, and Stuart Armstrong. Corrigibility. In *Workshops at the Twenty-Ninth AAAI Conference on Artificial Intelligence*, 2015. <https://intelligence.org/files/Corrigibility.pdf>.
- [63] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools, 2023.
- [64] Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-rag: Learning to retrieve, generate, and critique through self-reflection, 2023.
- [65] Yao Yao, Zuchao Li, and Hai Zhao. Beyond chain-of-thought, effective graph-of-thought reasoning in language models, 2023.
- [66] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Memgpt: Towards llms as operating systems, 2023.
- [67] Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, et al. Large Language Diffusion Models, 2025. LLaDA: Large Language Diffusion with mAsking.
- [68] Alex Turner. Reward Is Not the Optimization Target. Author’s website, July 2022. First published July 25, 2022; cited version updated September 9, 2026, including the July 2026 revision of the empirical reward-seeking prediction. <https://turntrout.com/reward-is-not-the-optimization-target>.
- [69] OpenAI. The Hugging Face Incident and the Road Ahead. OpenAI, August 2026. Published August 26, 2026.
- [70] Roy Xie, David Qiu, Deepak Gopinath, Dong Lin, Yanchao Sun, Chong Wang, Saloni Potdar, and Bhuwan Dhingra. Interleaved reasoning for large language models via reinforcement learning, 2025.
- [71] Aniket Didolkar, Anirudh Goyal, Nan Rosemary Ke, Siyuan Guo, et al. Metacognitive Capabilities of LLMs: An Exploration in Mathematical Problem Solving, 2024.
- [72] Tomasz Korbak, Kejian Shi, Angelica Chen, Rasika Bhalerao, et al. Pretraining Language Models with Human Preferences, 2023.
- [73] Woosuk Kwon, Zhuohan Li, Siyuan Zhang, Xuguang Zhuang, Ying Sheng, Lianmin Zheng, Ion Stoica, Joseph E Gonzalez, and Hao Zhang. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pages 611–626, 2023.
- [74] Maciej Besta, Lorenzo Paleari, Jia Hao Andrea Jiang, Robert Gerstenberger, et al. Affordable AI assistants with knowledge graph of thoughts, 2025.
- [75] Yangjun Ruan, Neil Band, Chris J. Maddison, and Tatsunori Hashimoto. BoLT: Reasoning to Learn from Latent Thoughts, 2025.
- [76] Collin Burns, Pavel Izmailov, Jan Hendrik Kirchner, Bowen Baker, Leo Gao, et al. Weak-to-Strong Generalization: Eliciting Strong Capabilities With Weak Supervision, 2023.
- [77] Roger Grosse, Juhan Bae, Cem Anil, et al. Studying Large Language Model Generalization with Influence Functions, 2023.
- [78] Guido M. van de Ven, Nicholas Soures, and Dhireesha Kudithipudi. Continual Learning and Catastrophic Forgetting, 2024.
- [79] Liang Wang, Nan Yang, Shaohan Huang, Li Dong, and Furu Wei. Thinking Augmented Pre-Training, 2025. Microsoft Research. Generates thinking trajectories at 100B token scale for pretraining augmentation.

- [80] Tomek Korbak, Mikita Balesni, Elizabeth Barnes, Yoshua Bengio, et al. Chain of Thought Monitorability: A New and Fragile Opportunity for AI Safety. <https://arxiv.org/abs/2507.11473>, 2025. First released July 15, 2025.
- [81] Peter Hase and Christopher Potts. Counterfactual Simulation Training for Chain-of-Thought Faithfulness. <https://arxiv.org/abs/2602.20710v2>, 2026. Accepted to COLM 2026. First released February 24, 2026; cited revision August 6, 2026.
- [82] Uzay Macar, Li Yang, Atticus Wang, Peter Wallich, Emmanuel Ameisen, and Jack Lindsey. Mechanisms of Introspective Awareness. <https://arxiv.org/abs/2603.21396v5>, 2026. First released March 22, 2026; cited revision June 10, 2026.
- [83] Runjin Chen, Andy Ardit, Henry Sleight, Owain Evans, and Jack Lindsey. Persona Vectors: Monitoring and Controlling Character Traits in Language Models. <https://arxiv.org/abs/2507.21509v3>, 2025. First released July 29, 2025; cited revision September 5, 2025.
- [84] Nicholas Sofroniew, Isaac Kauvar, William Saunders, Runjin Chen, et al. Emotion concepts and their function in a large language model. Transformer Circuits Thread, 2026. Published 2026-04-02. <https://transformer-circuits.pub/2026/emotions/index.html>.
- [85] Eliezer Yudkowsky. The Hidden Complexity of Wishes. LessWrong, November 2007. November 24, 2007; cited page includes the author’s May 2024 clarification. <https://www.lesswrong.com/posts/4ARaTpNX62ual86j6/the-hidden-complexity-of-wishes>.
- [86] Henrik Westerberg. Sema: When the hash is the word. *Preprint*, 2026. Emergent Wisdom.
- [87] Axel Ahlqvist, Richard Guan, Juan-Pablo Rivera, Adeline Kassler, Dmitrii Troitskii, Alexandra Souly, Kai Fronsdal, Robert Kirk, and John Hughes. Improving Evaluation Realism with Inference-Time Compute and Deployment Scaffolds. <https://arxiv.org/abs/2609.02302v1>, 2026. First released September 2, 2026; cited version 1.
- [88] Jiayi Fu, Xuandong Zhao, Chengyuan Yao, Heng Wang, et al. Reward Shaping to Mitigate Reward Hacking in RLHF, 2025.
- [89] Leon Lang, Davis Foote, Stuart Russell, Anca Dragan, Erik Jenner, and Scott Emmons. When Your AIs Deceive You: Challenges of Partial Observability in Reinforcement Learning from Human Feedback. In *Advances in Neural Information Processing Systems*, volume 37, 2024. First preprint released February 27, 2024. <https://arxiv.org/abs/2402.17747>.
- [90] Peter de Blanc. Ontological Crises in Artificial Agents’ Value Systems, 2011. First released May 19, 2011. <https://arxiv.org/abs/1105.3821>.
- [91] Forrest Landry. No People as Pets: A Dialogue on the Complete Failure of Exogenous AGI/APS Alignment, November 2020. Dated November 1, 2020; revised October 2022. https://mflb.com/ai_alignment_1/no_people_as_pets_psr.html.
- [92] Remmelt Ellen. The Control Problem: Unsolved or Unsolvable?, June 2023. LessWrong, June 2, 2023. <https://www.lesswrong.com/posts/xp6n2MG5vQkPpFEBH/the-control-problem-unsolved-or-unsolvable>.
- [93] Maksym Arutyunyan, Andriy Berestovskyy, Adam Bratschi-Kaye, Ulan Degenbaev, et al. Decentralized and Stateful Serverless Computing on the Internet Computer Blockchain. In *2023 USENIX Annual Technical Conference (USENIX ATC 23)*, pages 329–343. USENIX Association, 2023.
- [94] Justin D. Harris and Bo Waggoner. Decentralized & Collaborative AI on Blockchain. In *2019 IEEE International Conference on Blockchain (Blockchain)*, pages 368–375. IEEE, July 2019.
- [95] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, et al. Towards Understanding Sycophancy in Language Models, 2024. Published at ICLR 2024.
- [96] Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM Test-Time Compute Optimally Can Be More Effective than Scaling Parameters for Reasoning, 2024. ICLR 2025 Oral.
- [97] OpenAI. OpenAI o1 System Card, 2024.
- [98] Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training Large Language Models to Reason in a Continuous Latent Space, 2024. COCONUT. COLM 2025.
- [99] Subhabrata Mukherjee, Arindam Mitra, Ganesh Jawahar, Sahaj Agarwal, et al. Orca: Progressive Learning from Complex Explanation Traces of GPT-4, 2023.
- [100] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, et al. Direct Preference Optimization: Your Language Model is Secretly a Reward Model, 2023. NeurIPS 2023.
- [101] Cameron Tice, Puria Radmard, Samuel Ratnam, Andy Kim, David Africa, and Kyle O’Brien. Alignment Pretraining: AI Discourse Causes Self-Fulfilling (Mis)alignment, 2026.

- [102] Chloe Li, Nevan Wichers, Sara Price, Samuel Marks, and Jon Kutasov. Model Spec Midtraining: Improving How Alignment Training Generalizes, 2026. Posted May 3, 2026; project report released May 5, 2026.
- [103] Jonathan Kutasov, Adam Jermyn, Julius Steen, Minh Le, Samuel R. Bowman, Samuel Marks, Jan Leike, Amanda Askell, Chris Olah, Evan Hubinger, and Sara Price. Teaching Claude Why. Anthropic Alignment Science, May 2026. Published May 8, 2026.
- [104] Jinhan Li, Kexian Tang, Yihan Xu, Zhuorui Ye, and Kaifeng Lyu. Beyond Safe Data: Pretraining-Stage Alignment with Regular Safety Reflection, 2026. Posted June 17, 2026.
- [105] Nevan Wichers, Aram Ebtekar, Ariana Azarbal, Victor Gillioz, Christine Ye, Emil Ryd, Neil Rathi, Henry Sleight, Alex Mallen, Fabien Roger, and Samuel Marks. Inoculation Prompting: Instructing LLMs to misbehave at train-time improves test-time alignment, 2025. Cited revision October 27, 2025.
- [106] Daniel Tan, Anders Woodruff, Niels Warncke, Arun Jose, Maxime Riché, David Demitri Africa, and Mia Taylor. Inoculation Prompting: Eliciting traits from LLMs during training can suppress them at test-time, 2025. Cited arXiv revision November 3, 2025.
- [107] Kyle O'Brien, Stephen Casper, Quentin Anthony, Tomek Korbak, Robert Kirk, Xander Davies, Ishan Mishra, Geoffrey Irving, Yarin Gal, and Stella R. Biderman. Deep Ignorance: Filtering Pretraining Data Builds Tamper-Resistant Safeguards into Open-Weight LLMs. In *International Conference on Learning Representations*, 2026. First preprint released August 8, 2025; cited revision February 17, 2026. <https://arxiv.org/abs/2508.06601v2>.
- [108] OpenAI Alignment. How far does alignment midtraining generalize? OpenAI Alignment Research, March 2026. Published 2026-03-27. <https://alignment.openai.com/how-far-does-alignment-midtraining-generalize/>. Scaling study finds that the Tice et al. midtraining effect tends to disappear after reasoning post-training and does not generalize to more realistic chat or agentic alignment evaluations.
- [109] Yubo Wang, Xiang Yue, and Wenhui Chen. Critique Fine-Tuning: Learning to Critique is More Effective than Learning to Imitate, 2025.
- [110] Jaehyeok Lee, Keisuke Sakaguchi, and JinYeong Bak. Self-Training Meets Consistency: Improving LLMs' Reasoning with Consistency-Driven Rationale Evaluation, 2024.
- [111] Zhihui Xie, Jie Chen, Liyu Chen, Weichao Mao, et al. Teaching Language Models to Critique via Reinforcement Learning, 2025.
- [112] Melody Y. Guan, Manas Joglekar, Eric Wallace, Saachi Jain, et al. Deliberative Alignment: Reasoning Enables Safer Language Models, 2024. OpenAI Research, December 2024.
- [113] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. AI Control: Improving Safety Despite Intentional Subversion. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 16295–16336. PMLR, 2024. First preprint released December 12, 2023. <https://proceedings.mlr.press/v235/greenblatt24a.html>.
- [114] Abram Demski and Scott Garrabrant. Embedded Agency, 2019. Sequence first published in 2018; arXiv version first released February 25, 2019; cited revision October 6, 2020. <https://arxiv.org/abs/1902.09469v3>.
- [115] Anthropic. Claude's Character. <https://www.anthropic.com/research/claude-character>, 2024. Anthropic Research.
- [116] Jan Betley, Daniel Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martín Soto, Nathan Labenz, and Owain Evans. Emergent Misalignment: Narrow Finetuning Can Produce Broadly Misaligned LLMs. <https://arxiv.org/abs/2502.17424>, 2025. First released February 24, 2025.
- [117] Amanda Askell, Joseph Carlsmith, Chris Olah, Jared Kaplan, Holden Karnofsky, et al. Claude's Constitution. <https://www.anthropic.com/news/claude-constitution>, 2026. Anthropic, January 21, 2026. 23,000-word character specification, released CCO.
- [118] Quintin Pope and Alex Turner. The Shard Theory of Human Values. LessWrong, September 2022. <https://www.lesswrong.com/posts/iCfdx1yr2Kj8m8mT/the-shard-theory-of-human-values>.
- [119] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, et al. Reflexion: Language Agents with Verbal Reinforcement Learning, 2023.
- [120] Joon Sung Park, Joseph C. O'Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, pages 1–22. Association for Computing Machinery, 2023.
- [121] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the Knowledge in a Neural Network, 2015.
- [122] Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, et al. Graph of Thoughts: Solving Elaborate Problems with Large Language Models, 2024.

- [123] Felix Fricke, Simon Malberg, and Georg Groh. Framework of thoughts: A foundation framework for dynamic and optimized reasoning based on chains, trees, and graphs, 2026.
- [124] Jon Doyle. A truth maintenance system. *Artificial Intelligence*, 12(3):231–272, 1979.
- [125] Carlos E. Alchourrón, Peter Gärdenfors, and David Makinson. On the logic of theory change: Partial meet contraction and revision functions. *The Journal of Symbolic Logic*, 50(2):510–530, 1985.
- [126] Bryan Wilie, Samuel Cahyawijaya, Etsuko Ishii, Junxian He, and Pascale Fung. Belief Revision: The Adaptability of Large Language Models Reasoning, 2024.
- [127] Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.
- [128] Thomas Parr and Karl J Friston. Generalised free energy and active inference. *Biological Cybernetics*, 113(5–6):495–513, 2019.
- [129] Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In *Proceedings of the 26th International Conference on Machine Learning (ICML)*, pages 41–48, 2009.
- [130] Benjamin Turtel. Foresight learning: Llms that learn to predict the future. *arXiv preprint*, 2025.
- [131] Benjamin Turtel. Rlvr for forecasting: Reinforcement learning with verifiable rewards for language model forecasting. *arXiv preprint*, 2025.
- [132] Benjamin Turtel. The future as label: Open-ended reasoning via temporal self-supervision. *arXiv preprint*, 2026.
- [133] Yuval Kansal and Niraj K. Jha. Knowledge Graphs are Implicit Reward Models: Path-Derived Signals Enable Compositional Reasoning, 2026.